

# Pipelined CPU with Caches and TLBs

***Seungwan Noh***

*Pusan National University*

*Department of Electronics Engineering*

# Contents

---

- Memory Hierarchy
- Virtual Memory and TLB
- TLB-Based MIPS Memory Management
- Design of Pipelined CPU with Caches and TLB in Verilog HDL
- Simulation
- References

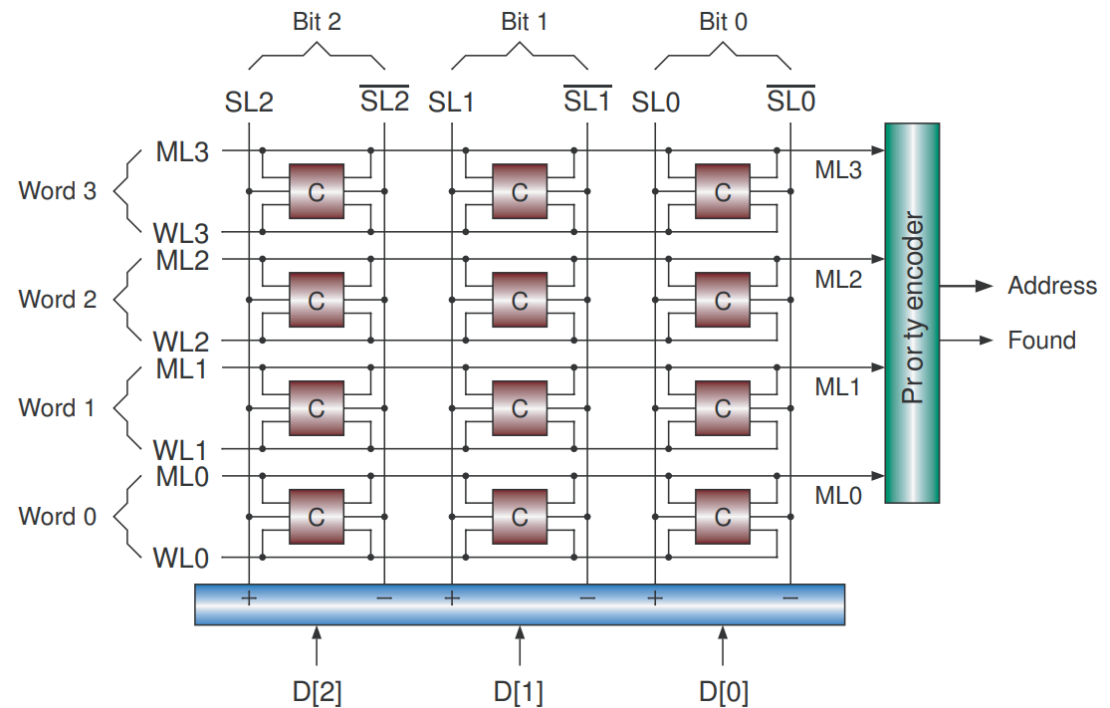

# Memory

---

- Memory is a temporary place for storing programs.
- SRAM
  - ✓ Fast and expensive (caches and TLBs).
- DRAM
  - ✓ Large and inexpensive (main memory).
- ROM
  - ✓ Read only memory (initial program or firmware).
- CAM
  - ✓ Content addressable memory.
  - ✓ Mainly used to design fully associative cache or TLB.

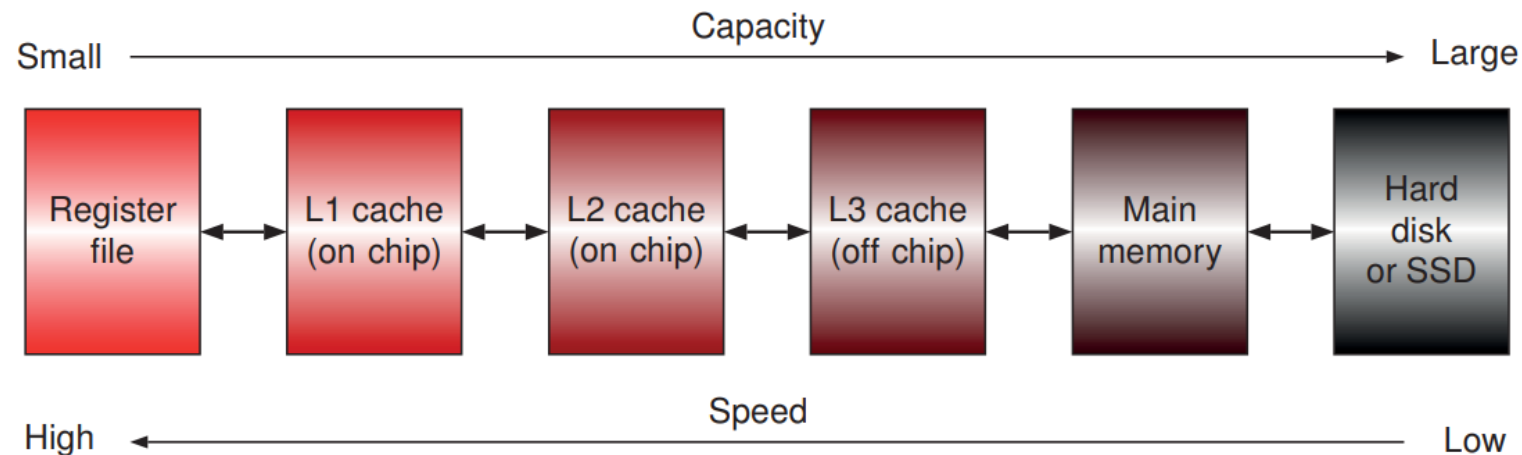
# Content Addressable Memory

- CAM is a very special memory.
- CAM searches the entire memory to see given data word is stored anywhere in it.
- If the data word is found, CAM returns the address.



# Memory Hierarchy

- The register file can be considered as the fastest memory.
- Caches store partial data of memory.
- L1 and L2 caches are on-chip, L3 is an off-chip.
- Main memory stores programs that are being executed.
- Hard disk or SSD has the largest capacity but is lowest speed.
- It store files with main memory or provide users virtual memory.



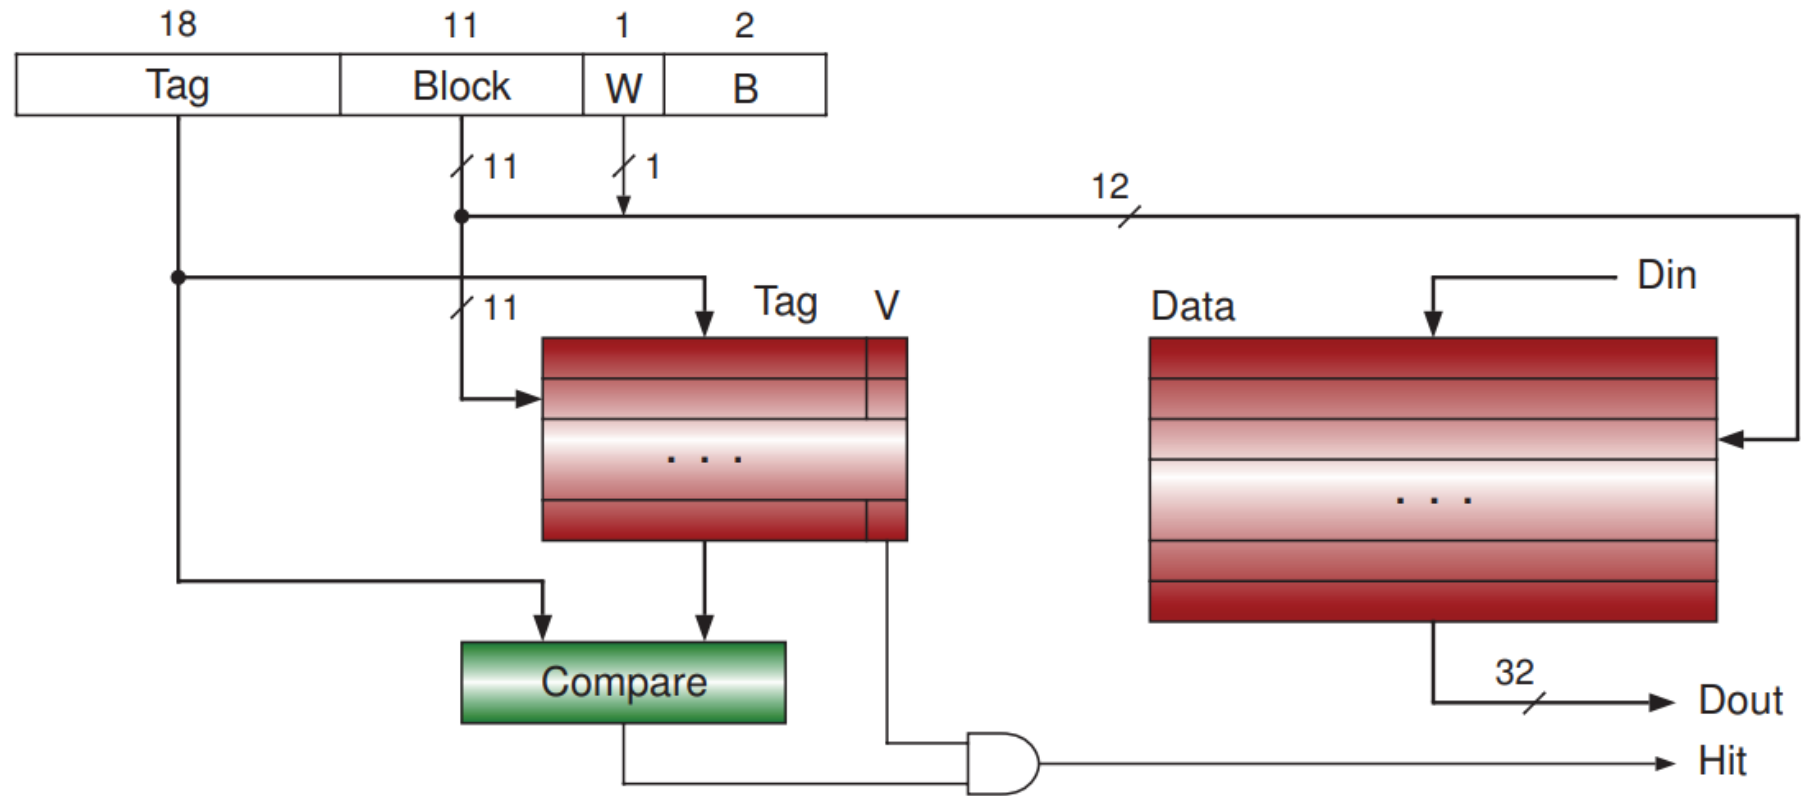
# Cache

---

- Caches in the computer field is a fast storage for storing data that are likely to be used again.
- The benefit of using a cache comes from the program's temporal locality and spatial locality.


# Cache

- Direct mapping cache



# Cache

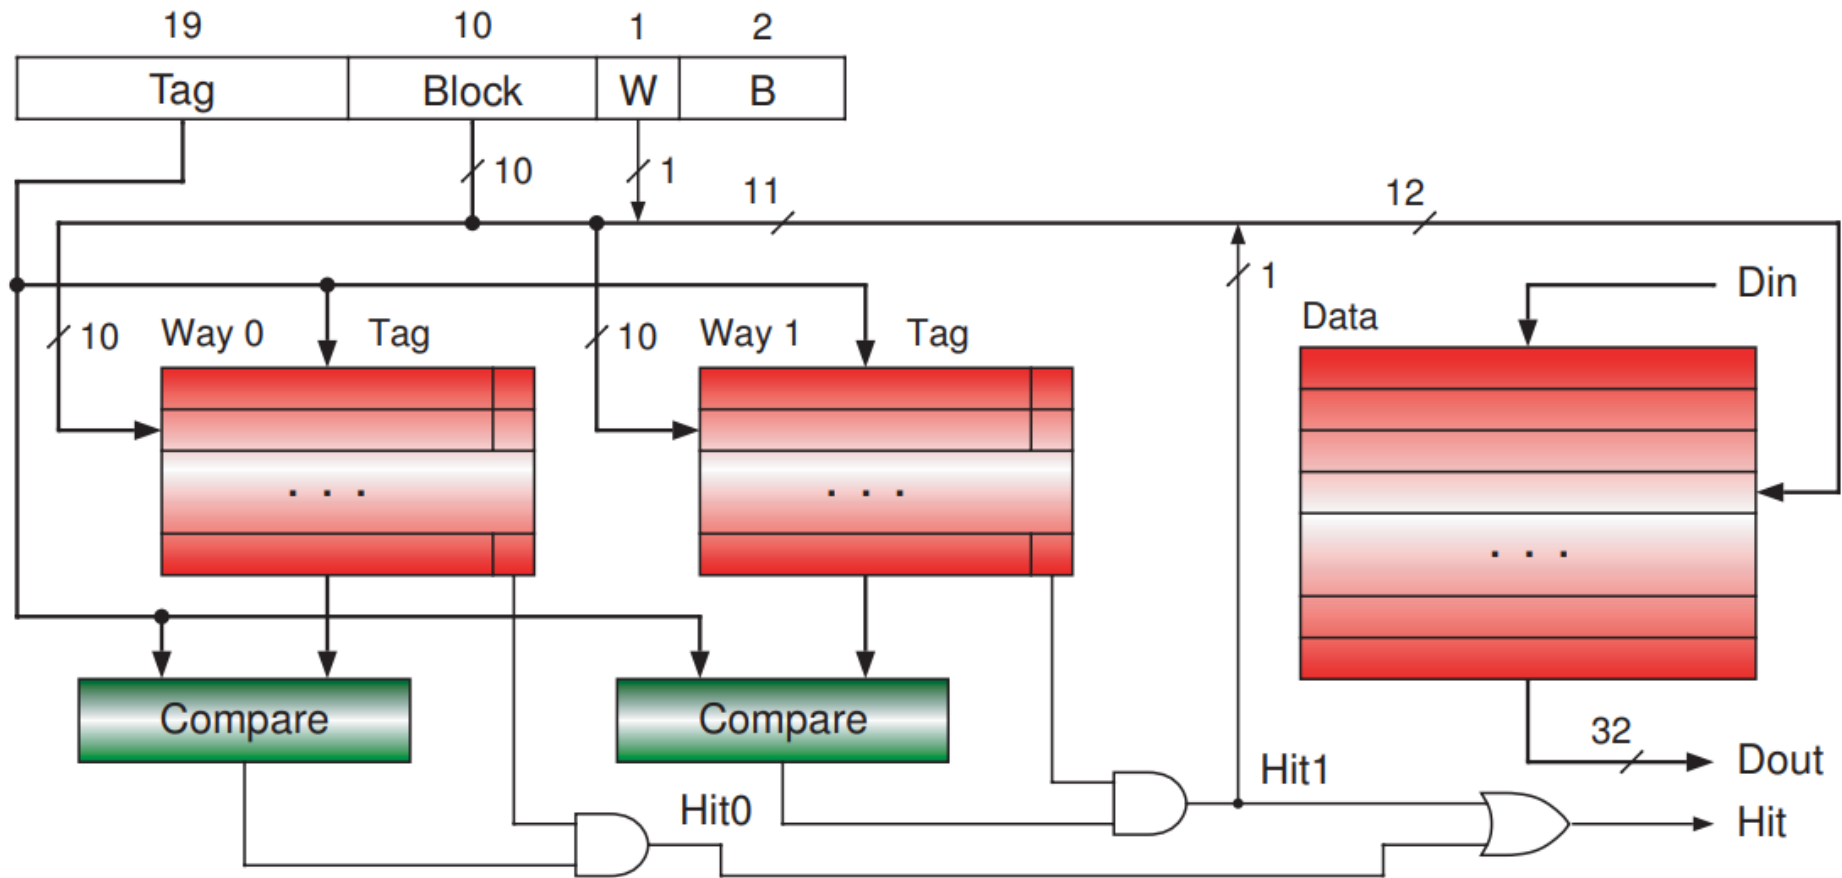
- Fully associative mapping





# Cache

- Set associative mapping



# Cache

---

- Cache block replacement algorithms
- LRU replacement algorithm
- Random replacement algorithm
- FIFO replacement algorithm

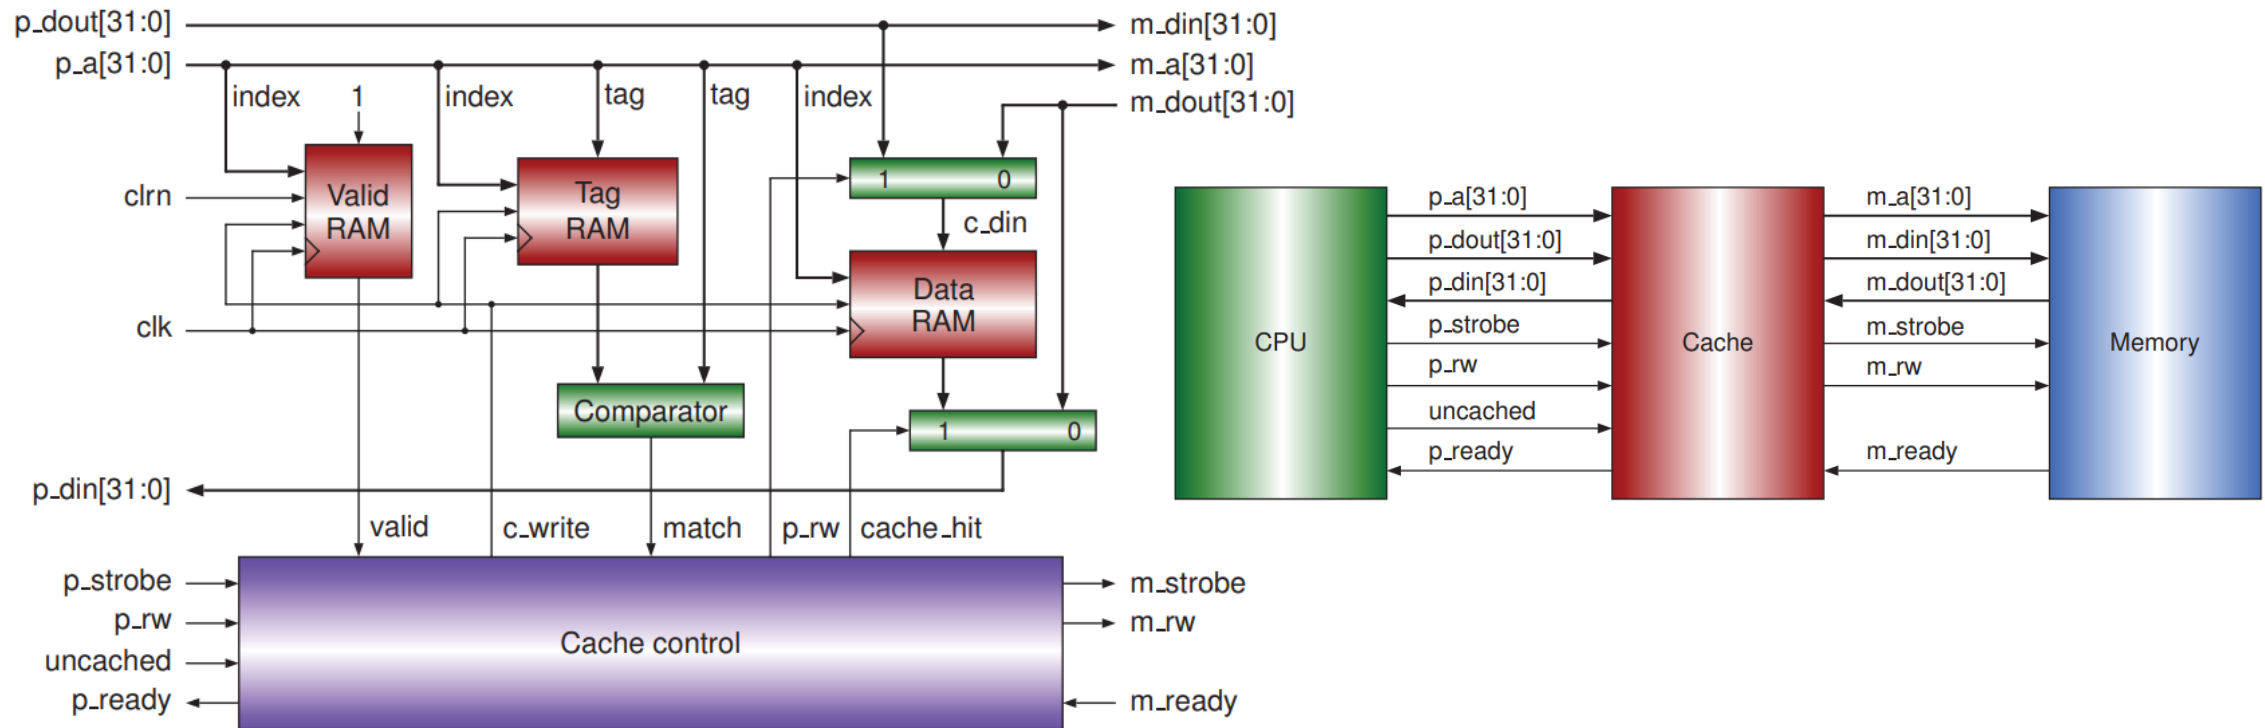
# Cache

---

- Cache write policies
- Cache hit
  - ✓ write through; update both the cache and the memory
  - ✓ write back; update the cache only, when it replaced update the memory
- Cache miss
  - ✓ write allocate
  - ✓ no write allocate

# Data Cache Design in Verilog HDL

- Direct mapping, write through, and write allocate



# Data Cache Design in Verilog HDL

---

- Virtual memory is not a real memory.
- It allows processes to use more memory than the real memory.
- When process access a virtual memory, MMU translates the virtual address into a physical address.
  - ✓ To speed up this translation, TLB is fabricated in CPU.

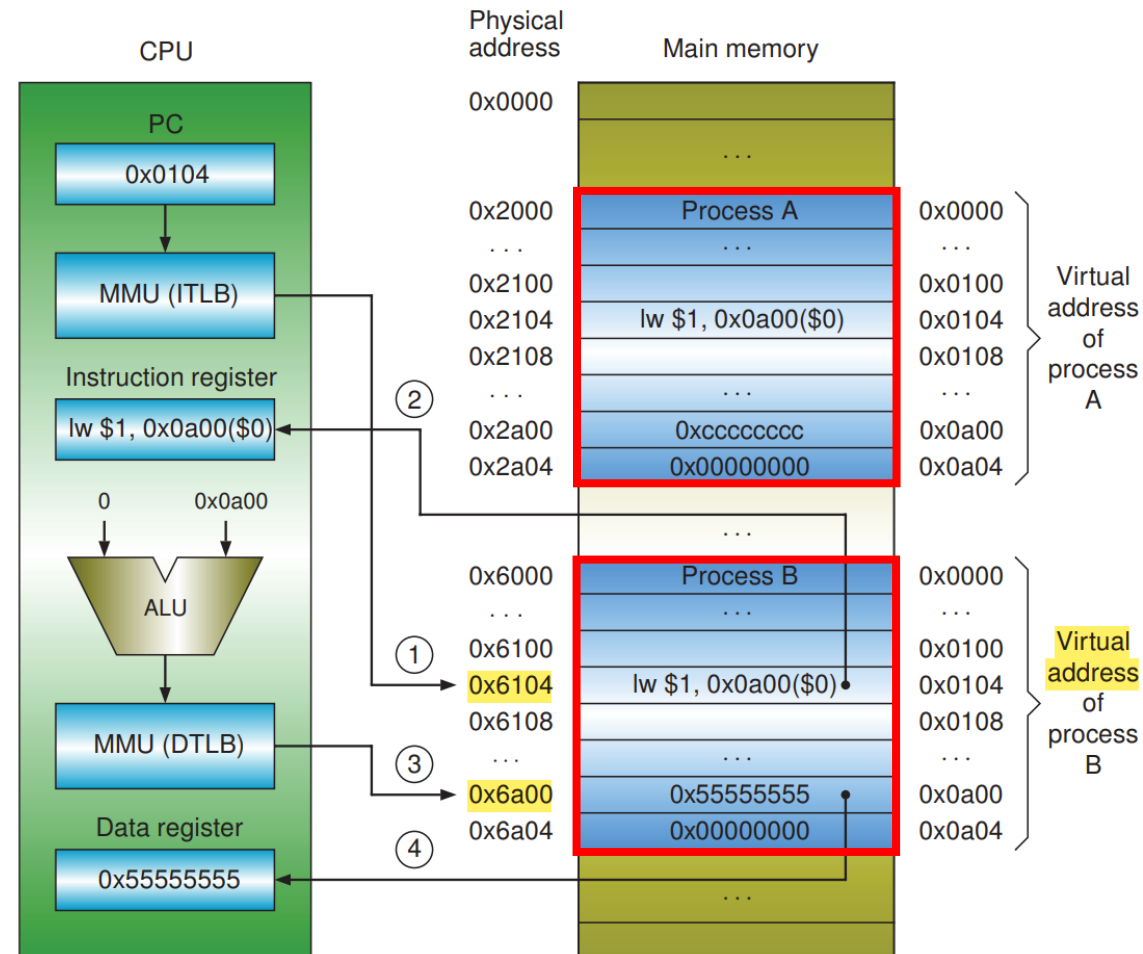
# Virtual Memory and TLB

- The main memory is a real storage in which programs stored.
- Main memory is also called physical memory or real memory.
- Multiple processes can be resided in the main memory simultaneously.
- The process use a virtual address to access its virtual memory.

Process A		Process B	
Virtual address	Instruction or data	Virtual address	Instruction or data
0x0000:	...	0x0000:	...
...	...	...	...
0x0104:	lw \$1, 0x0a00(\$0)	0x0104:	lw \$1, 0x0a00(\$0)
...	...	...	...
0x0a00:	0xcccccccc	0x0a00:	0x55555555
0x0a04:	0x00000000	0x0a04:	0x00000000

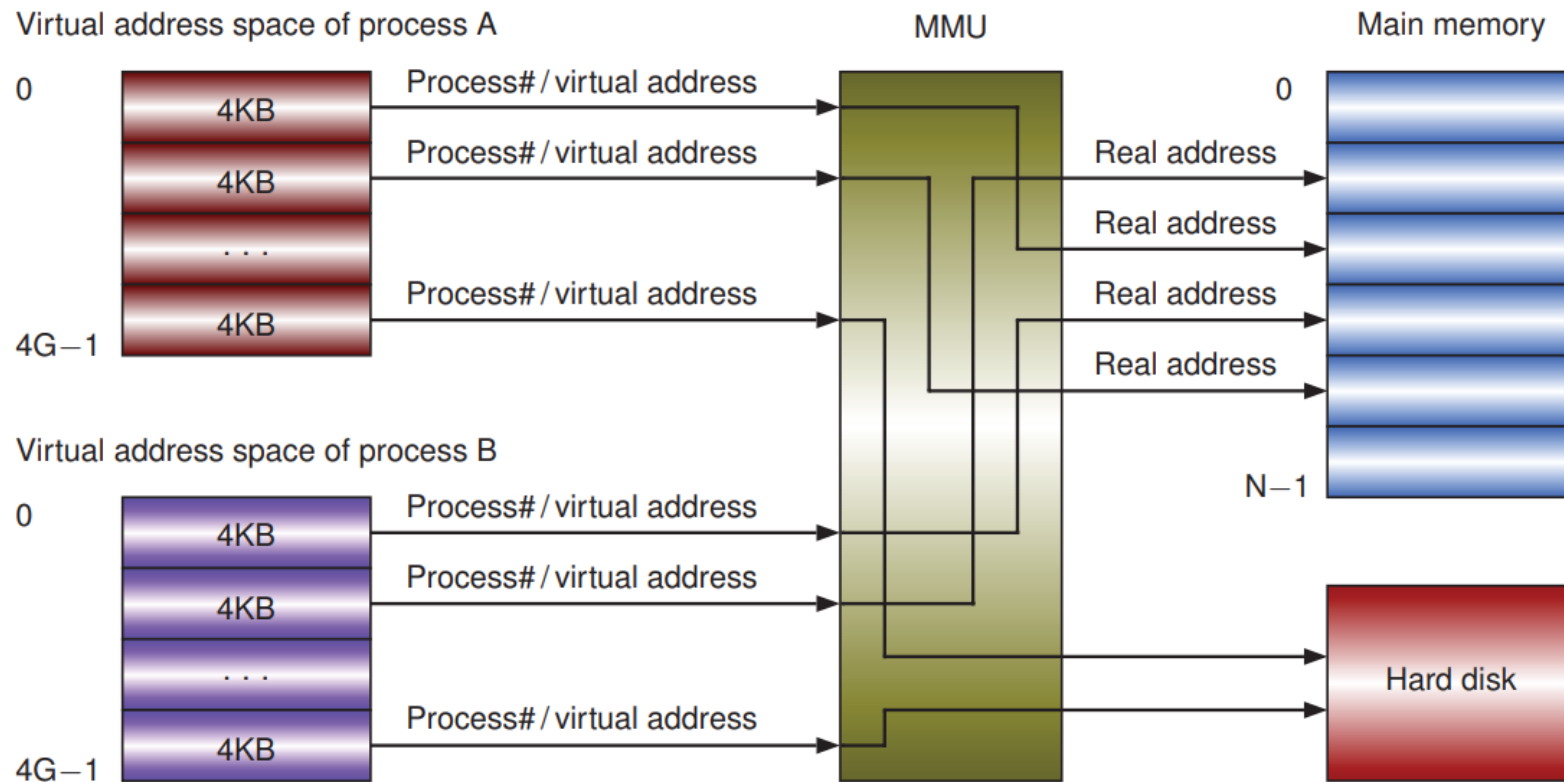
# Virtual Memory and TLB

- Two processes reside in main memory



# Virtual Memory and TLB

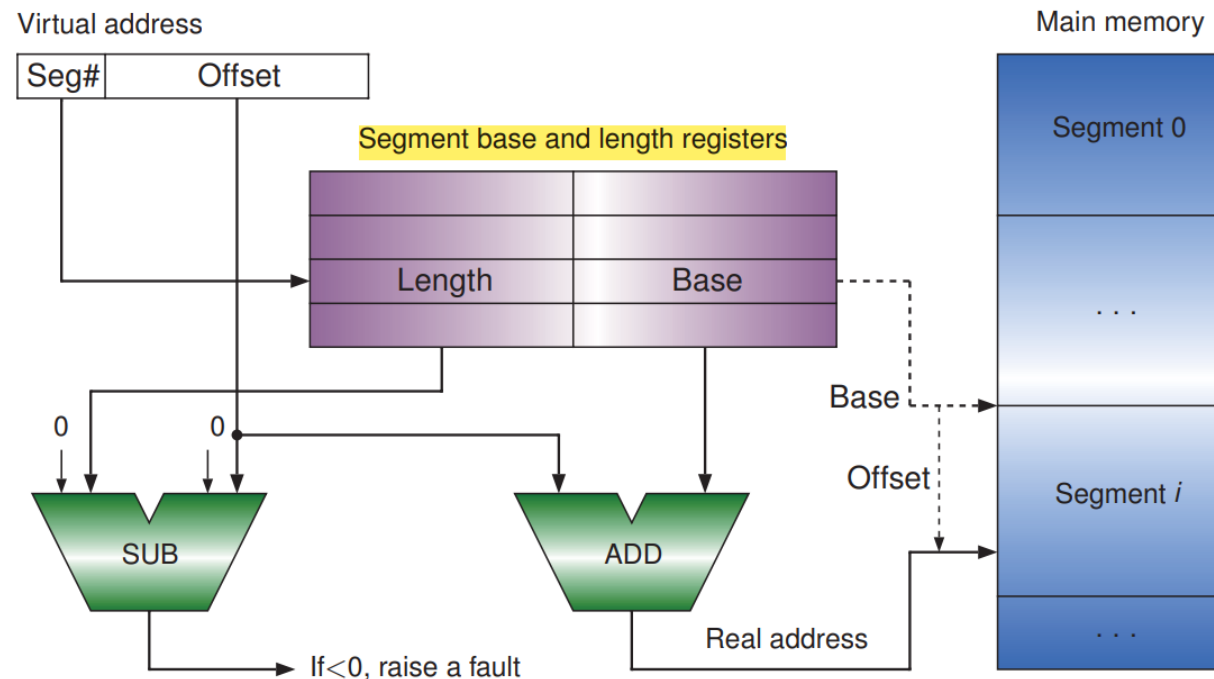
- Concept of virtual memory management





# Virtual Memory and TLB

- **Segmentation management**
- With segmentation management, the main memory is divided into segments based on the natural divisions of a program.
- The size of a memory segment is generally not fixed.



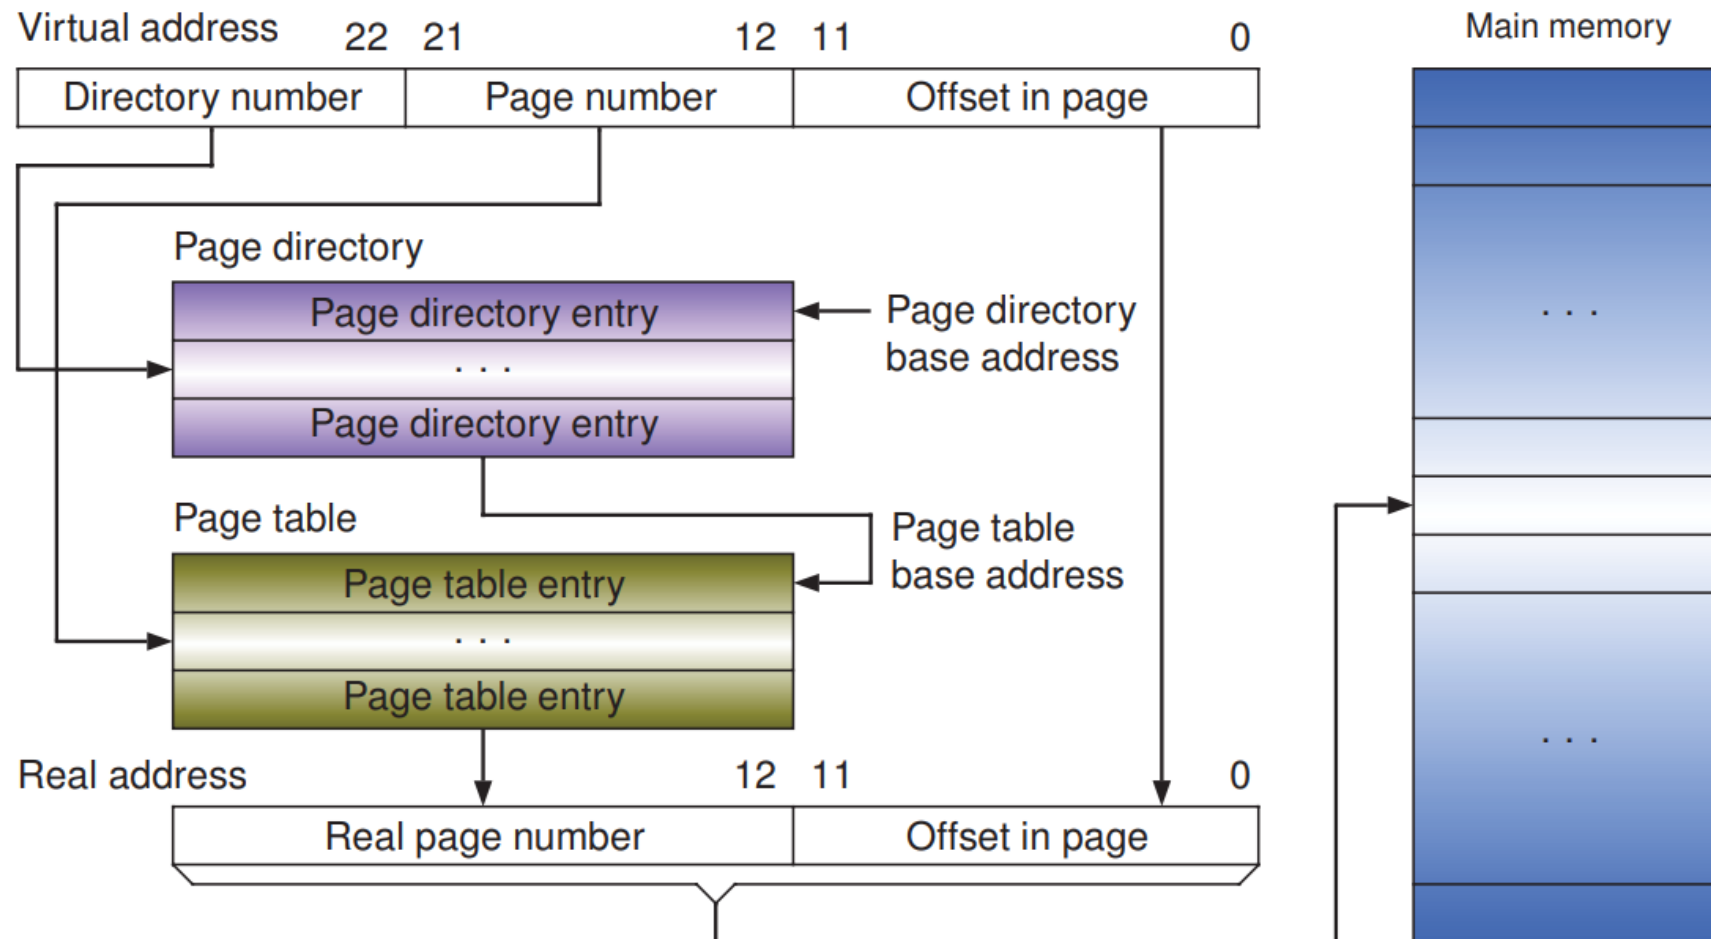
# Virtual Memory and TLB

---

- **Paging management**
- With paging management, both the virtual memory and real memory are divided into pages of fixed size.
- Page is the smallest unit with which the MMU maps a virtual page to a physical page.
  - ✓ Translate VPN to RPN
- Translation is done with a page table.
- To prevent page table takes a large memory, we use two-level page tables.
- The first-level page table is called page directory; it stores the physical base address of the (second-level) page table.

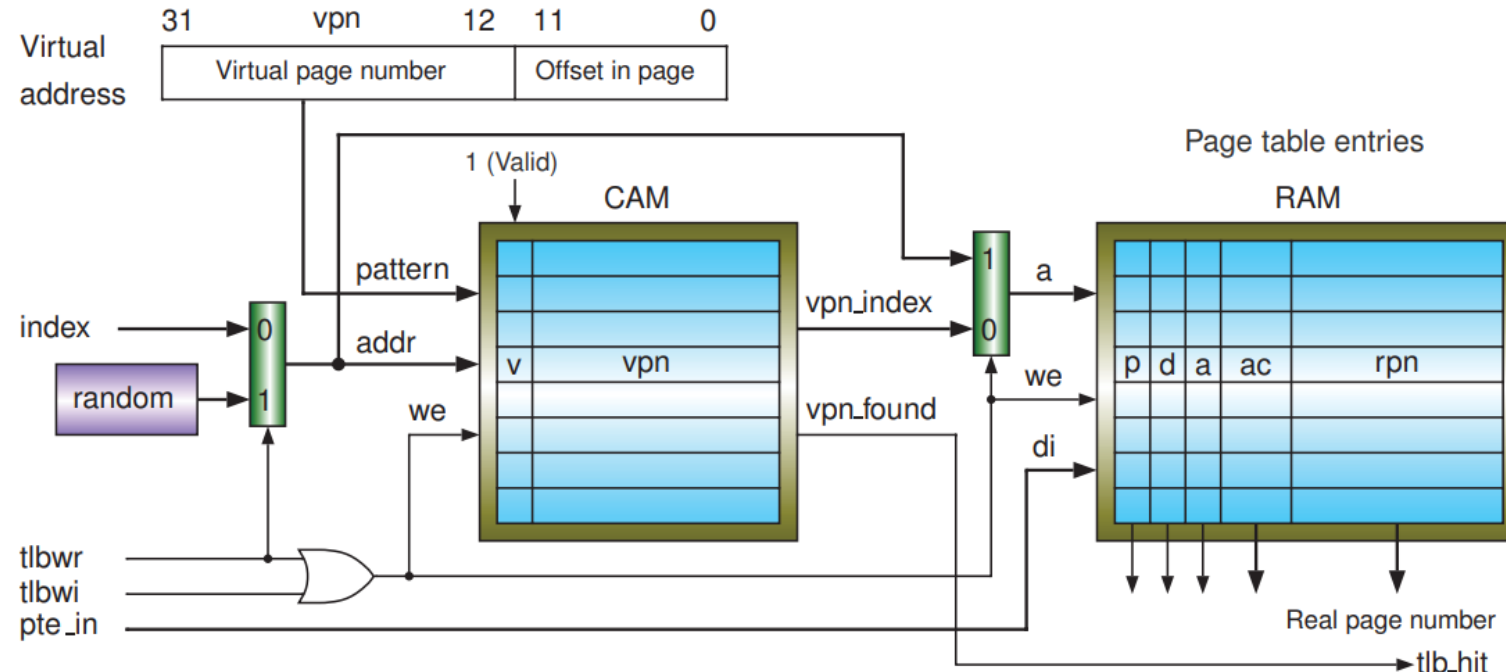
# Virtual Memory and TLB

- Two-level paging management

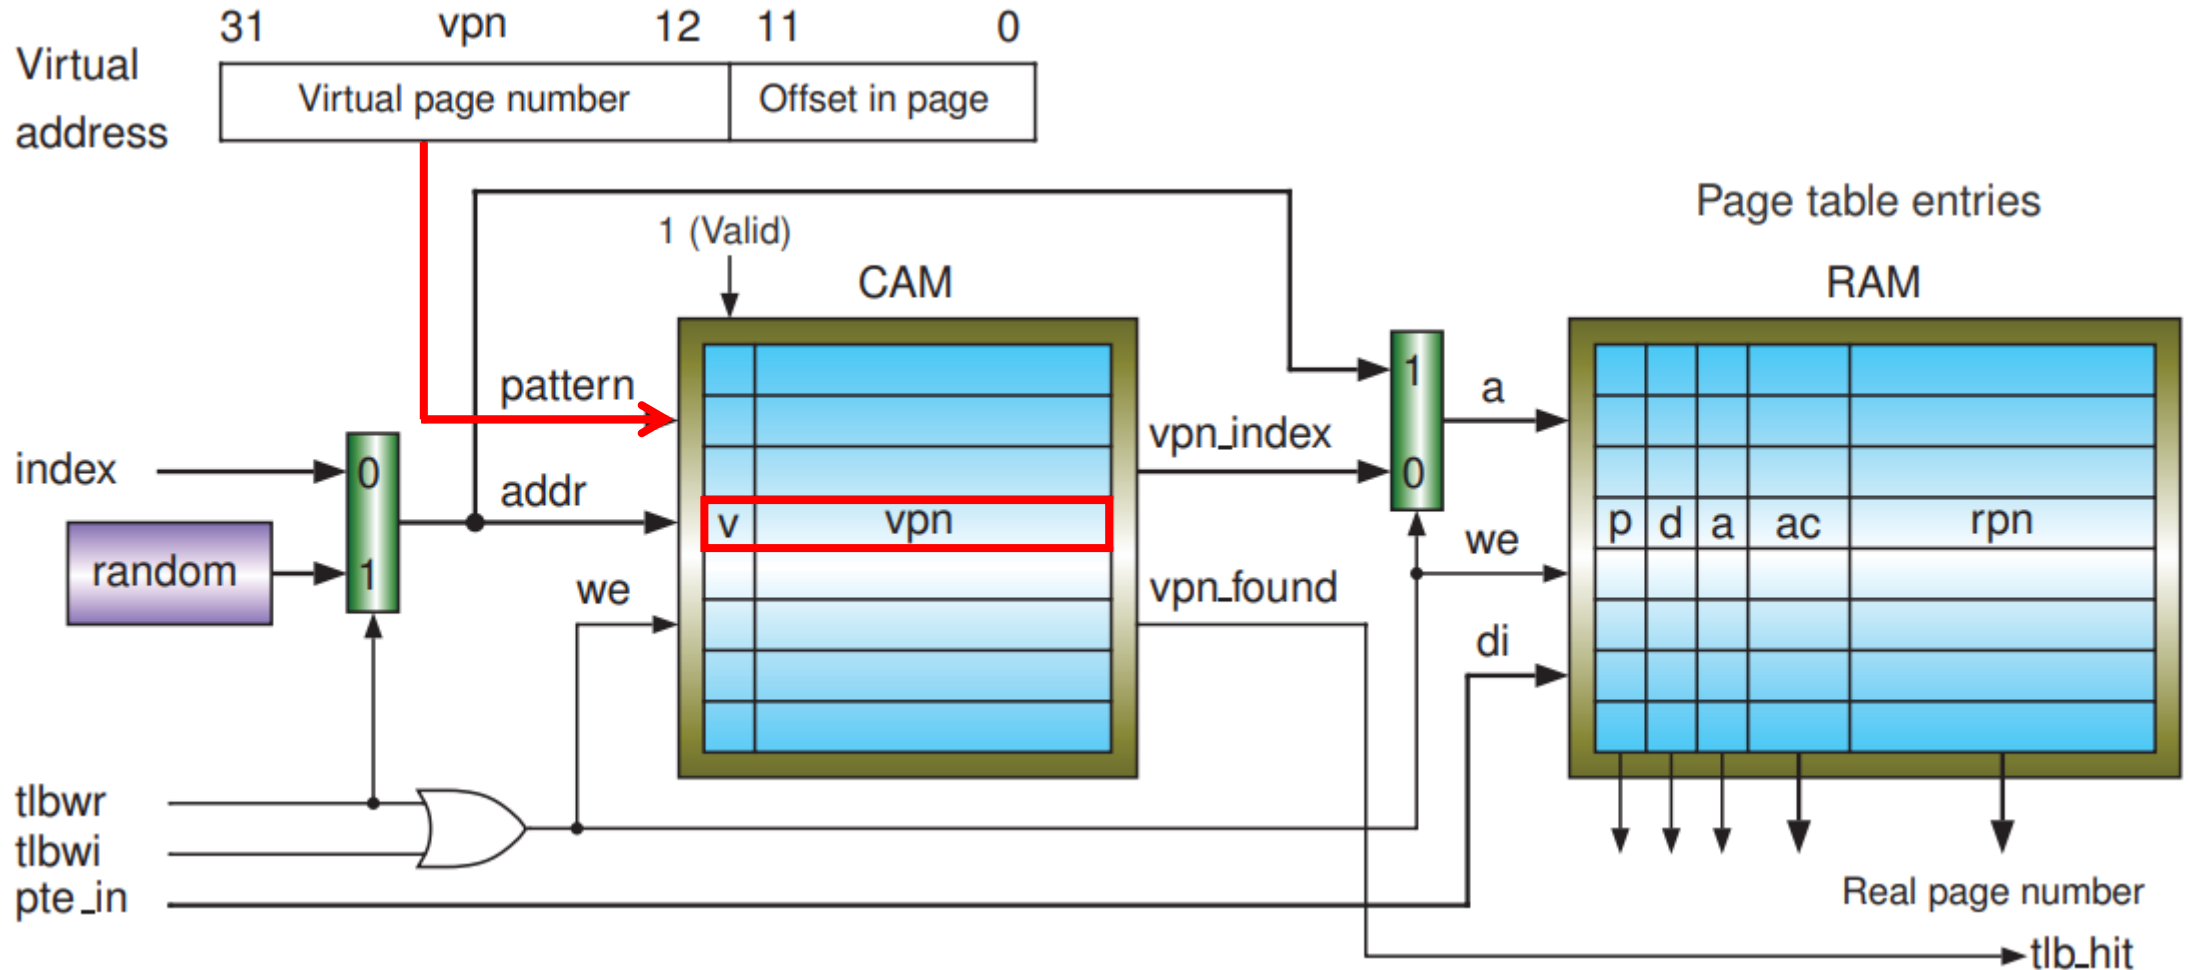


# Virtual Memory and TLB

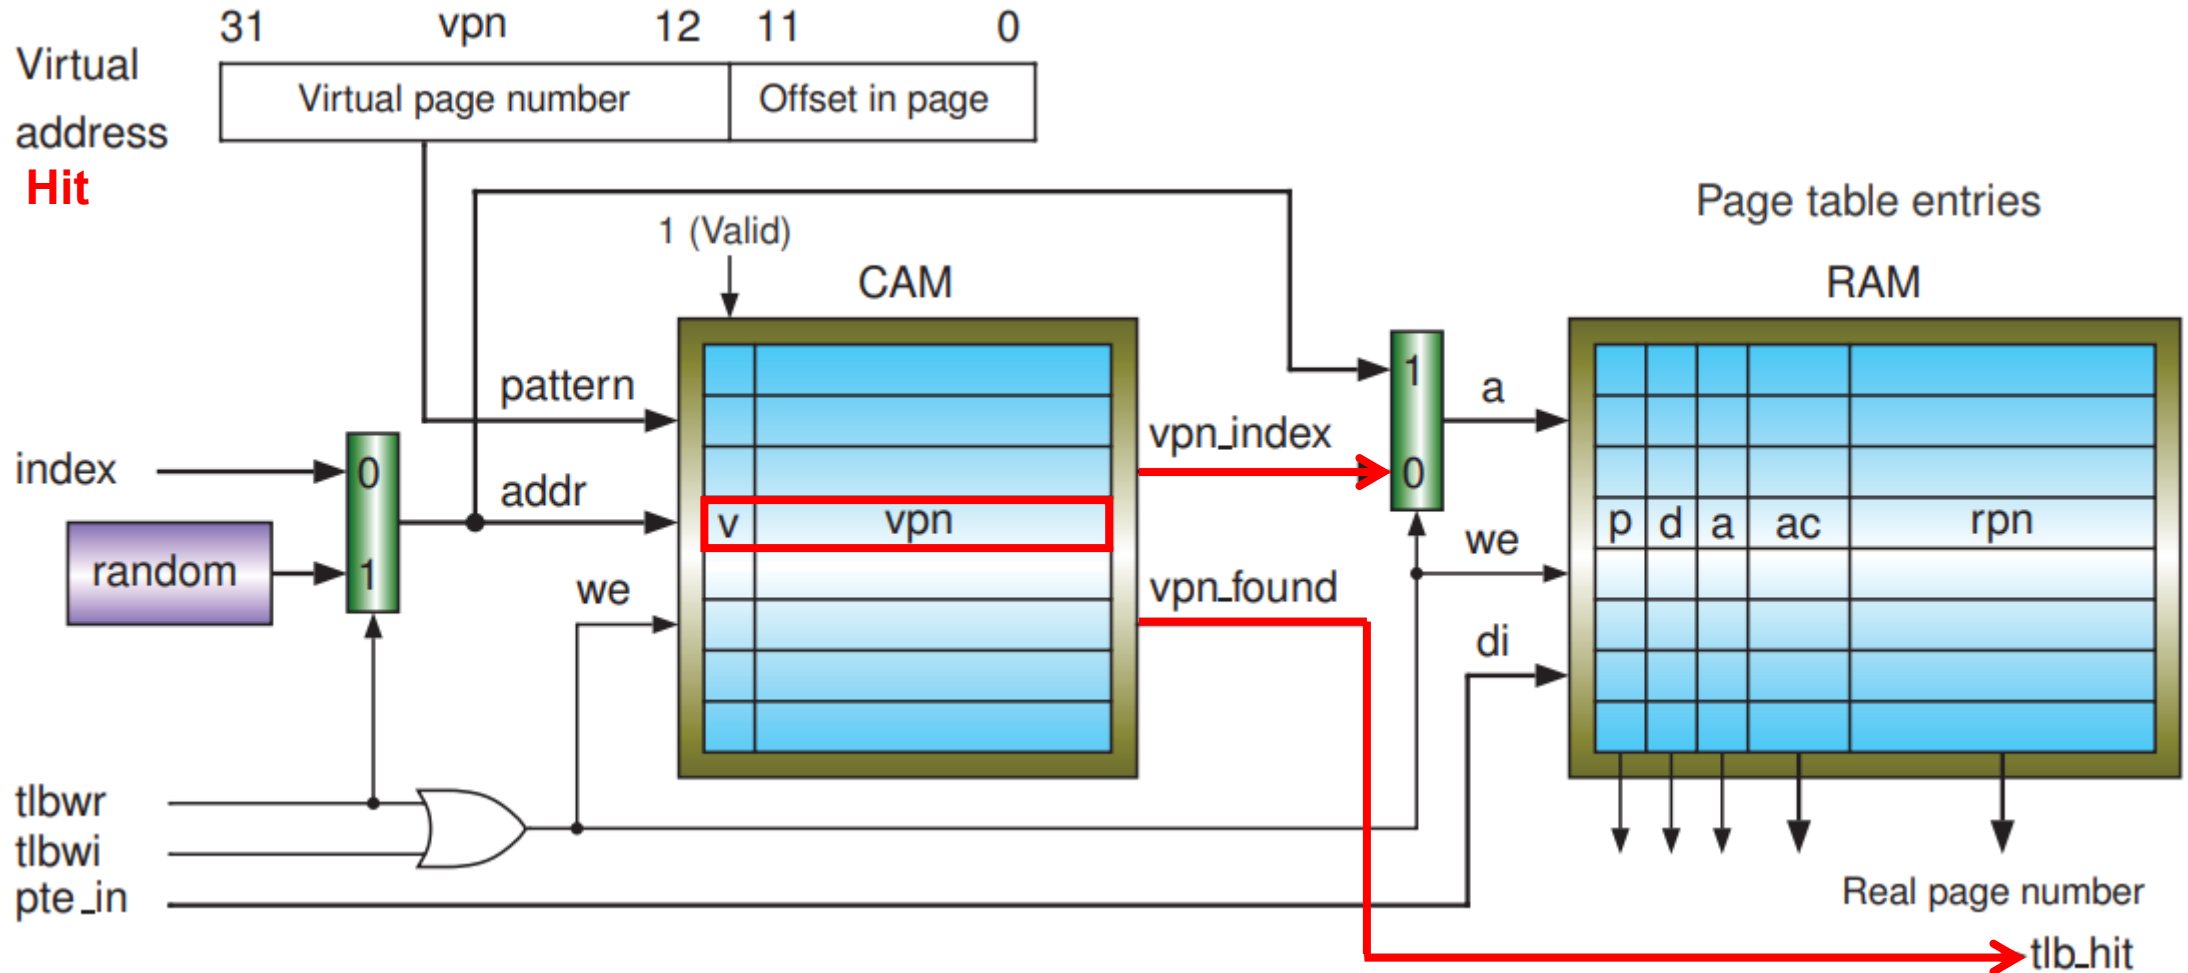
- To speed up the translation, CPUs fabricate **TLBs**.
- The data cache stores data blocks in cache RAM, while the TLB stores RPNs.
- CAM is used to search for the VPN of the virtual address.



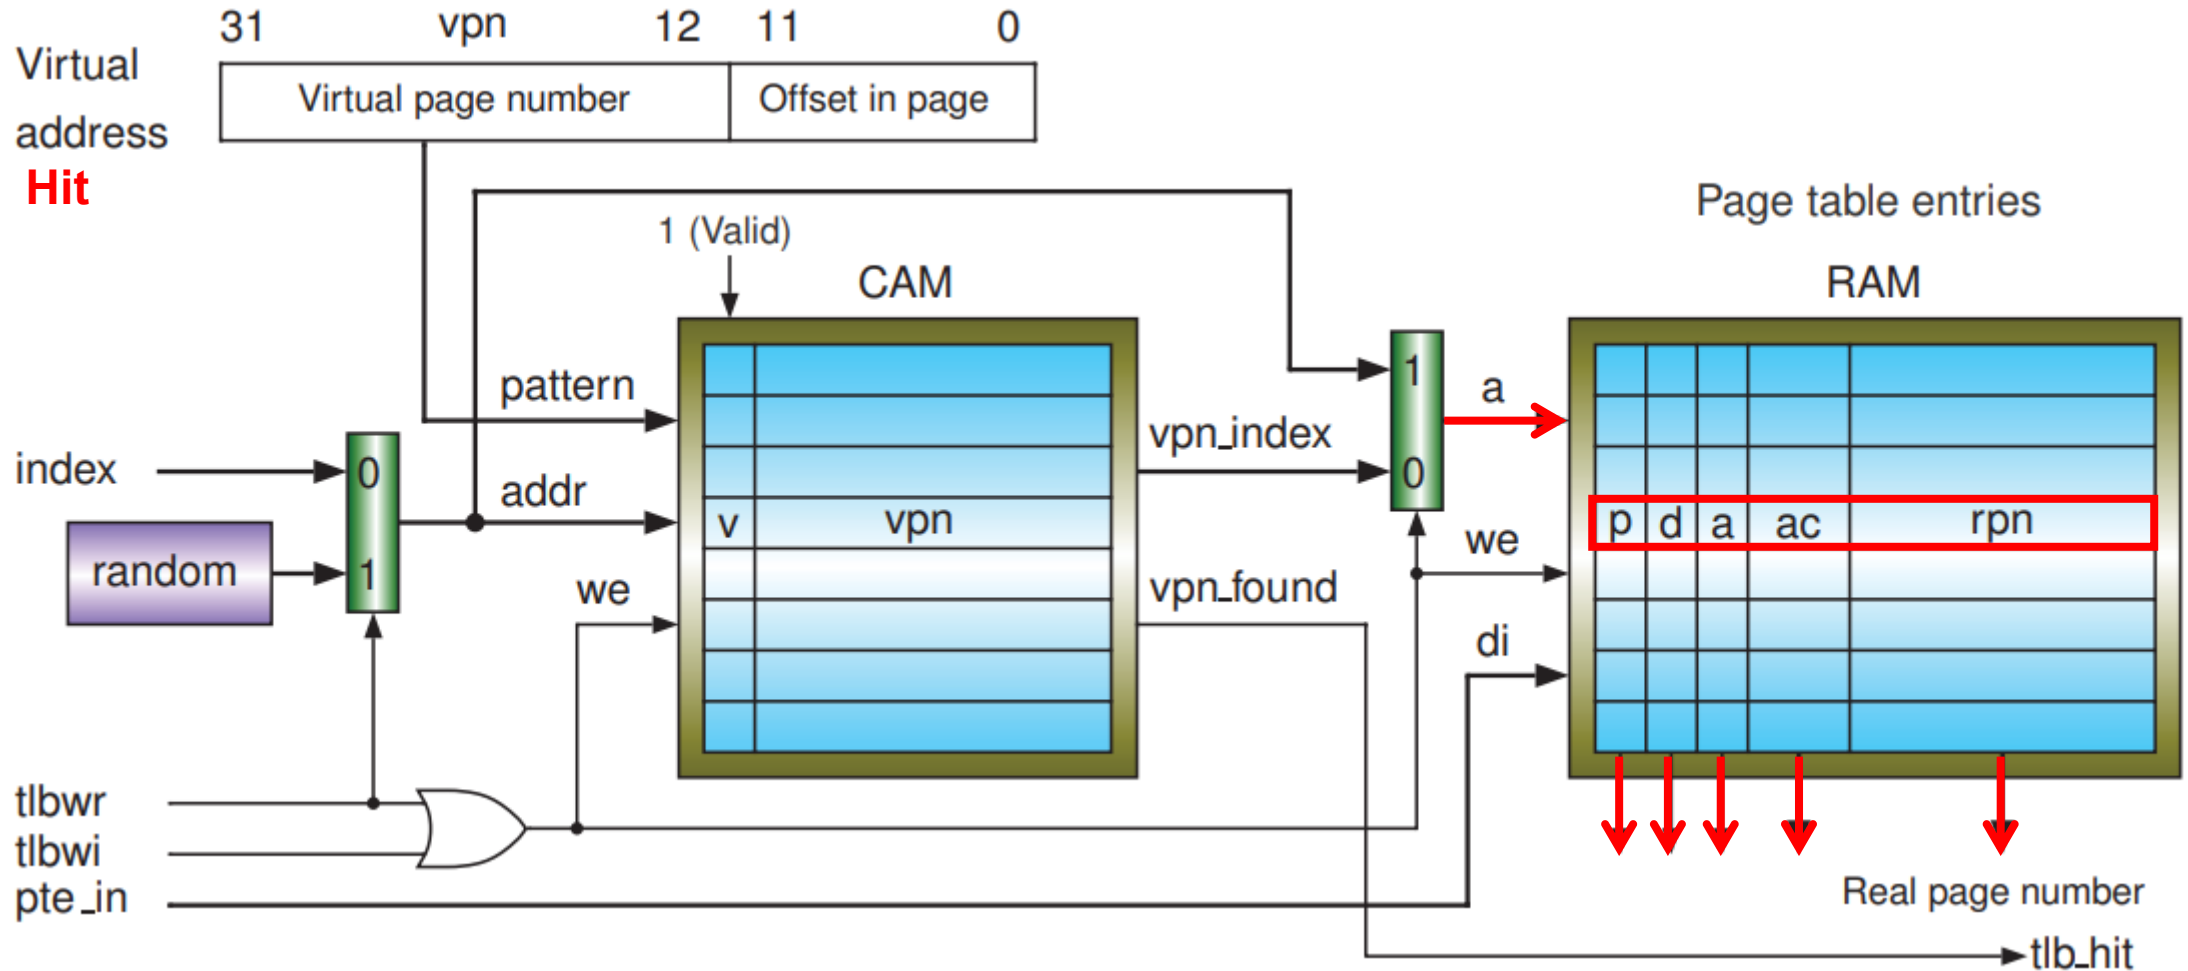
# Virtual Memory and TLB



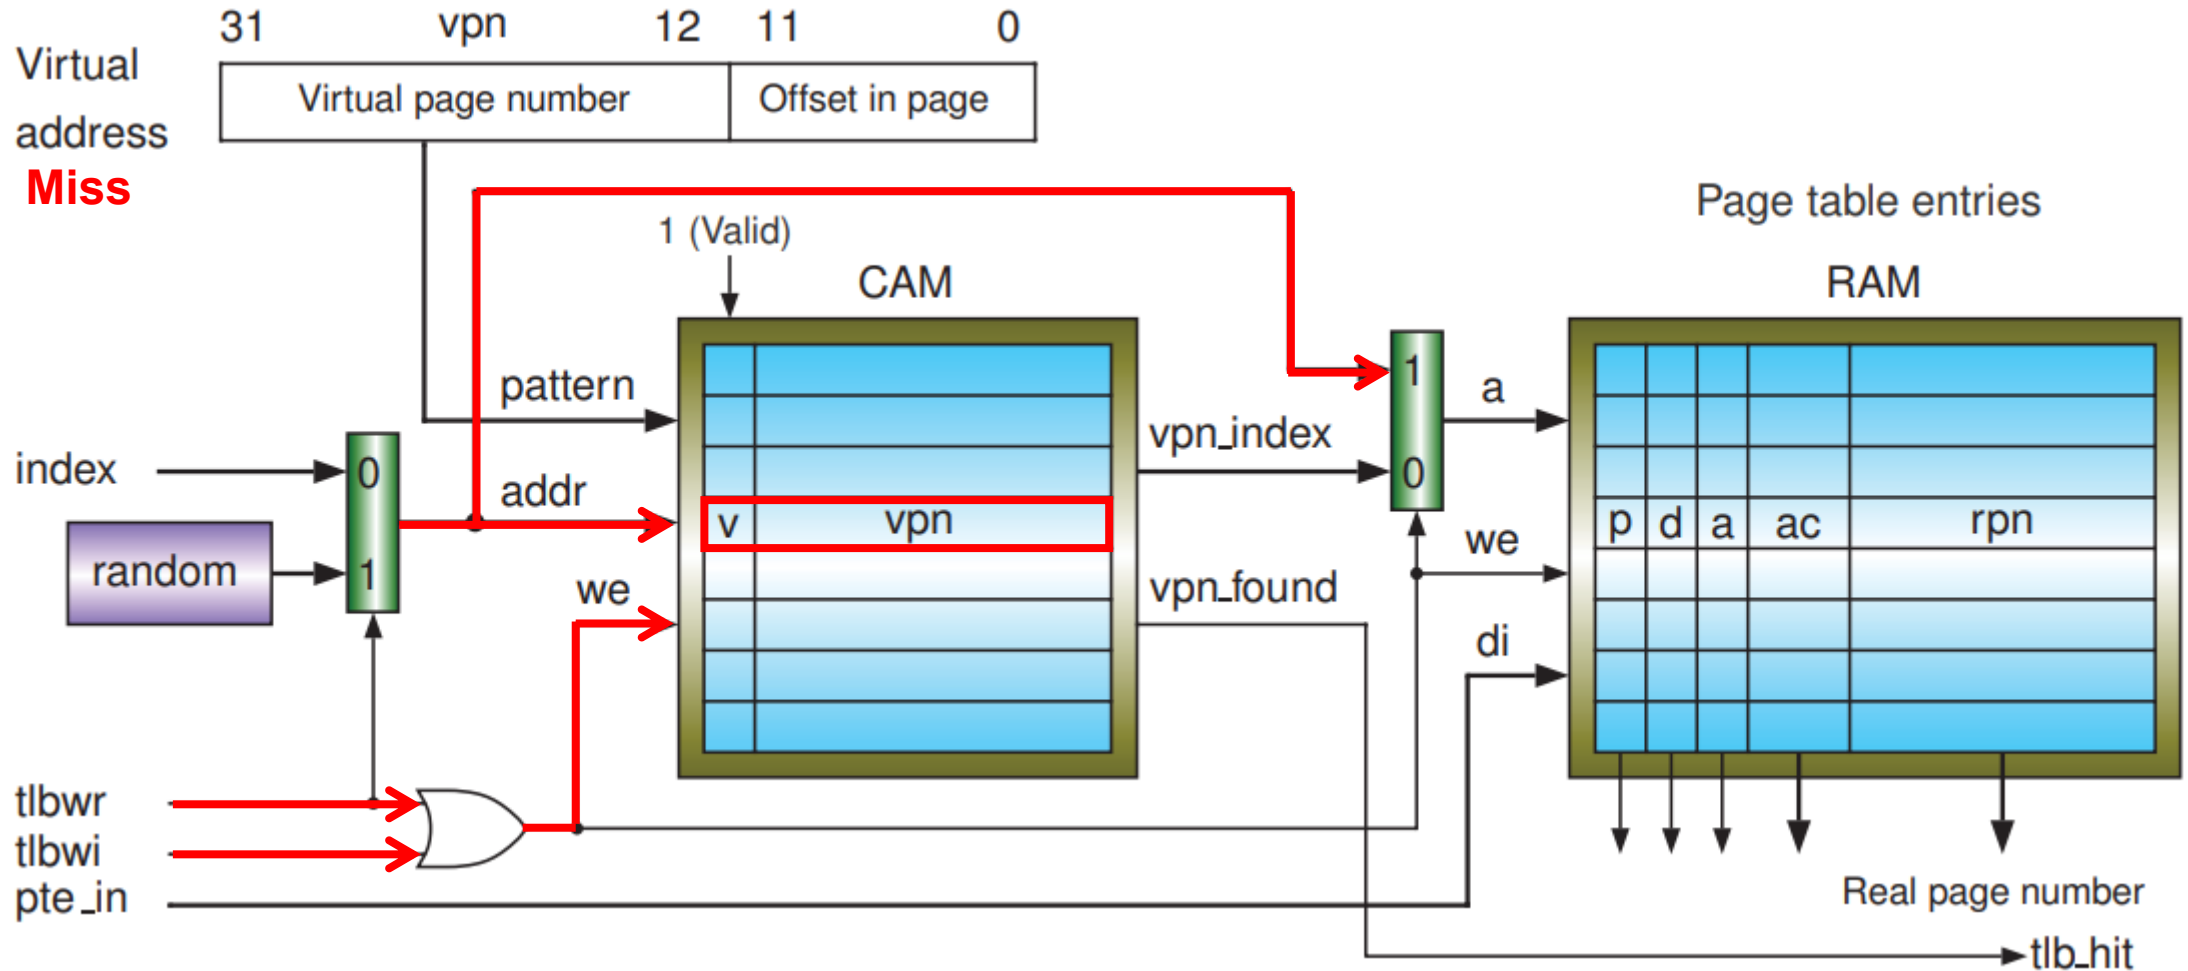
# Virtual Memory and TLB



# Virtual Memory and TLB

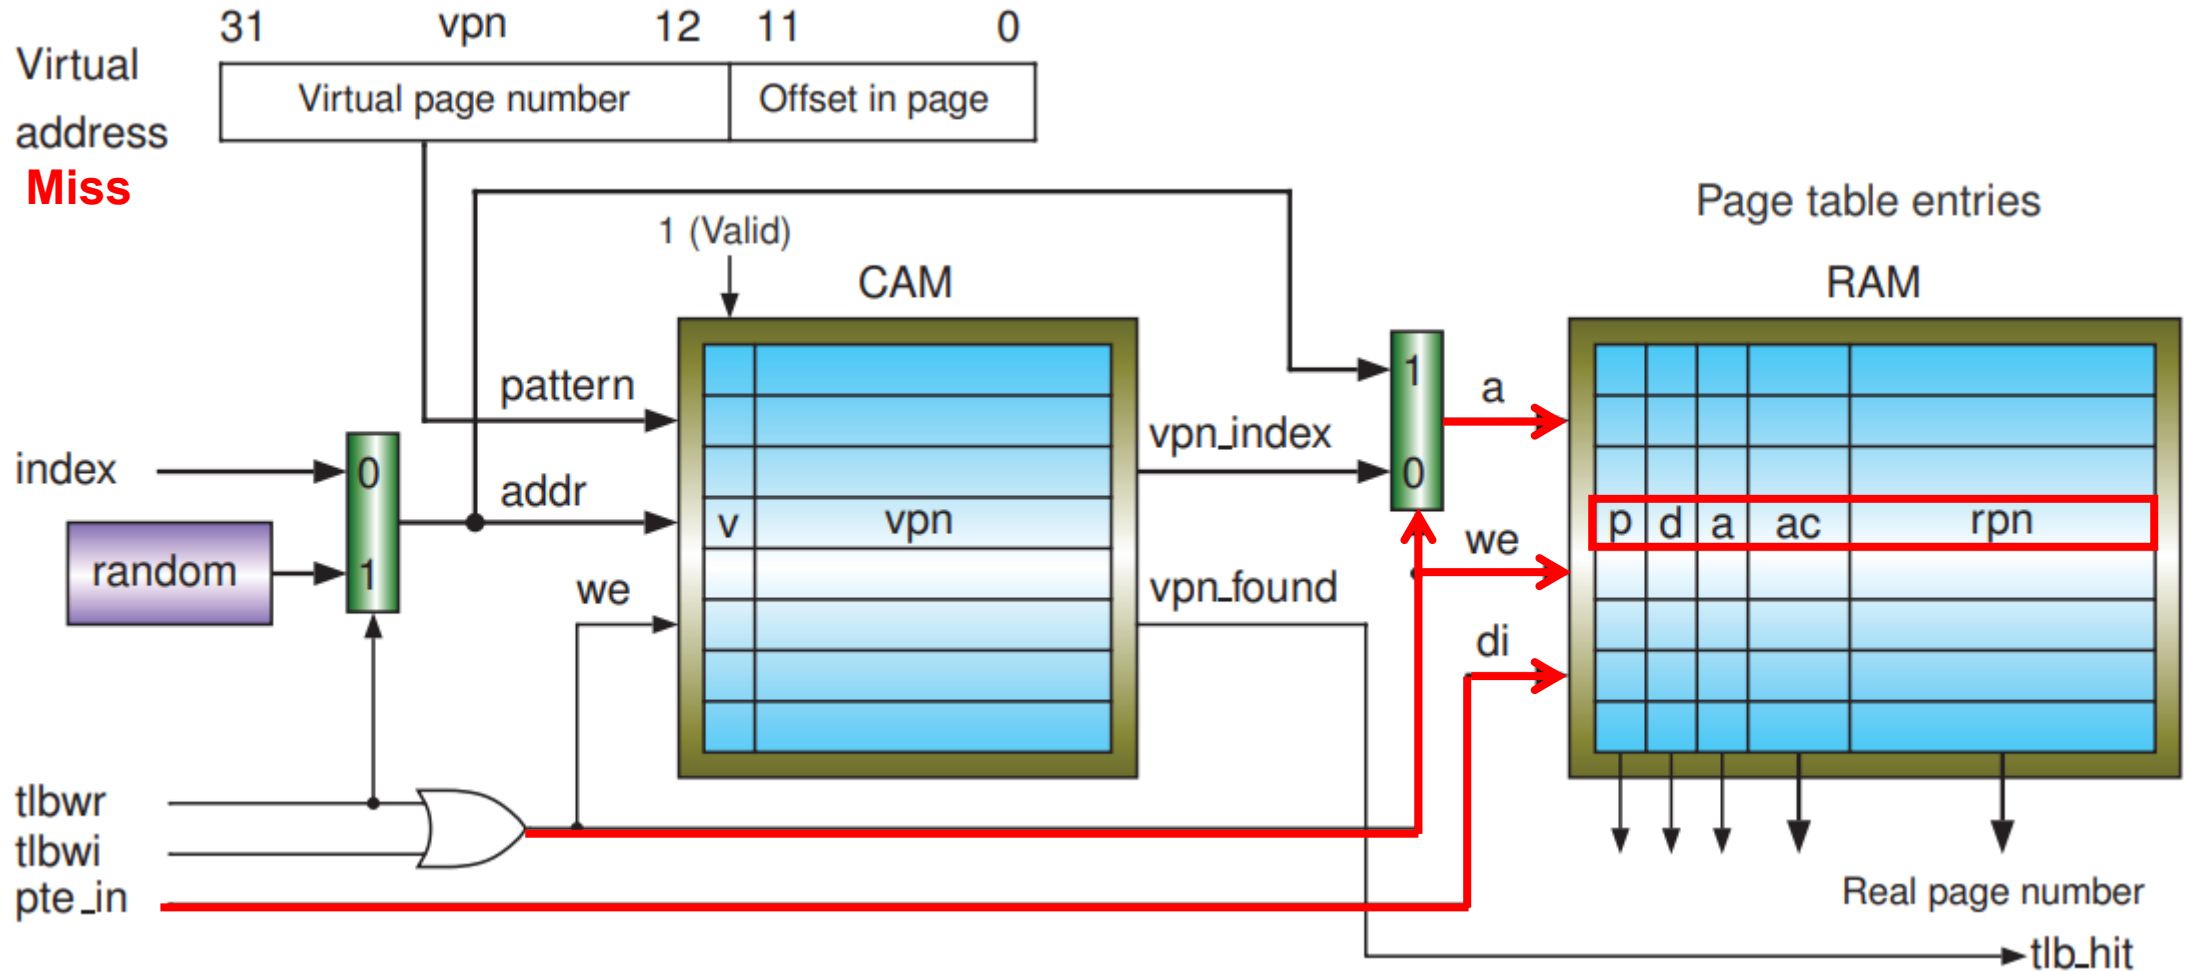


# Virtual Memory and TLB





# Virtual Memory and TLB



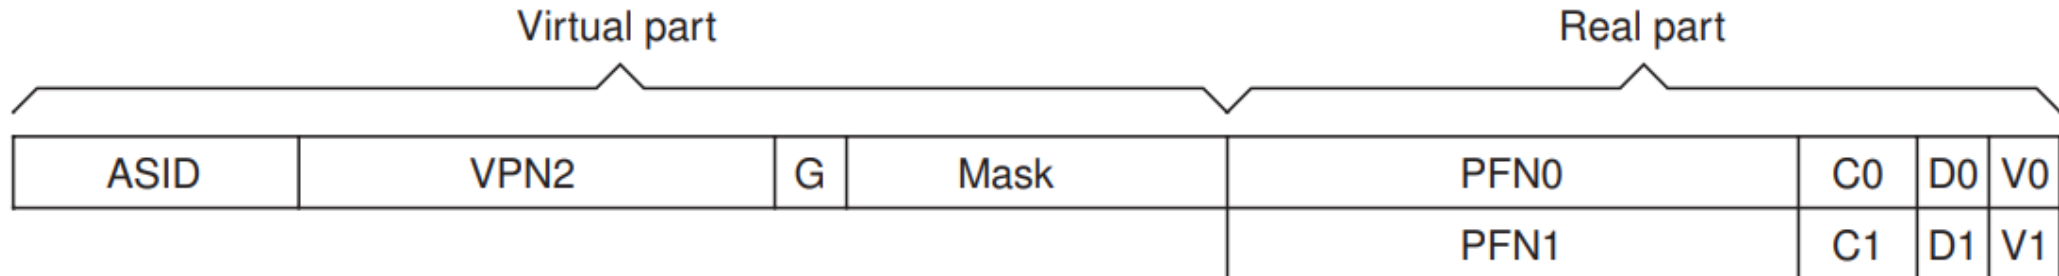
# TLB-Based MIPS Memory Management

---

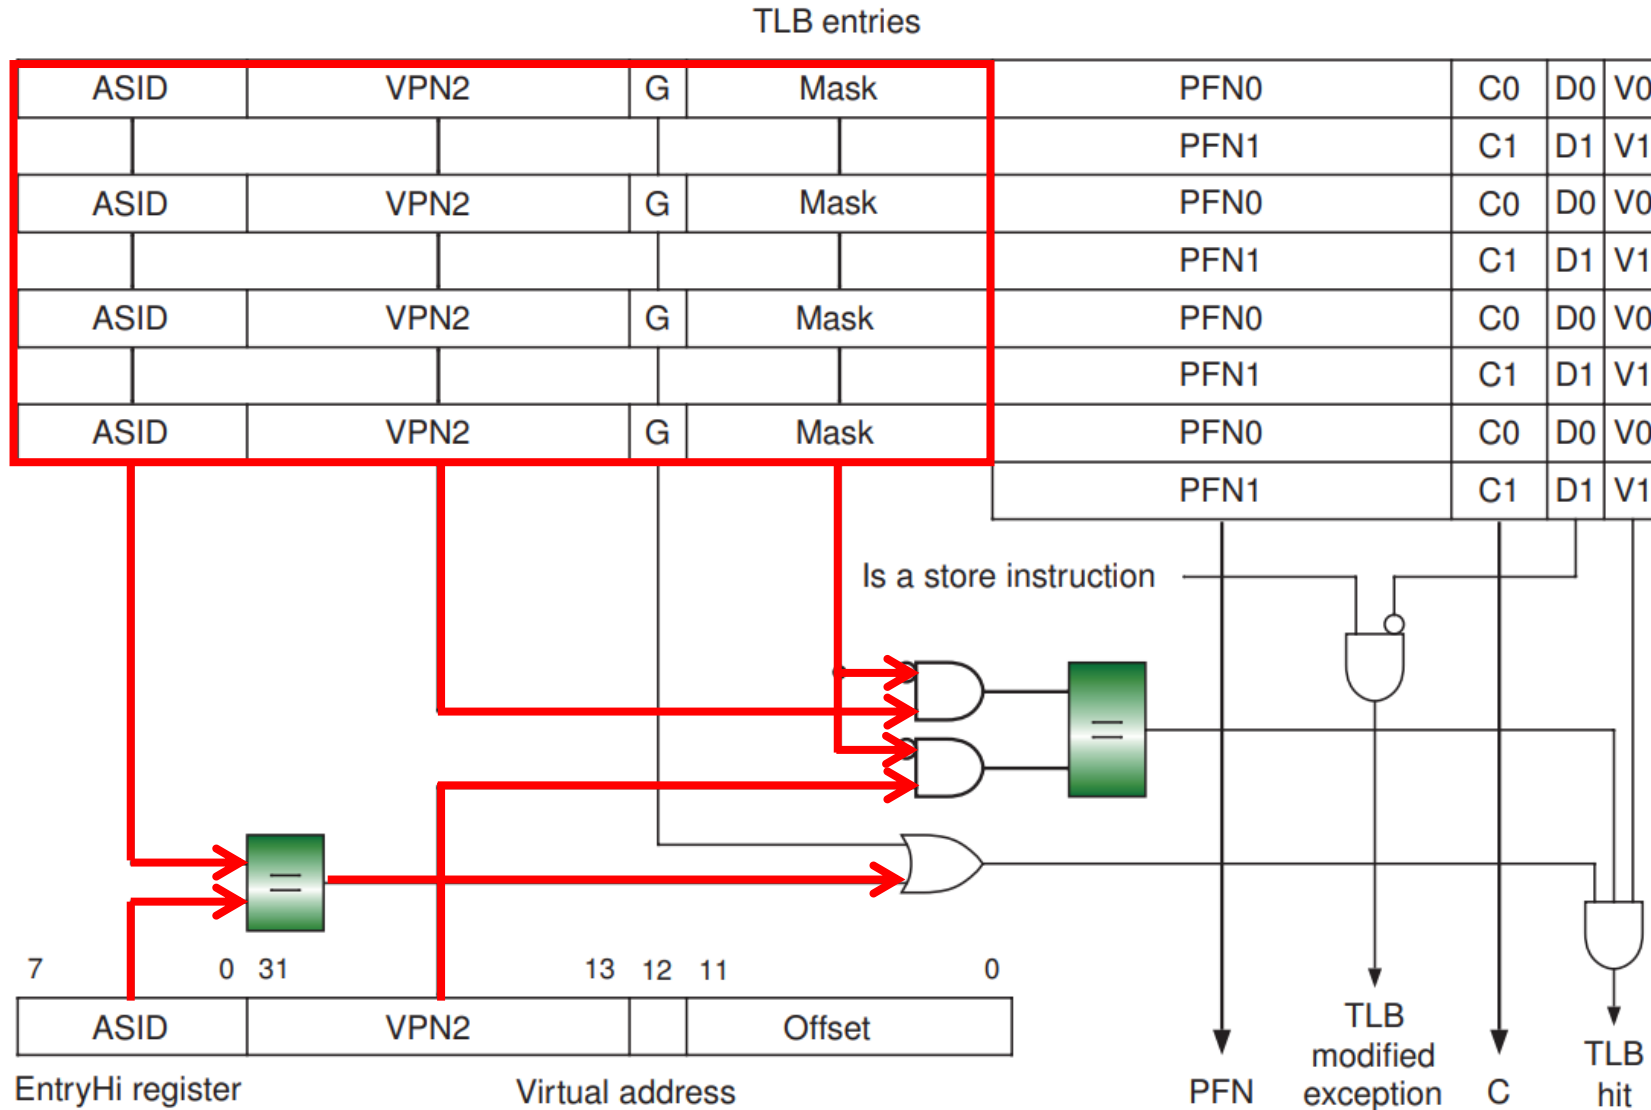
- The MIPS CPU uses a TLB-based address translation.
- MIPS CPU runs at one of two privilege modes
  - User mode
    - ✓The program has access only to the general-purpose register file, floating-point register file, and up to 2GB virtual memory
  - Kernal mode
    - ✓The program has access to 4GB virtual memory, and all register files include **CP0 registers**.
    - ✓Privileged instructions can be executed.

# TLB-Based MIPS Memory Management

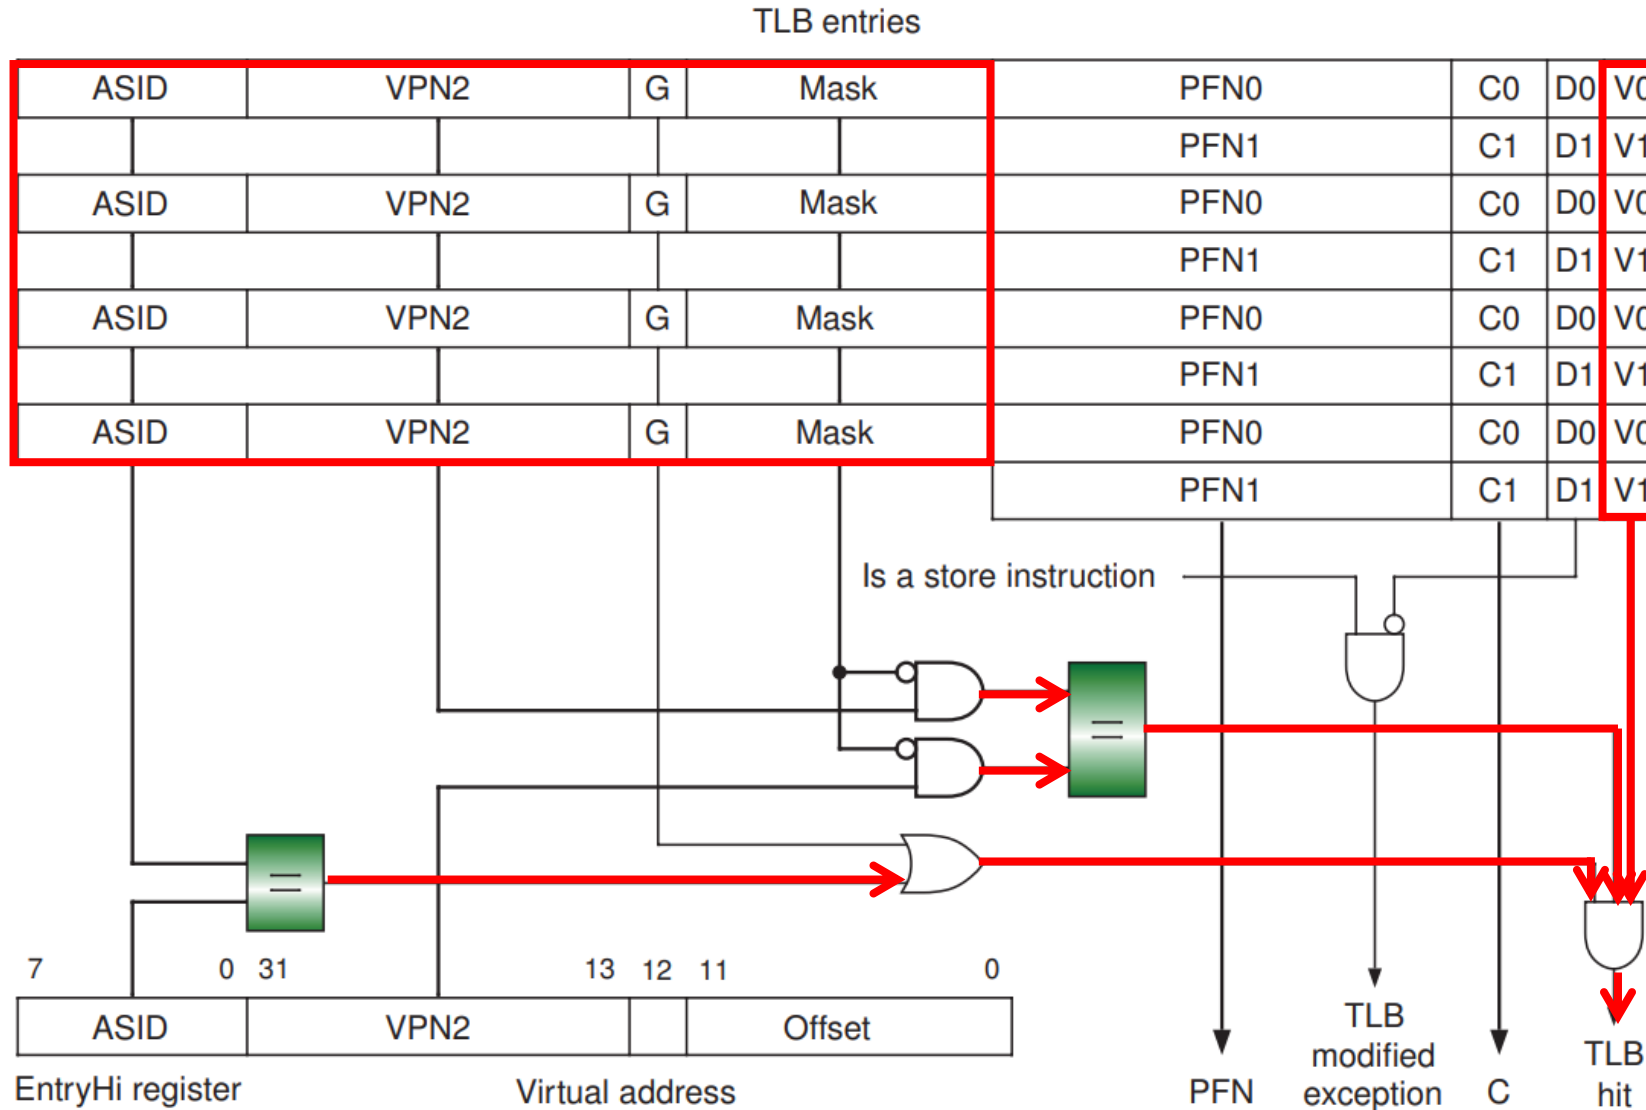
- **Organization of MIPS TLB**
- In MIPS architecture, virtual addresses are extended with an 8-bit **address space identifier (ASID)** to distinguish between processes.
- The **ASID** is stored in every TLB entry.
- Each entry contains a virtual part and a real part.
  - ✓ Virtual part (ASID, VPN2, G, and Mask)
  - ✓ Real part (entry 0, entry 1, PFN, C, D, and V)



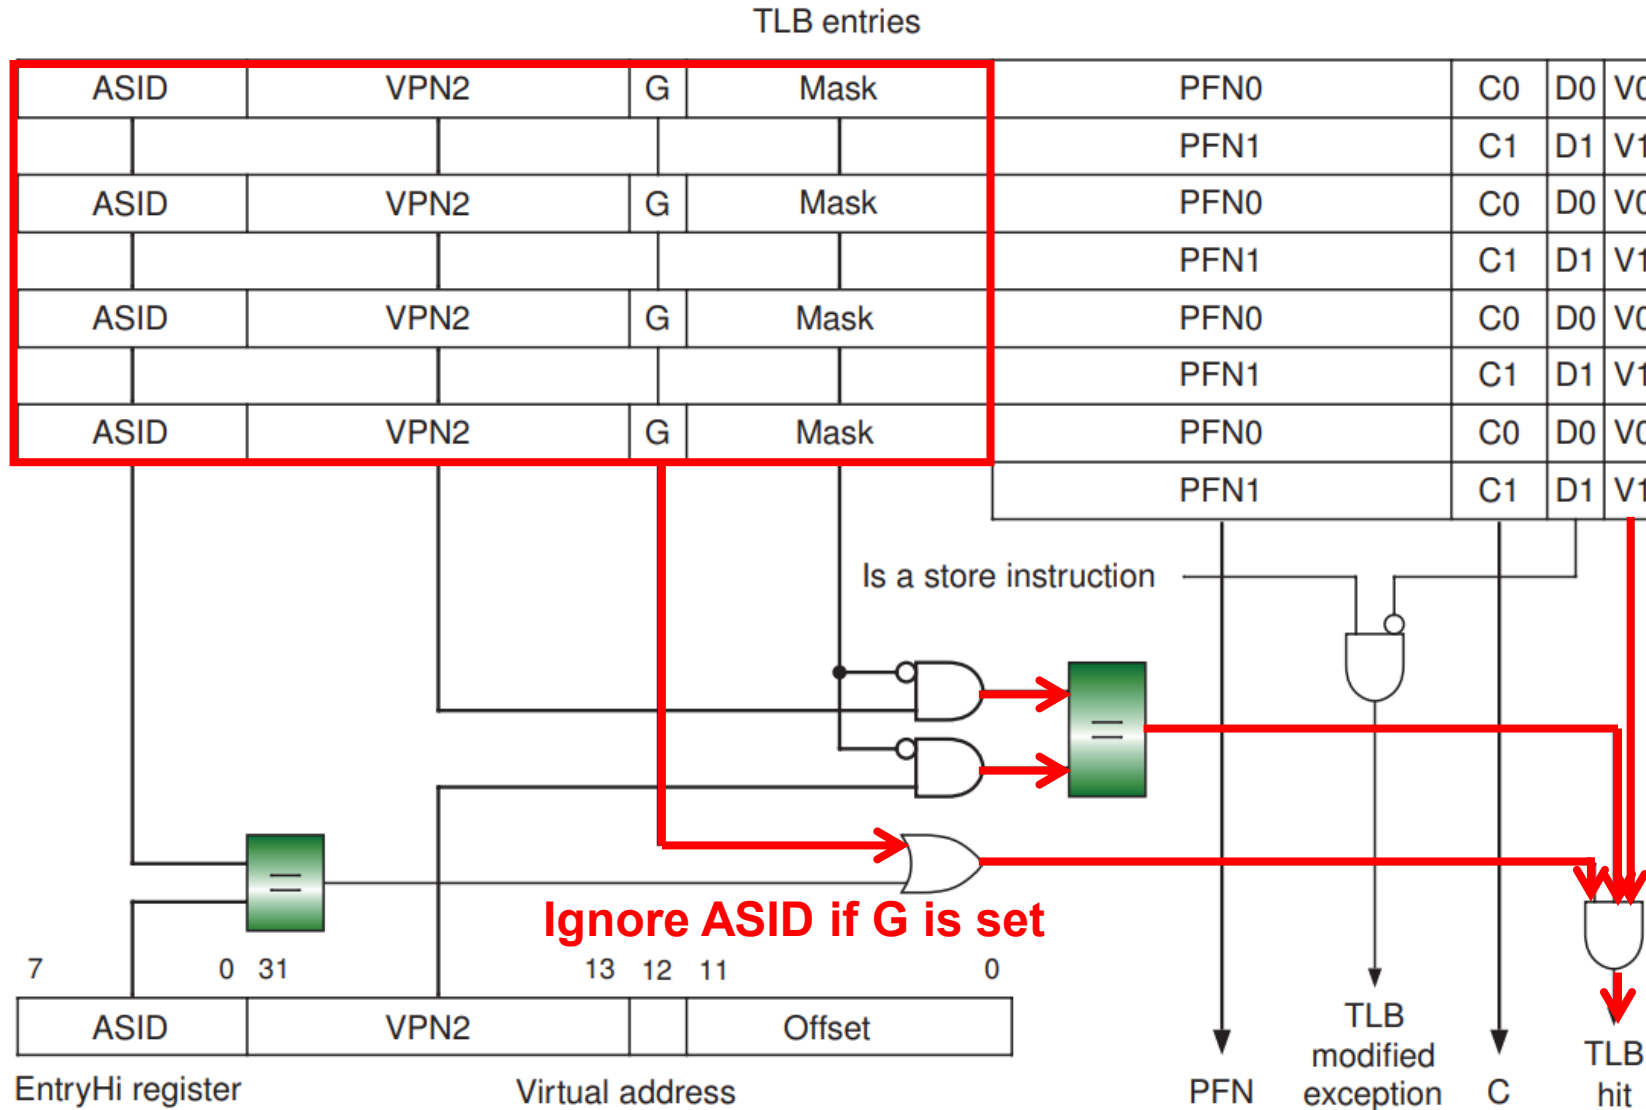
# TLB-Based MIPS Memory Management



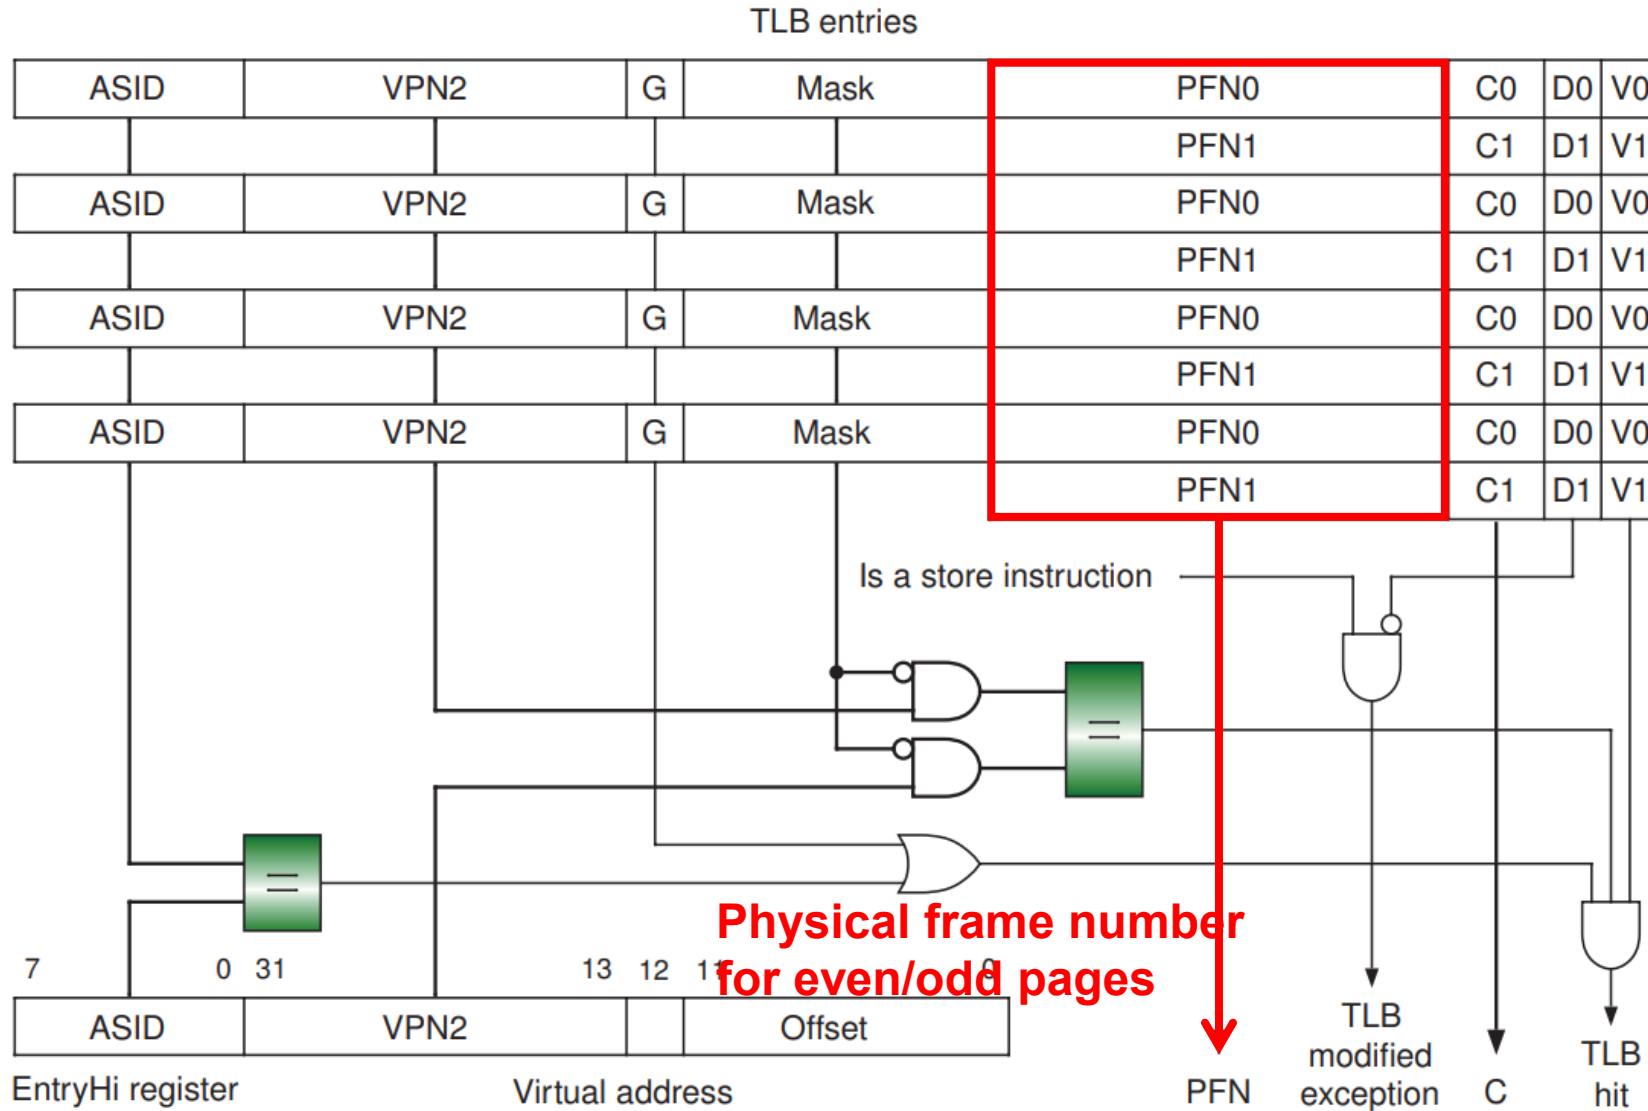
# TLB-Based MIPS Memory Management



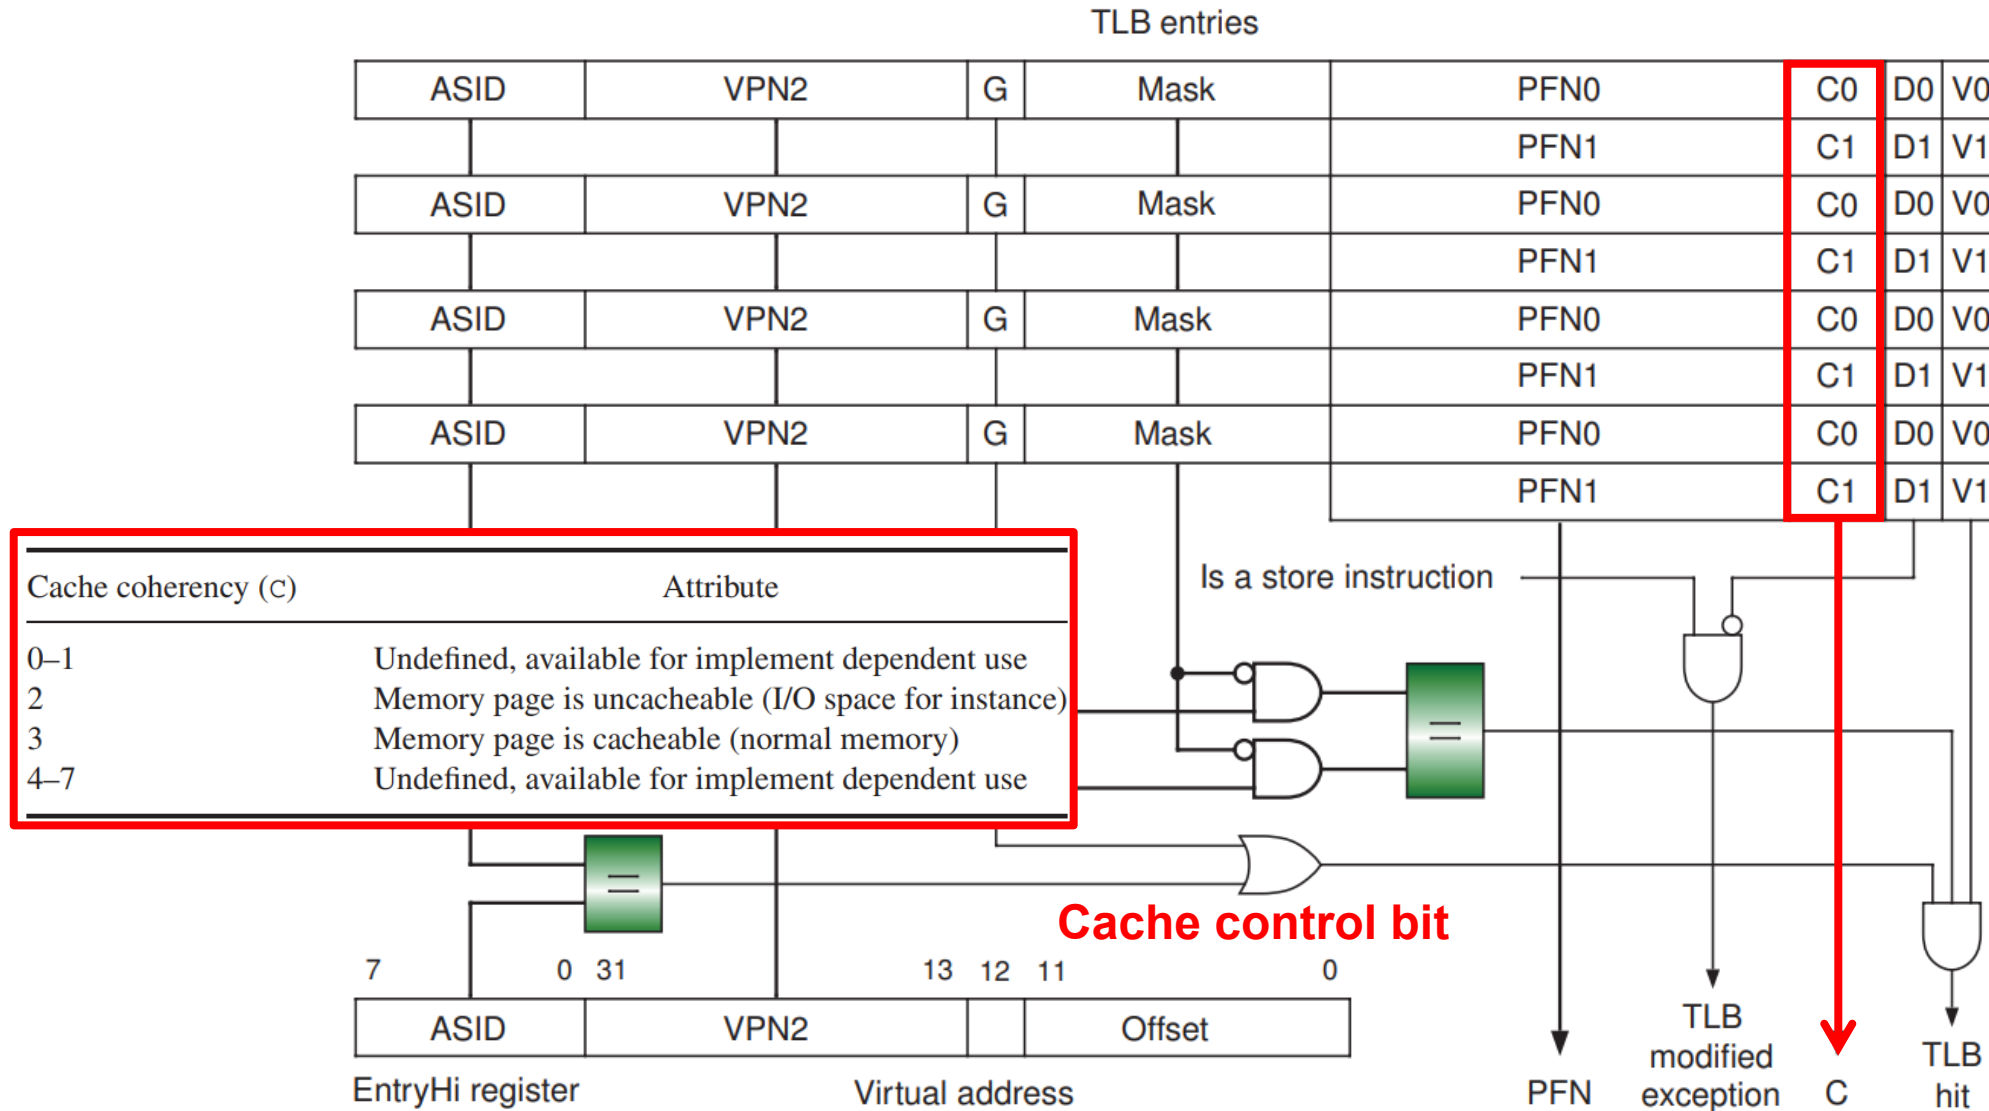
# TLB-Based MIPS Memory Management



# TLB-Based MIPS Memory Management

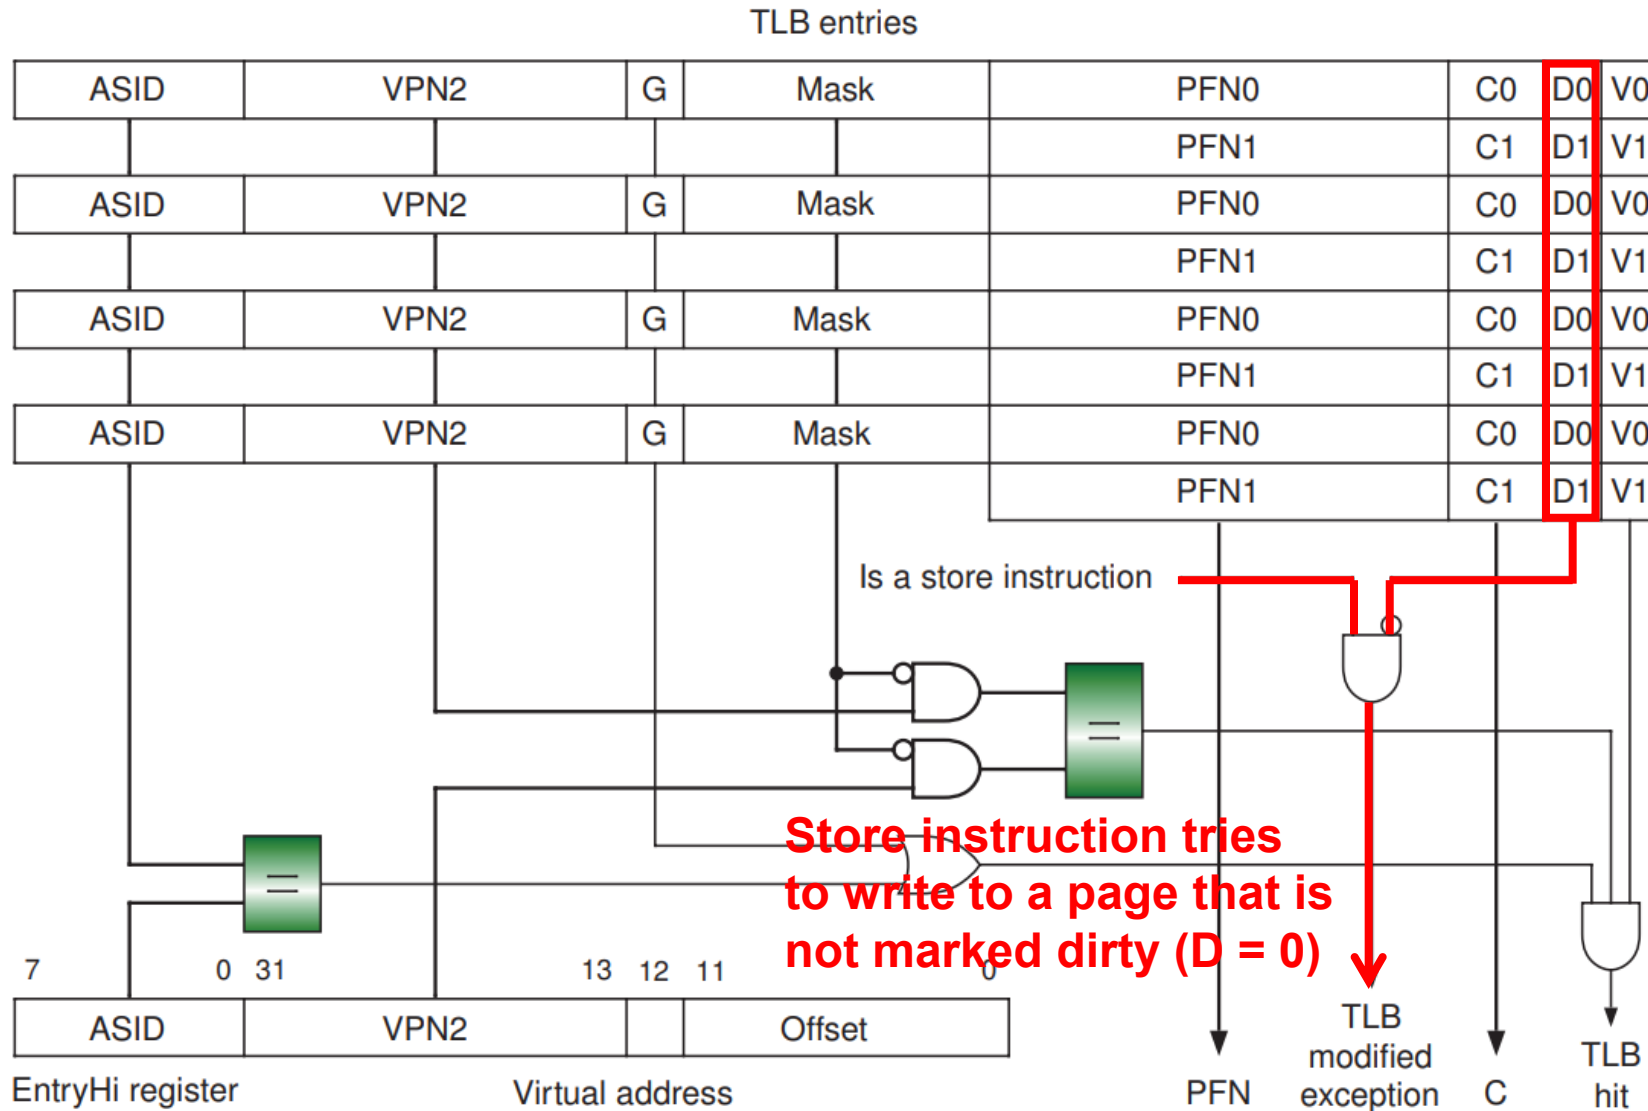


# TLB-Based MIPS Memory Management





# TLB-Based MIPS Memory Management



# TLB-Based MIPS Memory Management

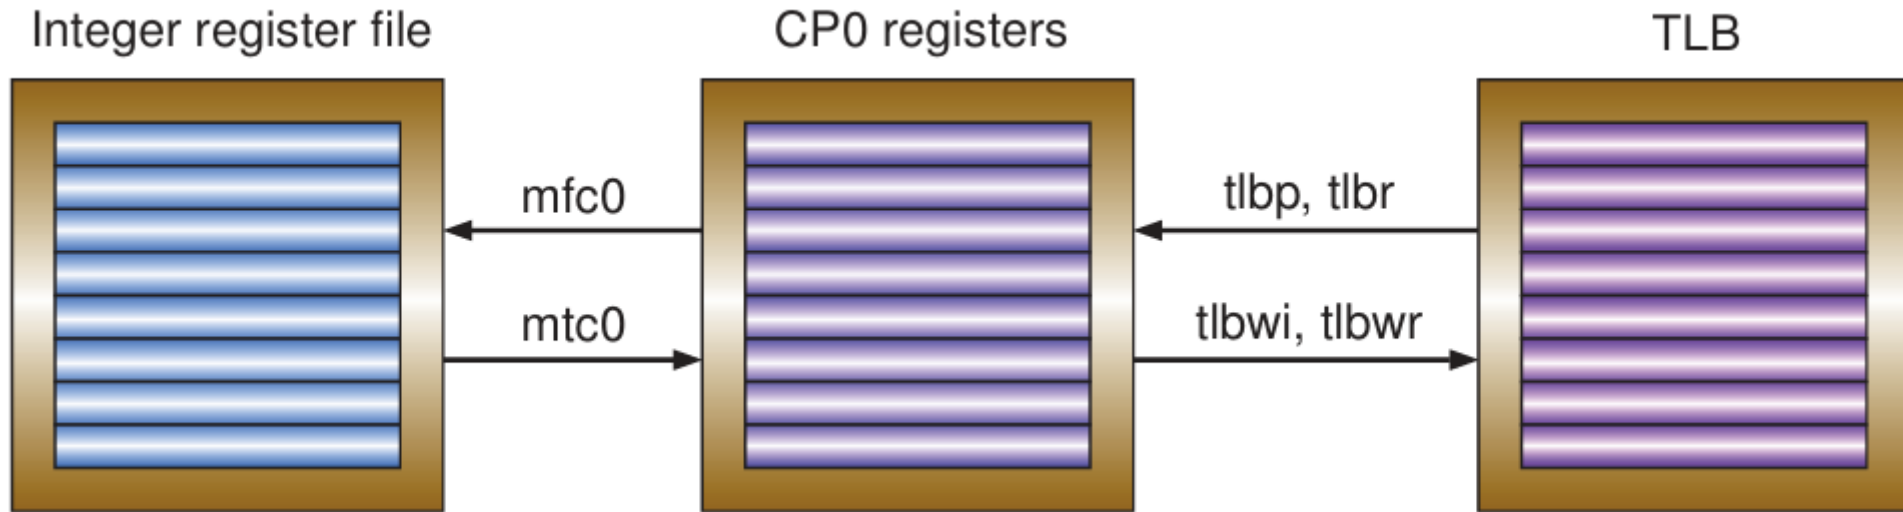
---

- The MIPS TLB is software-managed; there are instructions to manage the TLB's contents.
  - ✓ ***tlbp***
    - probes the TLB to see if a particular translation is in there and puts the search result on a CP0 register.
  - ✓ ***tlbr***
    - reads the contents of a TLB entry into CP0 registers.
  - ✓ ***tlbwi***
    - writes a specific TLB entry with the contents of CP0 registers.
  - ✓ ***tlbwr***
    - writes a random TLB entry with the contents of CP0 registers.
  - ✓ ***mfc0, mtc0***
    - transfer data between a CP0 register and an IU register in the general-purpose register file.

# TLB-Based MIPS Memory Management

---

- CPU manipulating TLB through CP0 registers



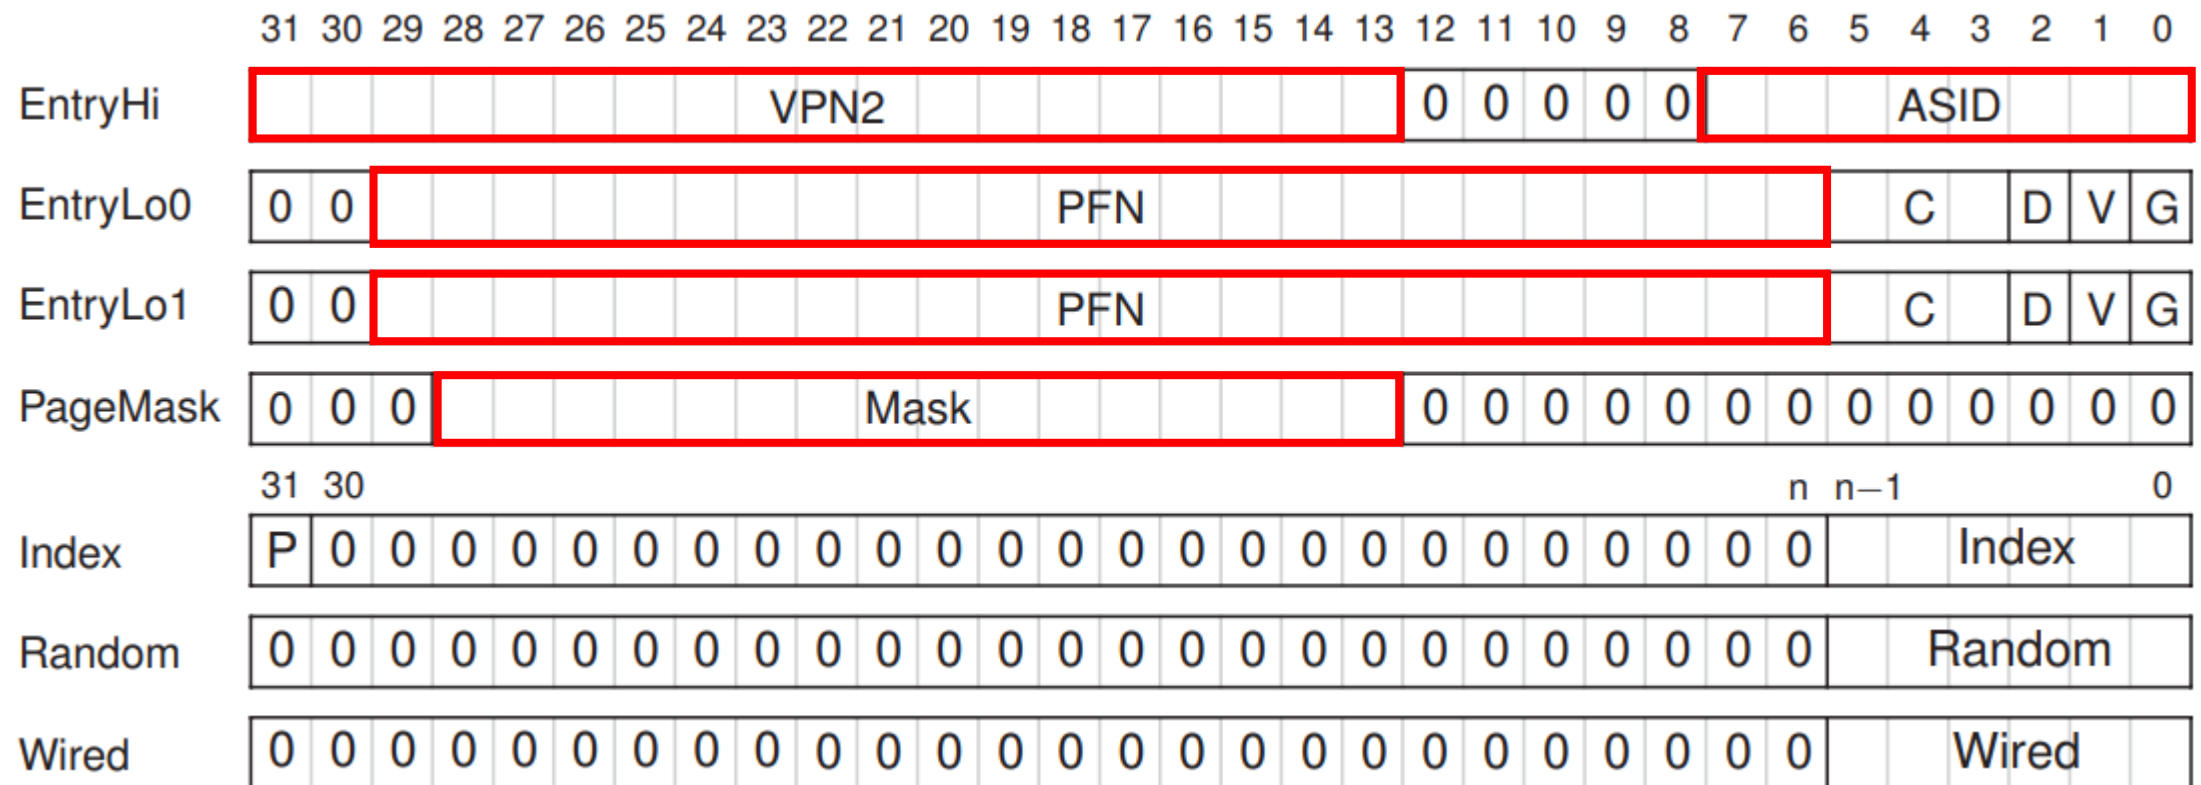
# TLB-Based MIPS Memory Management

- MIPS TLB manipulation CP0 registers

	31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
EntryHi																					0	0	0	0	0							
EntryLo0	0	0																														
EntryLo1	0	0																														
PageMask	0	0	0																		0	0	0	0	0	0	0	0	0	0	0	0
	31	30																														
Index	P	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0				
Random	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0				
Wired	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0					

# TLB-Based MIPS Memory Management

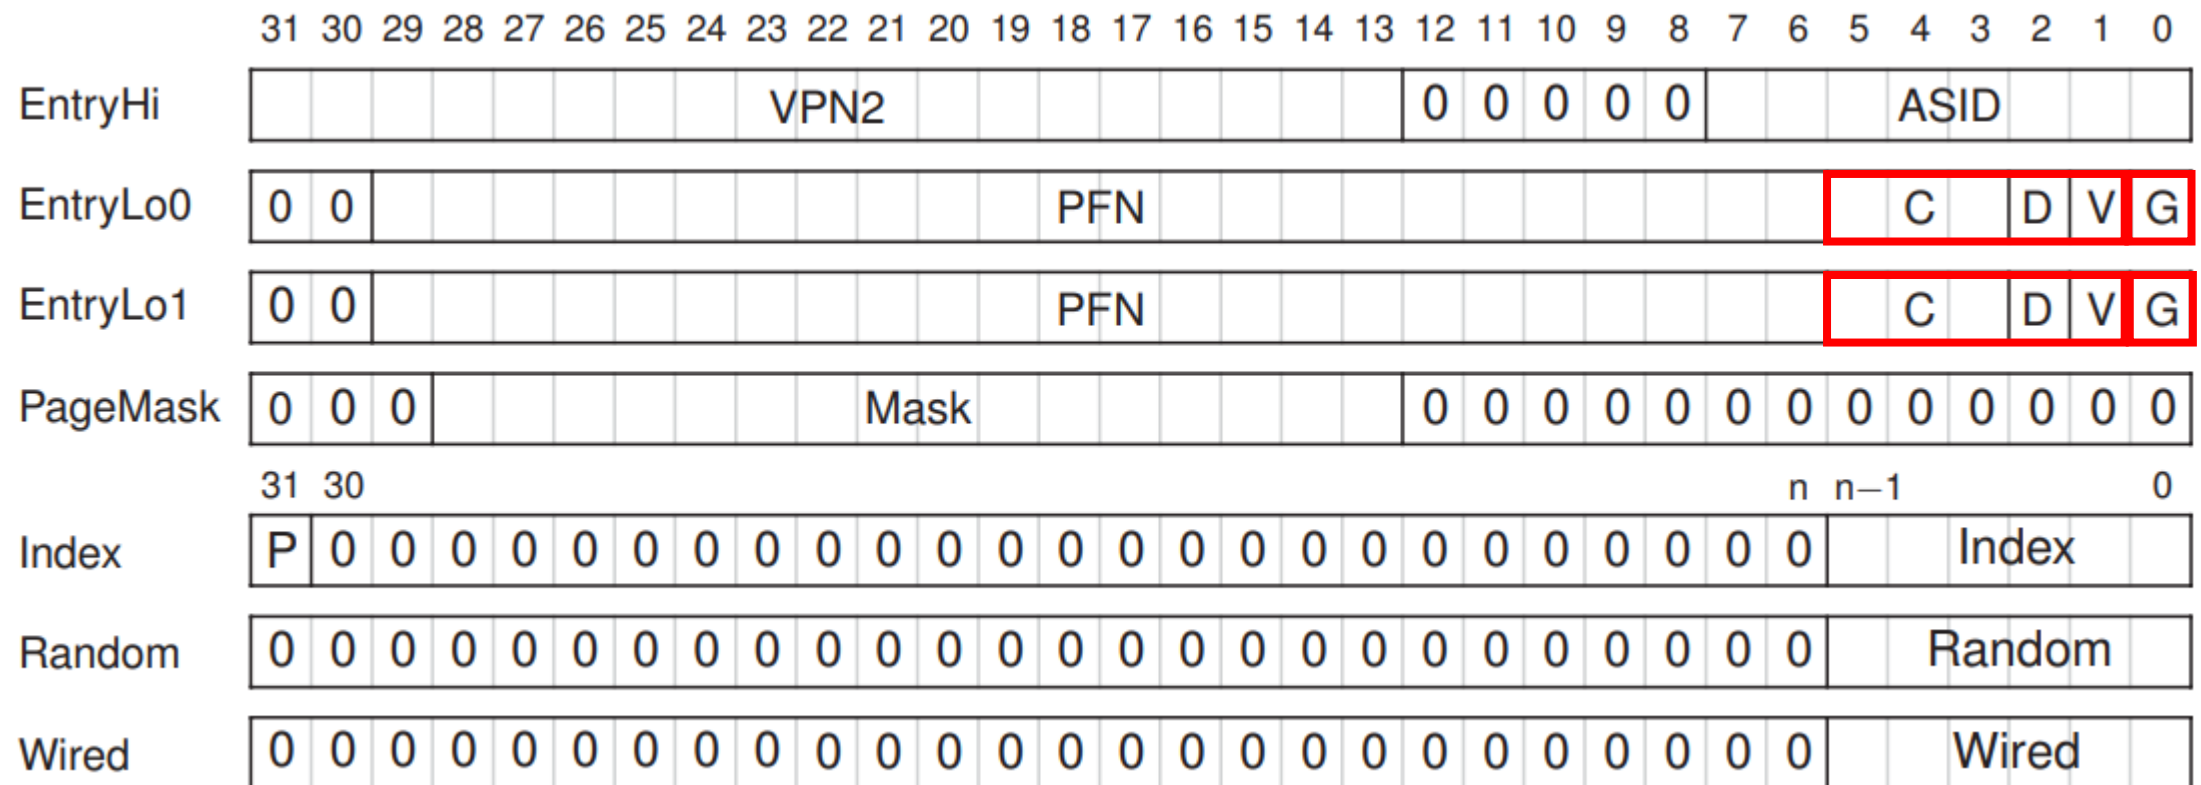
- MIPS TLB manipulation CP0 registers



Correspond exactly to the fields of the TLB entry

# TLB-Based MIPS Memory Management

- MIPS TLB manipulation CP0 registers



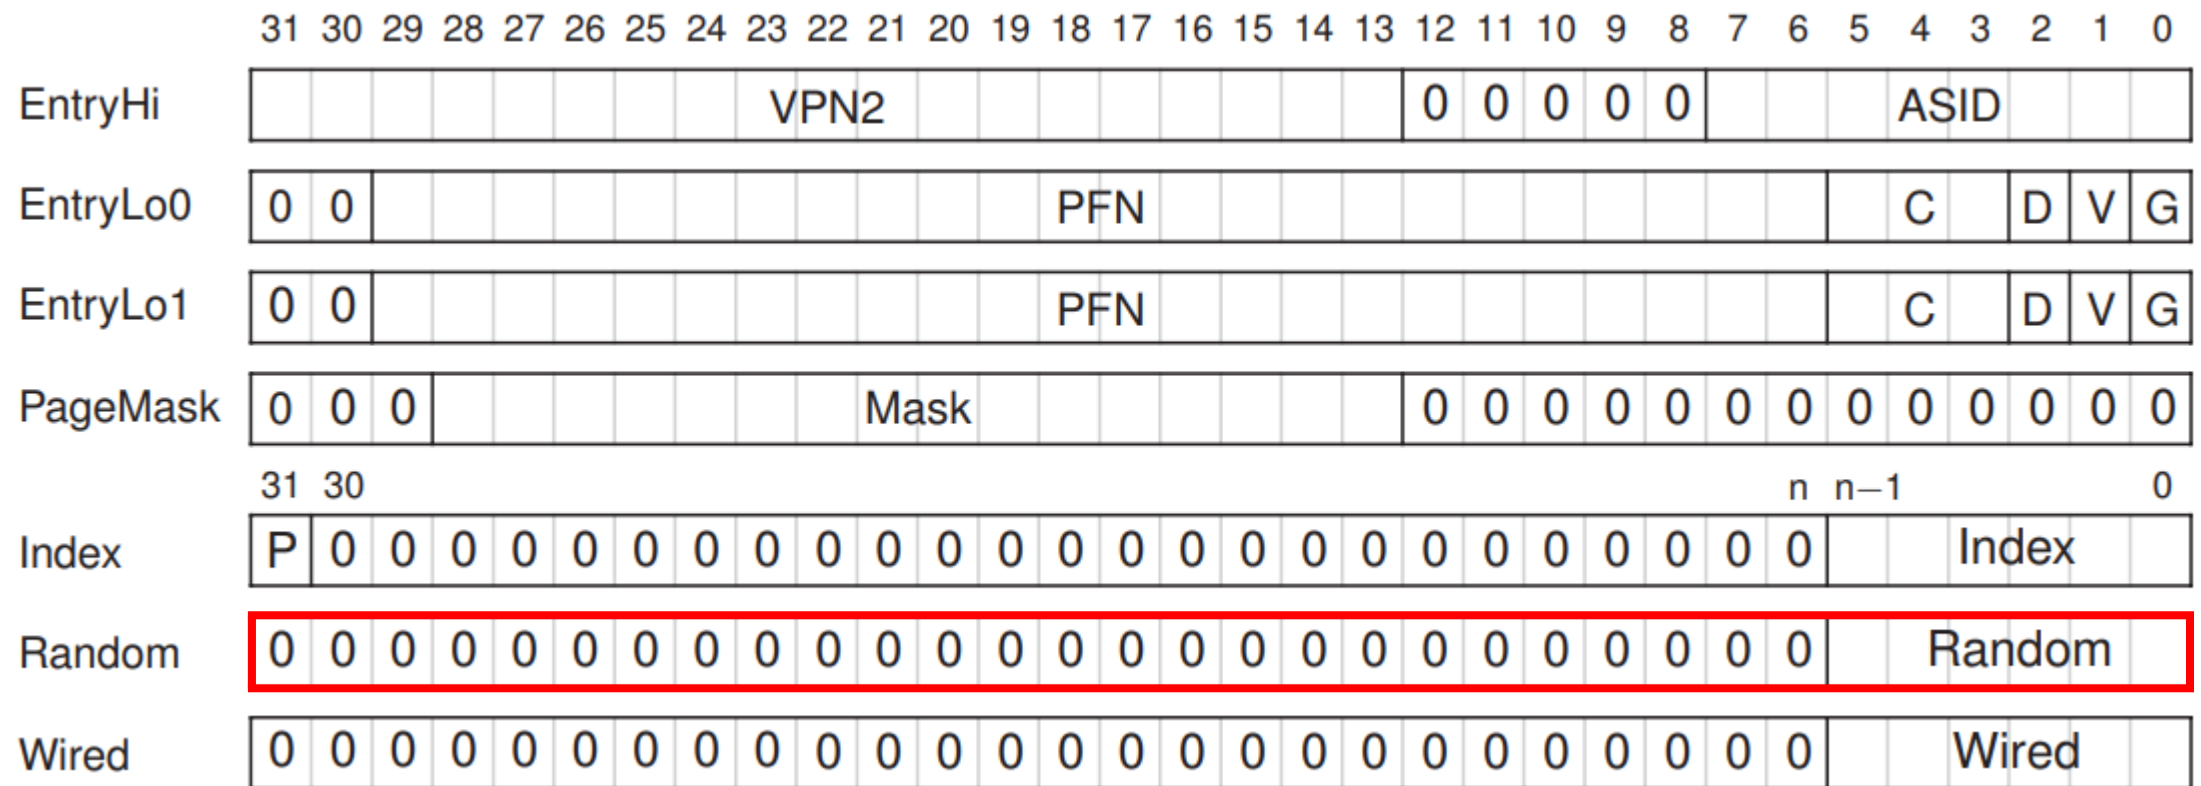
Two G bits (G0 and G1)

TLB write  $G = G0 \& G1$ , TLB read  $G0 = G1 = G$



# TLB-Based MIPS Memory Management

- MIPS TLB manipulation CP0 registers

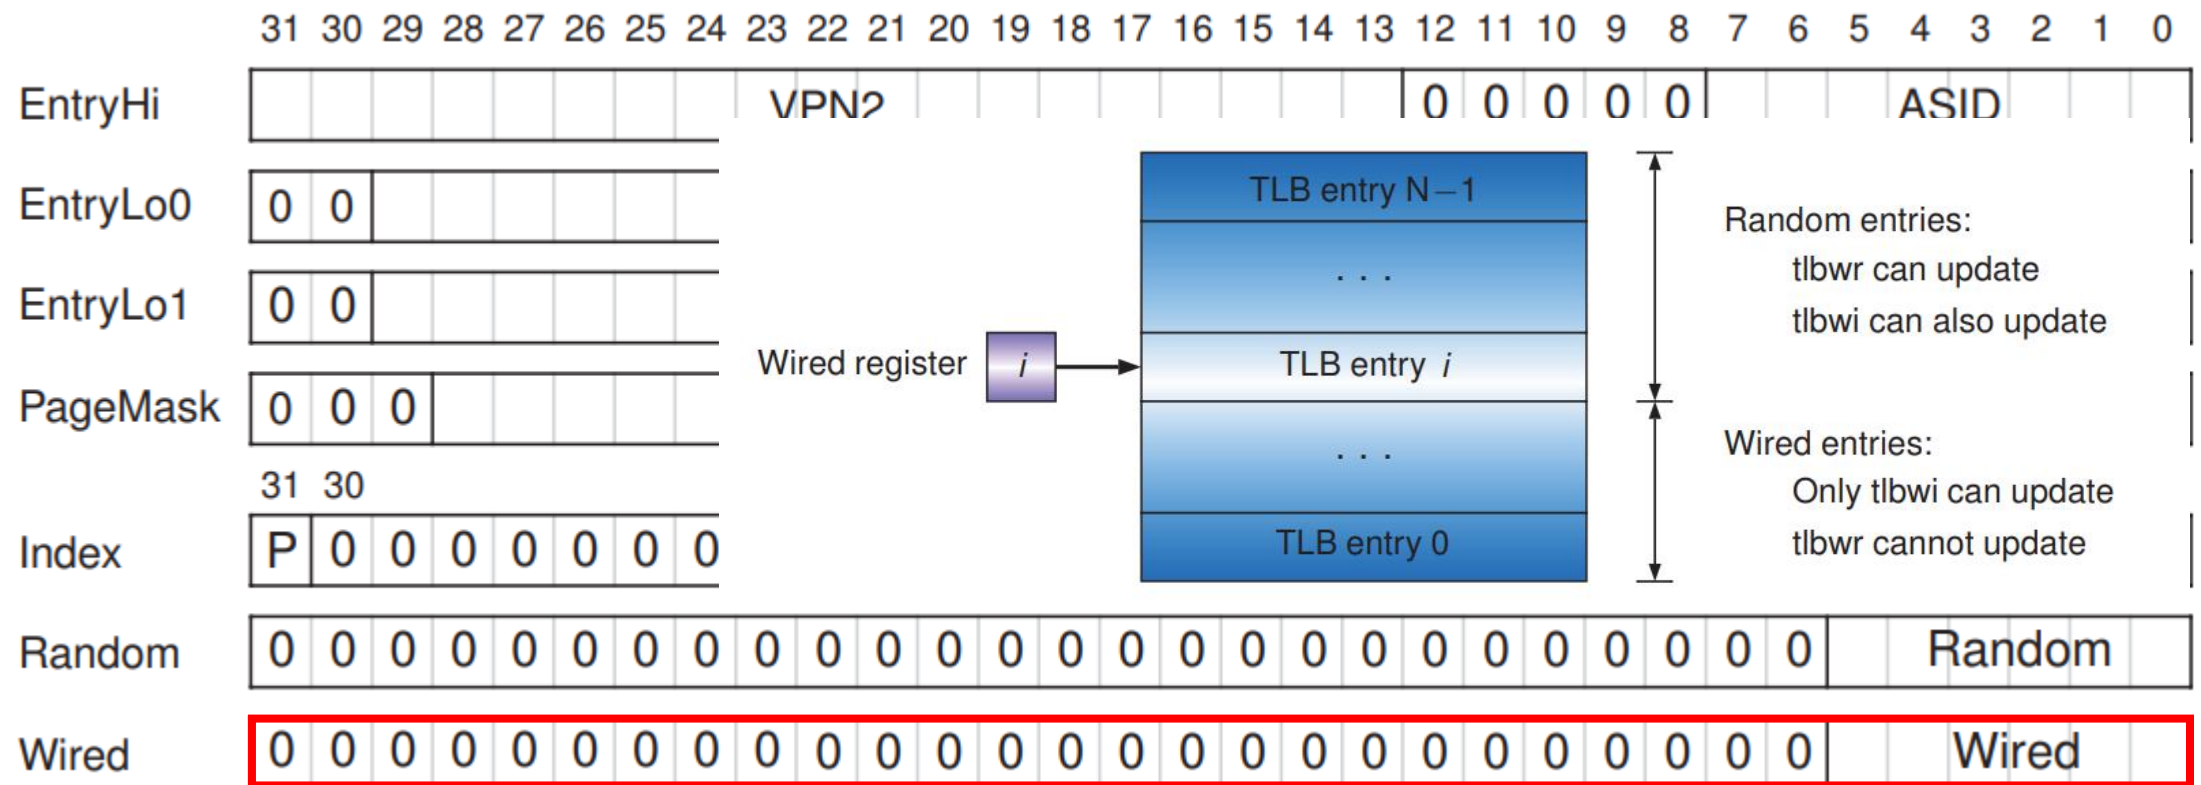


Random register holds a random number generated by HW that is used to index the TLB for the *tlbwr* instruction



# TLB-Based MIPS Memory Management


- MIPS TLB manipulation CP0 registers



**Wired register is used to specify the boundary between the wired (reserved) and random entries in the TLB**

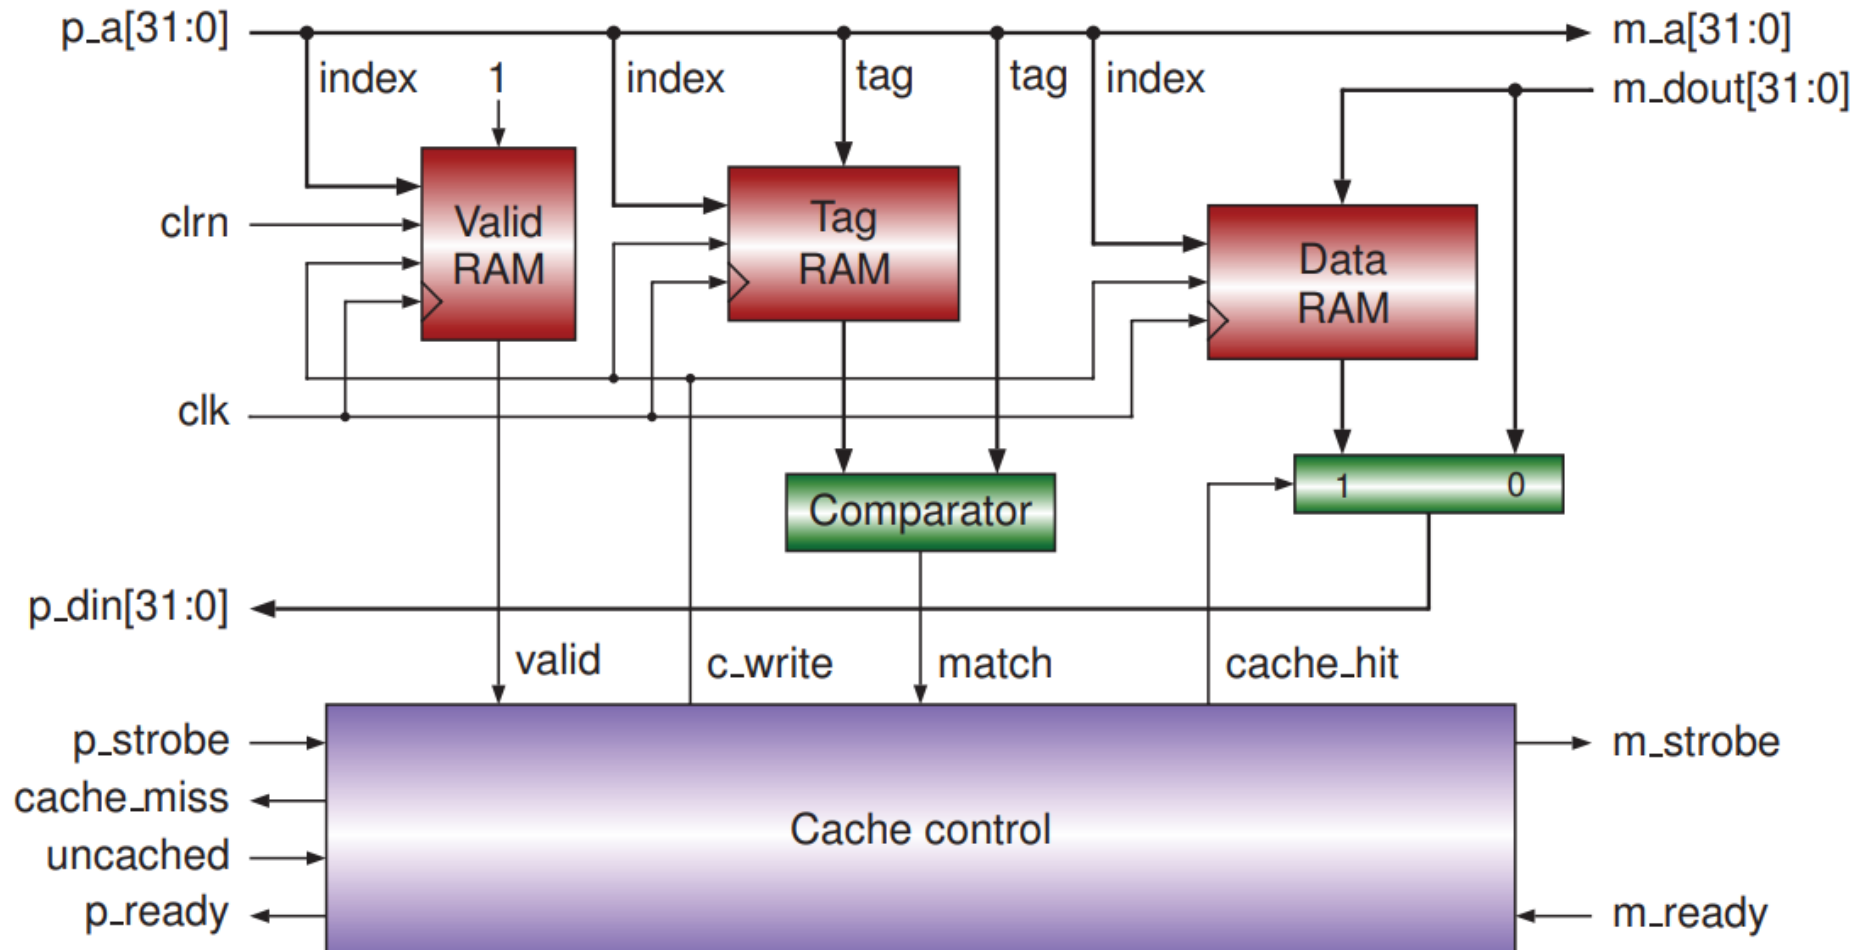
# Design of Pipelined CPU with Caches and TLB

- Overall structure of caches and TLBs



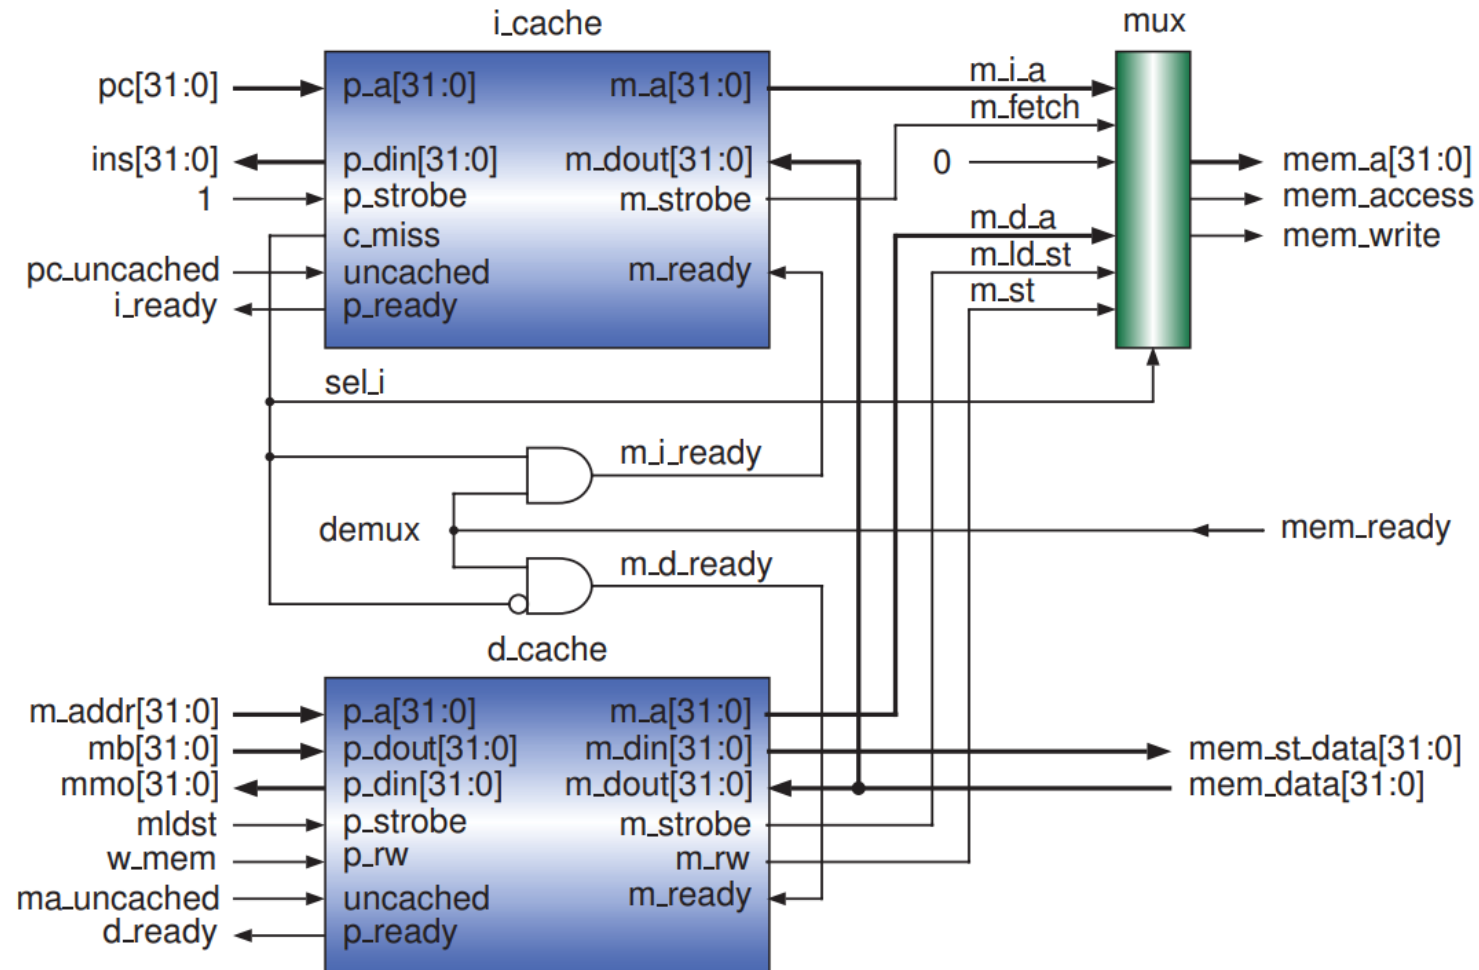
# Design of Pipelined CPU with Caches and TLB

- Instruction cache



# Design of Pipelined CPU with Caches and TLB

- Memory interface to instruction and data caches



# Design of Pipelined CPU with Caches and TLB

---

- **Instruction cache**
- **Processor interface**
  - ✓ p\_a (address), p\_din, p\_strobe (request), and p\_ready.
  - ✓ cache\_miss
- **Memory interface**
  - ✓ m\_a (address), m\_dout, m\_strobe (request), and m\_ready.

```
1  module i_cache (  
2      input    [31:0] p_a,  
3      output   [31:0] p_din,  
4      input    p_strobe,  
5      input    uncached,  
6      output   p_ready,  
7      output   cache_miss,  
8      input    clk, clrn,  
9      output   [31:0] m_a,  
10     input    [31:0] m_dout,  
11     output   m_strobe,  
12     input    m_ready  
13 );
```

# Design of Pipelined CPU with Caches and TLB

- **Instruction cache**
- d\_valid, d\_tags, d\_data are indexed from 0 to 63.
  - ✓ 64 cache lines.
- d\_data stores 32-bit instruction.
  - ✓ cache line size is 1 word (32 bits).
- p\_a (processor address) is divided into 3 parts.
  - ✓ index (p\_a[7:2]), tag (p\_a[31:8]), and byte offset (p\_a[1:0]).

```
14      reg          d_valid [0:63];
15      reg          [23:0] d_tags [0:63];
16      reg          [31:0] d_data [0:63];
17      wire          [5:0] index  = p_a[7:2];
18      wire          [23:0] tag    = p_a[31:8];
19      wire          c_write;
20      wire          [31:0] c_din;
```

# Design of Pipelined CPU with Caches and TLB

- **Instruction cache**
- d\_valid is cleared when clrn is asserted.
- c\_write is set, d\_valid bit for index is set, indicating the valid line.

```
22     integer i;  
23     always @(posedge clk or negedge clrn) begin  
24         if (!clrn) begin  
25             for (i=0; i<64; i=i+1) begin  
26                 d_valid[i] <= 0;  
27             end  
28         end else if (c_write) begin  
29             d_valid[index] <= 1;  
30         end  
31     end
```

# Design of Pipelined CPU with Caches and TLB

---

- **Instruction cache**
- `c_write` is set, tag and data are written specified by index.

```
33     always @(posedge clk) begin
34         if (c_write) begin
35             d_tags[index] <= tag;
36             d_data[index] <= c_din;
37         end
38     end
```



# Design of Pipelined CPU with Caches and TLB

## ▪ Instruction cache

- cache\_hit: p\_strobe is set, valid index, and tag match.
- cache\_miss: p\_strobe is set but invalid index or tag not match.
- m\_strobe: when cache\_miss, cache asserts request to main memory at the m\_a (= p\_a).
- p\_ready: cache\_hit or cache\_miss and m\_ready (main memory).
- c\_write: cache\_miss, not uncached, and m\_ready.
- p\_din: c\_dout if cahche\_hit or m\_dout if cache\_miss.

```
40 wire          valid  = d_valid[index];
41 wire [23:0]   tagout = d_tags[index];
42 wire [31:0]   c_dout = d_data[index];
43 wire          cache_hit = p_strobe & valid & (tagout == tag);
44 assign        cache_miss = p_strobe & (!valid | (tagout != tag));
45 assign        m_a       = p_a;
46 assign        m_strobe  = cache_miss;
47 assign        p_ready   = cache_hit | cache_miss & m_ready;
48 assign        c_write   = cache_miss & ~uncached & m_ready;
49 assign        c_din     = m_dout;
50 assign        p_din     = cache_hit ? c_dout : m_dout;
```

# Design of Pipelined CPU with Caches and TLB

## ▪ Data cache

```
1  module d_cache ( // direct mapping, 2^6 blocks, 1 word/block, write through
2      input  [31:0] p_a, // cpu address
3      input  [31:0] p_dout, // cpu data out to mem
4      output [31:0] p_din, // cpu data in from mem
5      input    p_strobe, // cpu strobe
6      input    p_rw, // cpu read/write command
7      input    uncached, // uncached
8      output    p_ready, // ready (to cpu)
9      input    clk, clrn, // clock and reset
10     output [31:0] m_a, // mem address
11     input  [31:0] m_dout, // mem data out to cpu
12     output [31:0] m_din, // mem data in from cpu
13     output    m_strobe, // mem strobe
14     output    m_rw, // mem read/write
15     input    m_ready // mem ready
16 );
17     reg        d_valid [0:63]; // 1-bit valid
18     reg [23:0] d_tags [0:63]; // 24-bit tag
19     reg [31:0] d_data [0:63]; // 32-bit data
20     wire [23:0] tag = p_a[31:8]; // address tag
21     wire [31:0] c_din; // data to cache
22     wire [5:0] index = p_a[7:2]; // block index
23     wire        c_write; // cache write
```

# Design of Pipelined CPU with Caches and TLB

- **Data cache**
- d\_valid is cleared when clrn is asserted.
- c\_write is set, valid, tags, and data are written.

```
25     integer i;
26     always @(posedge clk or negedge clrn) begin
27         if (!clrn) begin
28             for (i=0; i<64; i=i+1) begin
29                 d_valid[i] <= 0;
30             end
31         end else if (c_write) begin
32             d_valid[index] <= 1;
33         end
34     end
35
36     always @(posedge clk) begin
37         if (c_write) begin
38             d_tags[index] <= tag;
39             d_data[index] <= c_din;
40         end
41     end
```

```
17     reg          d_valid [0:63];          // 1-bit valid
18     reg [23:0]    d_tags  [0:63];          // 24-bit tag
19     reg [31:0]    d_data  [0:63];          // 32-bit data
20     wire [23:0]   tag = p_a[31:8];          // address tag
21     wire [31:0]   c_din;                    // data to cache
22     wire [5:0]    index = p_a[7:2];         // block index
23     wire          c_write;                  // cache write
```

# Design of Pipelined CPU with Caches and TLB

---

- **Data cache**
- **cache\_hit**: p\_strobe is set, valid index, and tag match.
- **cache\_miss**: p\_strobe is set but invalid index or tag not match.

```
43      wire      valid    = d_valid[index];
44      wire      [23:0] tagout = d_tags[index];
45      wire      [31:0] c_dout = d_data[index];
46      wire      cache_hit  = p_strobe & valid & (tagout == tag);
47      wire      cache_miss = p_strobe & ( !valid | (tagout != tag) );
```

# Design of Pipelined CPU with Caches and TLB

---

- **Data cache**
- m\_din: directly passes p\_dout (CPU's write data) to memory.
- m\_a: directly passes p\_a (CPU's address) to memory.
- m\_rw: directly passes p\_rw (CPU's rd/wr command) to memory.
- m\_strobe: Main memory is strobes for a transaction if
  - ✓ p\_rw is set, because of the write-through policy.
  - ✓ or a cache miss

```
48      assign  m_din      = p_dout;
49      assign  m_a        = p_a;
50      assign  m_rw       = p_rw;
51      assign  m_strobe   = p_rw | cache_miss;
```

# Design of Pipelined CPU with Caches and TLB

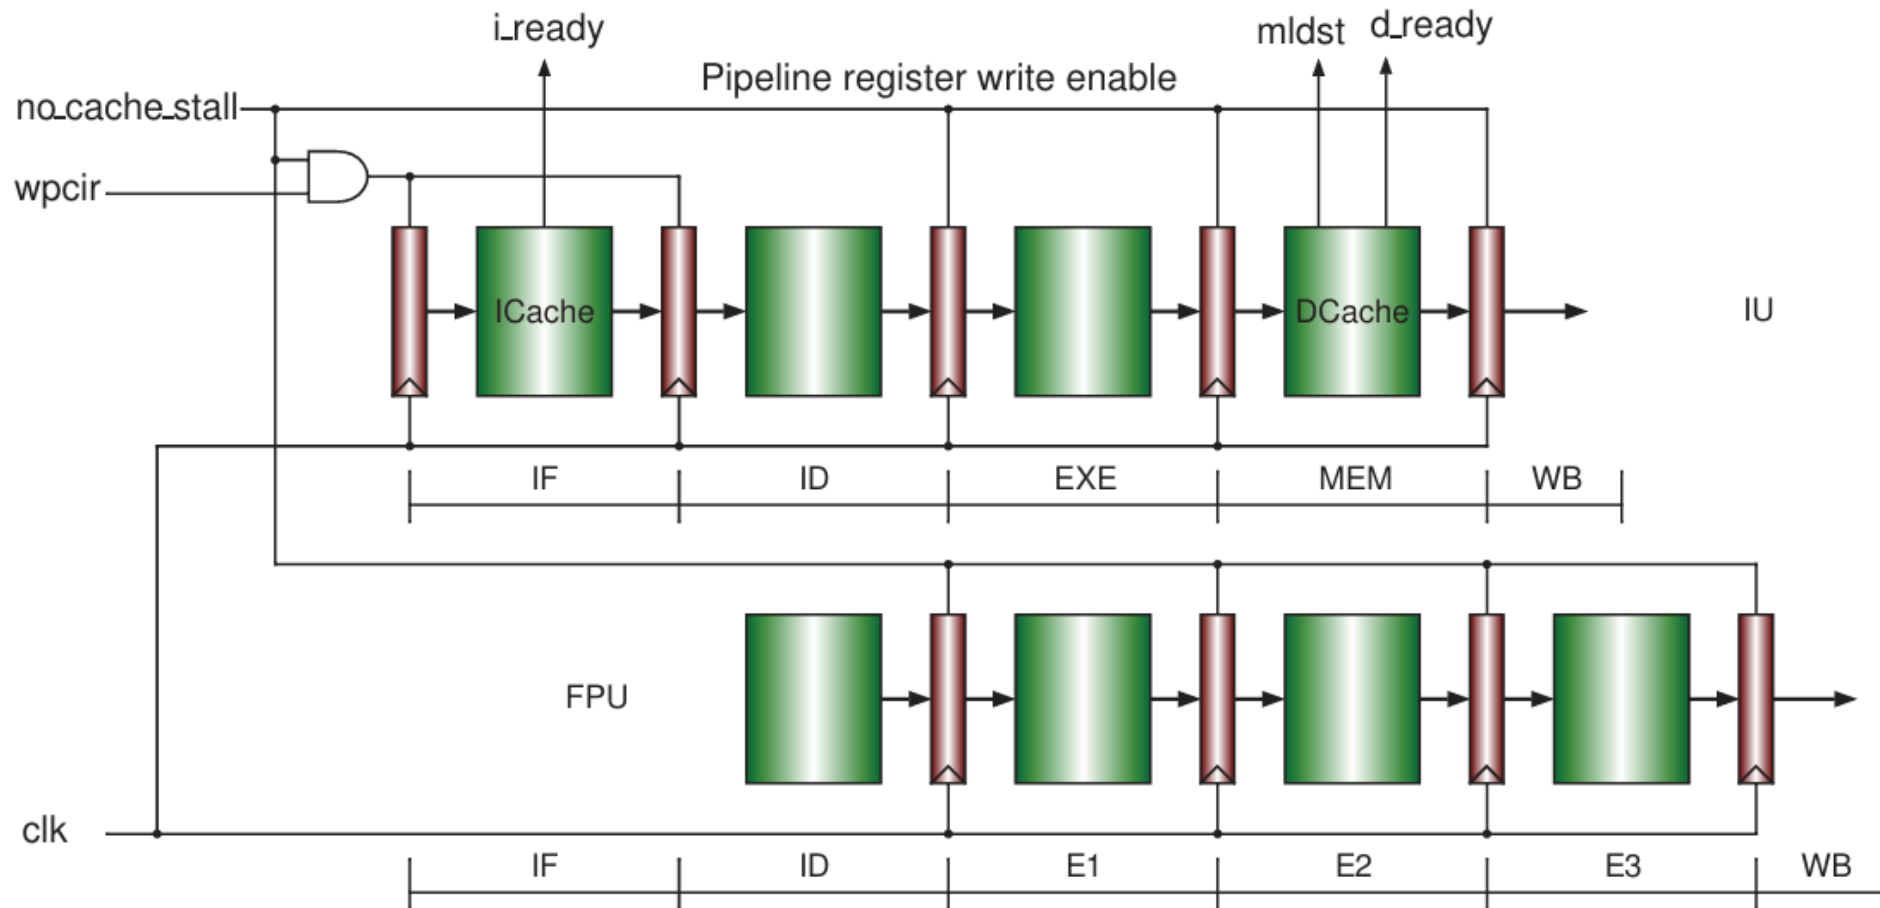
---

- **Data cache**
- For reads ( $\sim p\_rw$ ),
  - ✓  $p\_ready$  is asserted on a cache hit.
- For writes ( $p\_rw$ ) or read with miss,
  - ✓  $p\_ready$  is asserted when  $m\_ready$  is set (main memory).
- **c\_write**
  - ✓ uncached, CPU write ( $p\_rw$ ) or miss has been resolved.
- **c\_din**
  - ✓ For writing ( $p\_rw$ ),  $p\_dout$  (the data from the CPU).
  - ✓ For reading ( $\sim p\_rw$ ),  $m\_dout$  (the data from the main memory).

```
52    assign p_ready    = ~p_rw & cache_hit | (cache_miss | p_rw) & m_ready;
53    assign c_write    = ~uncached & (p_rw | cache_miss & m_ready);
54    assign c_din      = p_rw      ? p_dout : m_dout;
55    assign p_din      = cache_hit ? c_dout : m_dout;
```

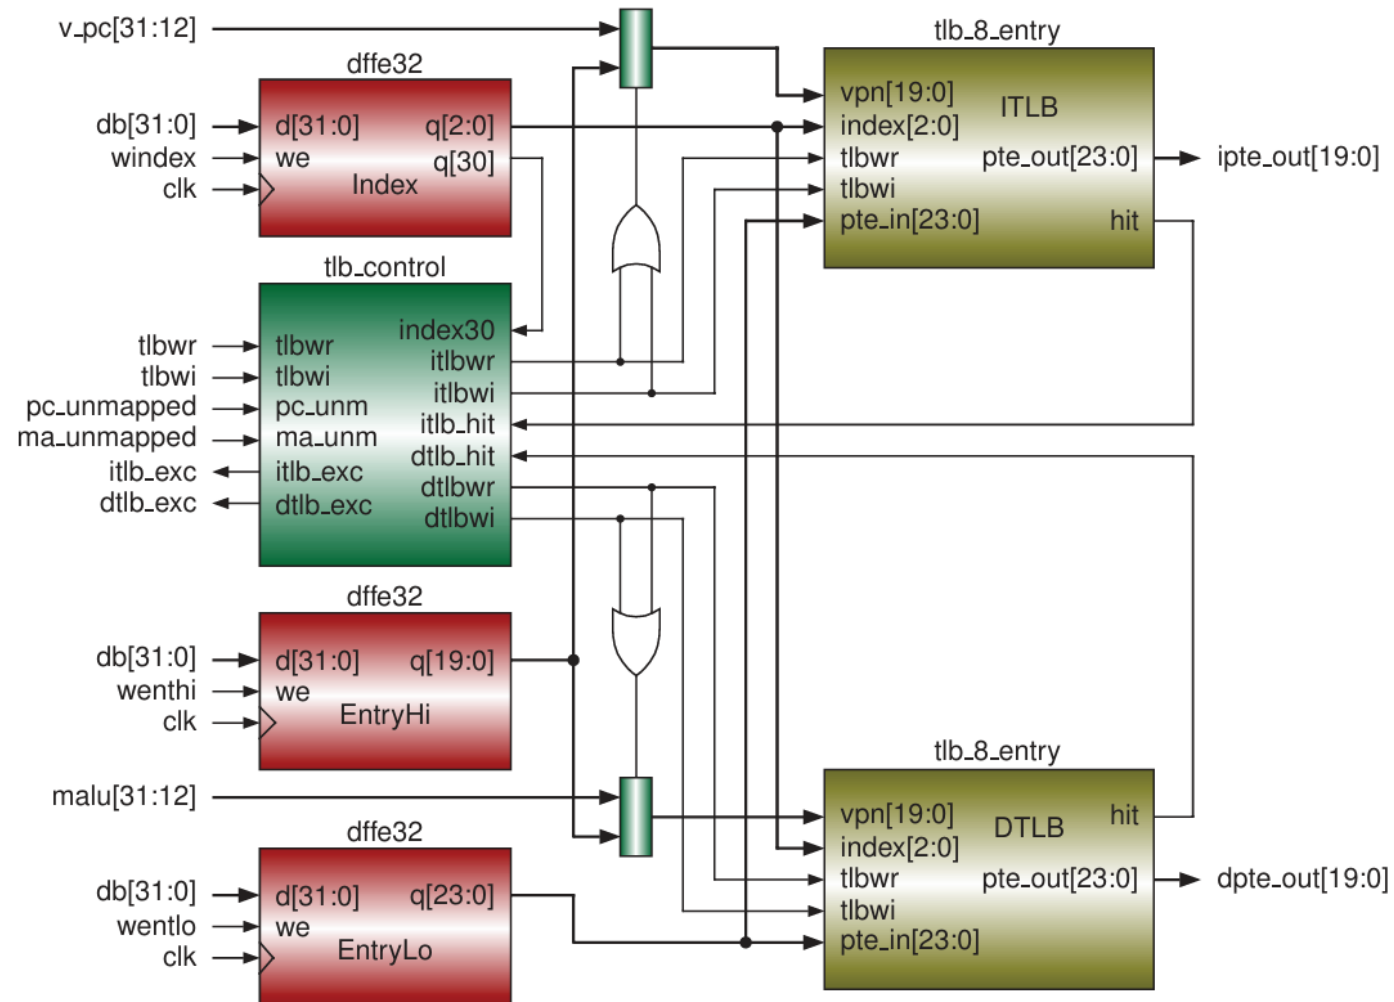
# Design of Pipelined CPU with Caches and TLB

- Pipeline halt circuit for cache misses



# Design of Pipelined CPU with Caches and TLB

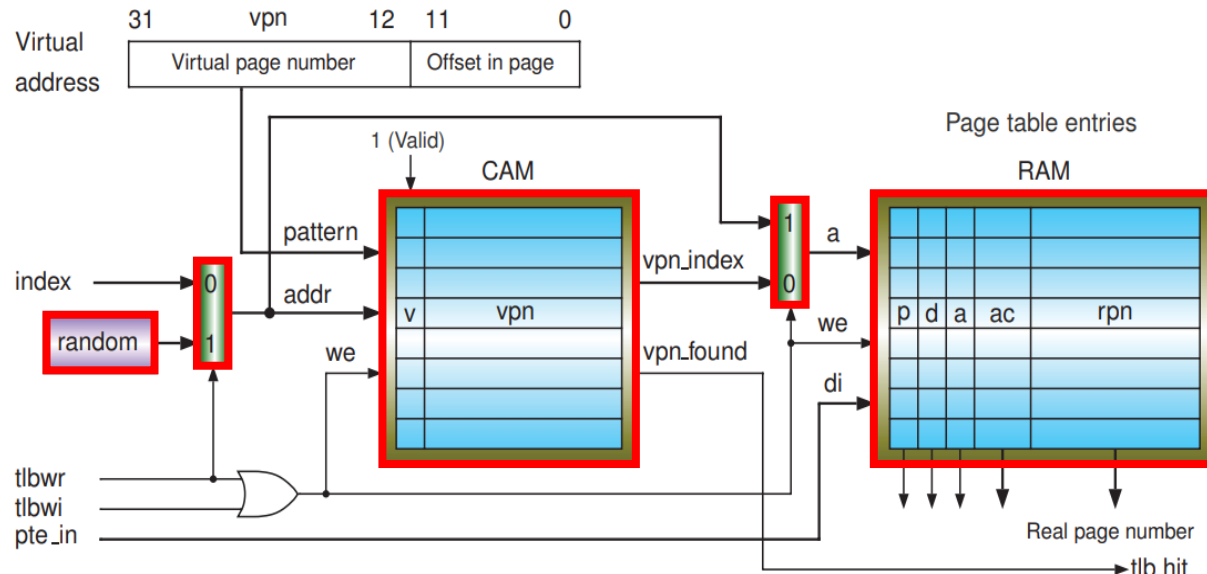
## ITLB and DTLB





# Design of Pipelined CPU with Caches and TLB

## ITLB and DTLB

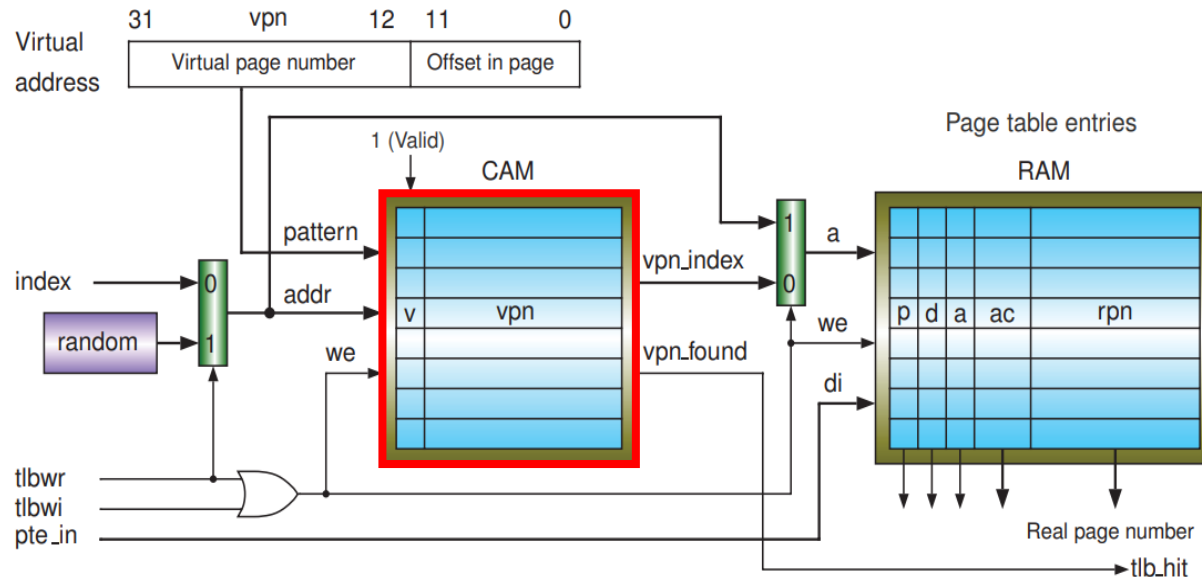


```

1 module tlb_8_entry (
2     input    [23:0] pte_in,
3     input    tlbwi,
4     input    tlbwr,
5     input    [2:0] index,
6     input    [19:0] vpn,
7     input    clk, clrn,
8     output   [23:0] pte_out,
9     output   tlb_hit
10 );
11 wire    [2:0] random;
12 wire    [2:0] w_idx;
13 wire    [2:0] ram_idx;
14 wire    [2:0] vpn_index;
15 wire      tlbw = tlbwi | tlbwr;
16 rand3    rdm (clk, clrn, random);
17 mux2x3    w_address (index, random, tlbwr, w_idx);
18 mux2x3    ram_address (vpn_index, w_idx, tlbw, ram_idx);
19 ram8x24    rpn (ram_idx, pte_in, clk, tlbw, pte_out);
20 cam8x21    valid_tag (clk, vpn, w_idx, tlbw, vpn_index, tlb_hit);
21 endmodule
    
```

# Design of Pipelined CPU with Caches and TLB

## ■ CAM

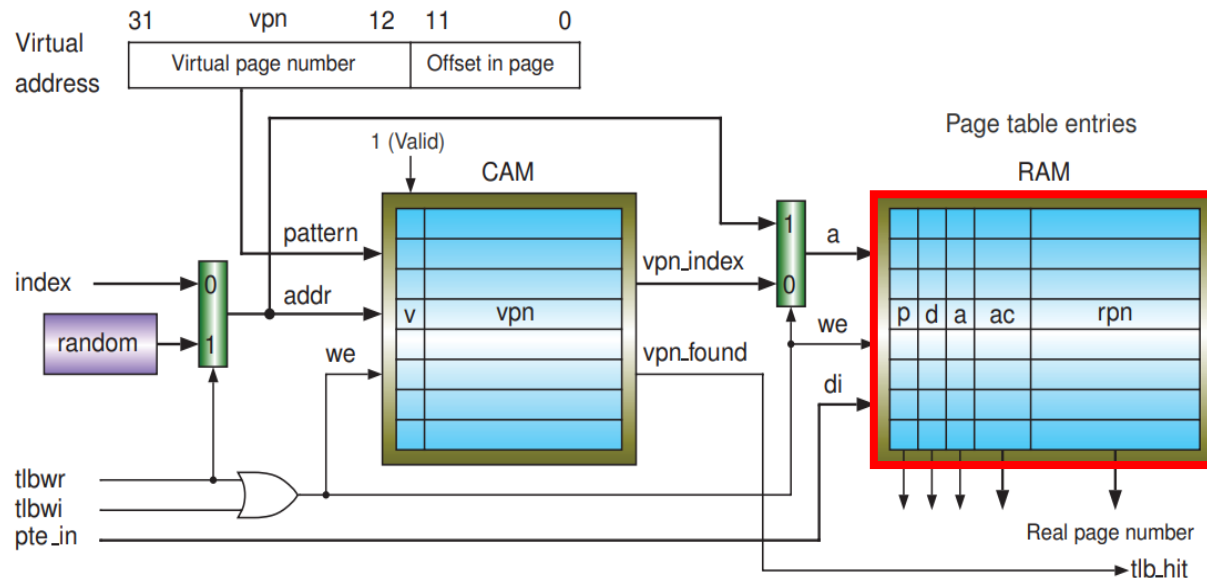


```

1 module cam8x21 (
2     input          clk,
3     input          wren,
4     input [19:0]   pattern,
5     input [2:0]    wraddress,
6     output [2:0]   maddress,
7     output         mfound
8 );
9
10 // write cam, update a line with pattern, valid bit <- 1
11 reg [20:0] ram [0:7];
12 always @(posedge clk) begin
13     if (wren) ram[waddress] <= {1'b1,pattern};
14 end
15
16 // fully associative search, should be implemented with CAM cells
17 wire [7:0] match_line;
18 assign match_line[7] = (ram[7] == {1'b1,pattern});
19 assign match_line[6] = (ram[6] == {1'b1,pattern});
20 assign match_line[5] = (ram[5] == {1'b1,pattern});
21 assign match_line[4] = (ram[4] == {1'b1,pattern});
22 assign match_line[3] = (ram[3] == {1'b1,pattern});
23 assign match_line[2] = (ram[2] == {1'b1,pattern});
24 assign match_line[1] = (ram[1] == {1'b1,pattern});
25 assign match_line[0] = (ram[0] == {1'b1,pattern});
26 assign mfound = |match_line;
27
28 // encoder for matched address, no multiple-match is allowed
29 assign maddress[2] = match_line[7] | match_line[6] |
30 match_line[5] | match_line[4];
31 assign maddress[1] = match_line[7] | match_line[6] |
32 match_line[3] | match_line[2];
33 assign maddress[0] = match_line[7] | match_line[5] |
34 match_line[3] | match_line[1];
35
36 // initialize cam, mainly clear valid bit of each line
37 integer i;
38 initial begin
39     for (i=0; i<8; i=i+1)
40         ram[i] = 0;
41 end
42 endmodule
  
```

# Design of Pipelined CPU with Caches and TLB

## ■ RAM

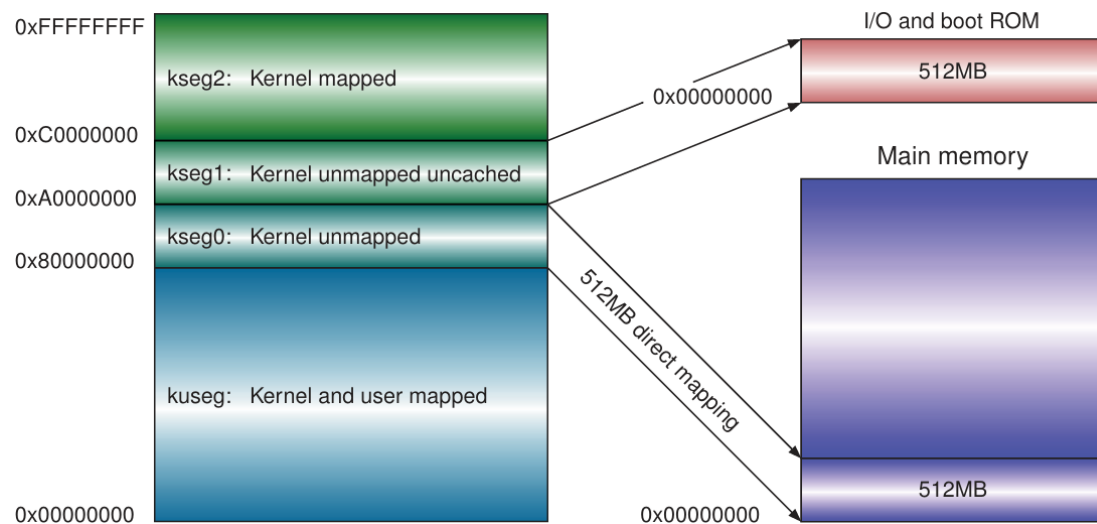


```

1  module ram8x24 (
2      input      [2:0]  address,
3      input      [23:0] data,
4      input      clk,
5      input      we,
6      output     [23:0] q
7  );
8      reg        [23:0] ram [0:7];
9      always @(posedge clk) begin
10         if (we) ram[address] <= data;
11     end
12     assign q = ram[address];
13
14     integer i;
15     initial begin
16         for (i=0 ; i<8; i=i+1)
17             ram[i] = 24'h0;
18     end
19 endmodule
    
```

# Design of Pipelined CPU with Caches and TLB

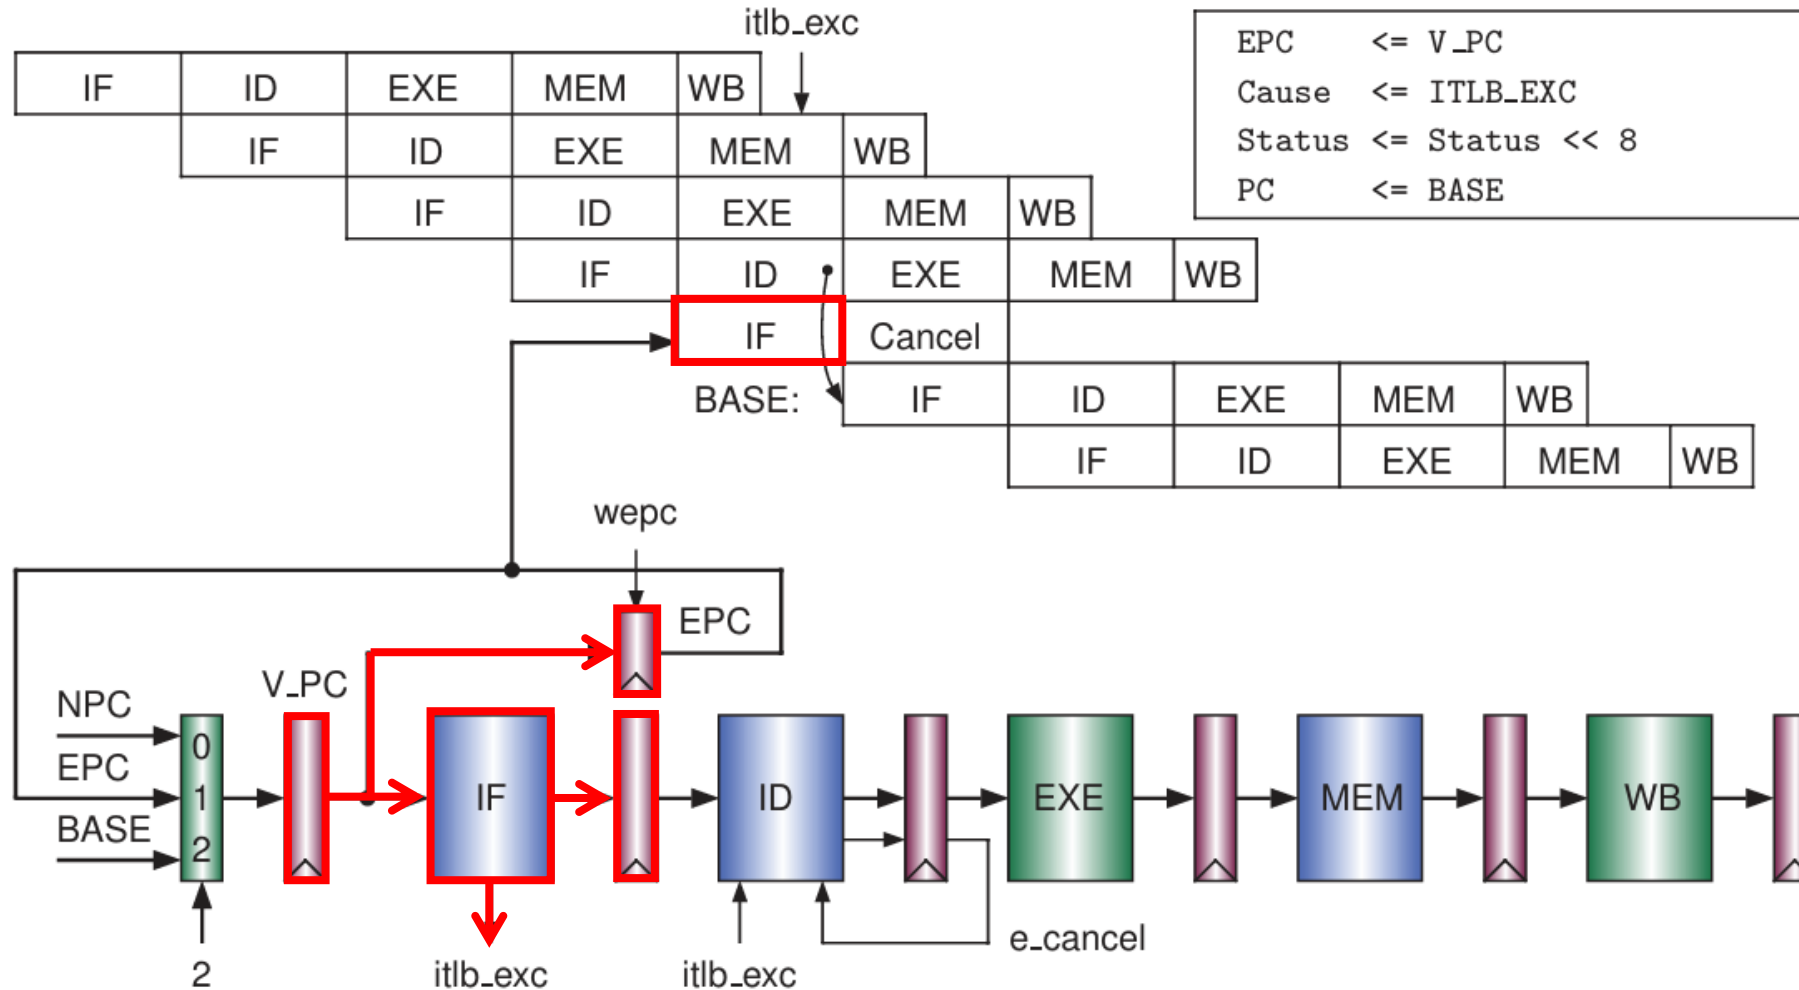
- Address mapping of the MIPS architecture



Segment	Virtual address	Physical address	TLBed	Cached	I/O
kuseg	00000000 ... 0000 01111111 ... 1111	00100000 ... 0000 –	Yes	Yes	No
kseg0	10000000 ... 0000 10011111 ... 1111	00000000 ... 0000 00011111 ... 1111	No	Yes	No
kseg1	10100000 ... 0000 10111111 ... 1111	00000000 ... 0000 00011111 ... 1111	No	No	Yes
kseg2	11000000 ... 0000 11111111 ... 1111	00100000 ... 0000 –	Yes	Yes	No

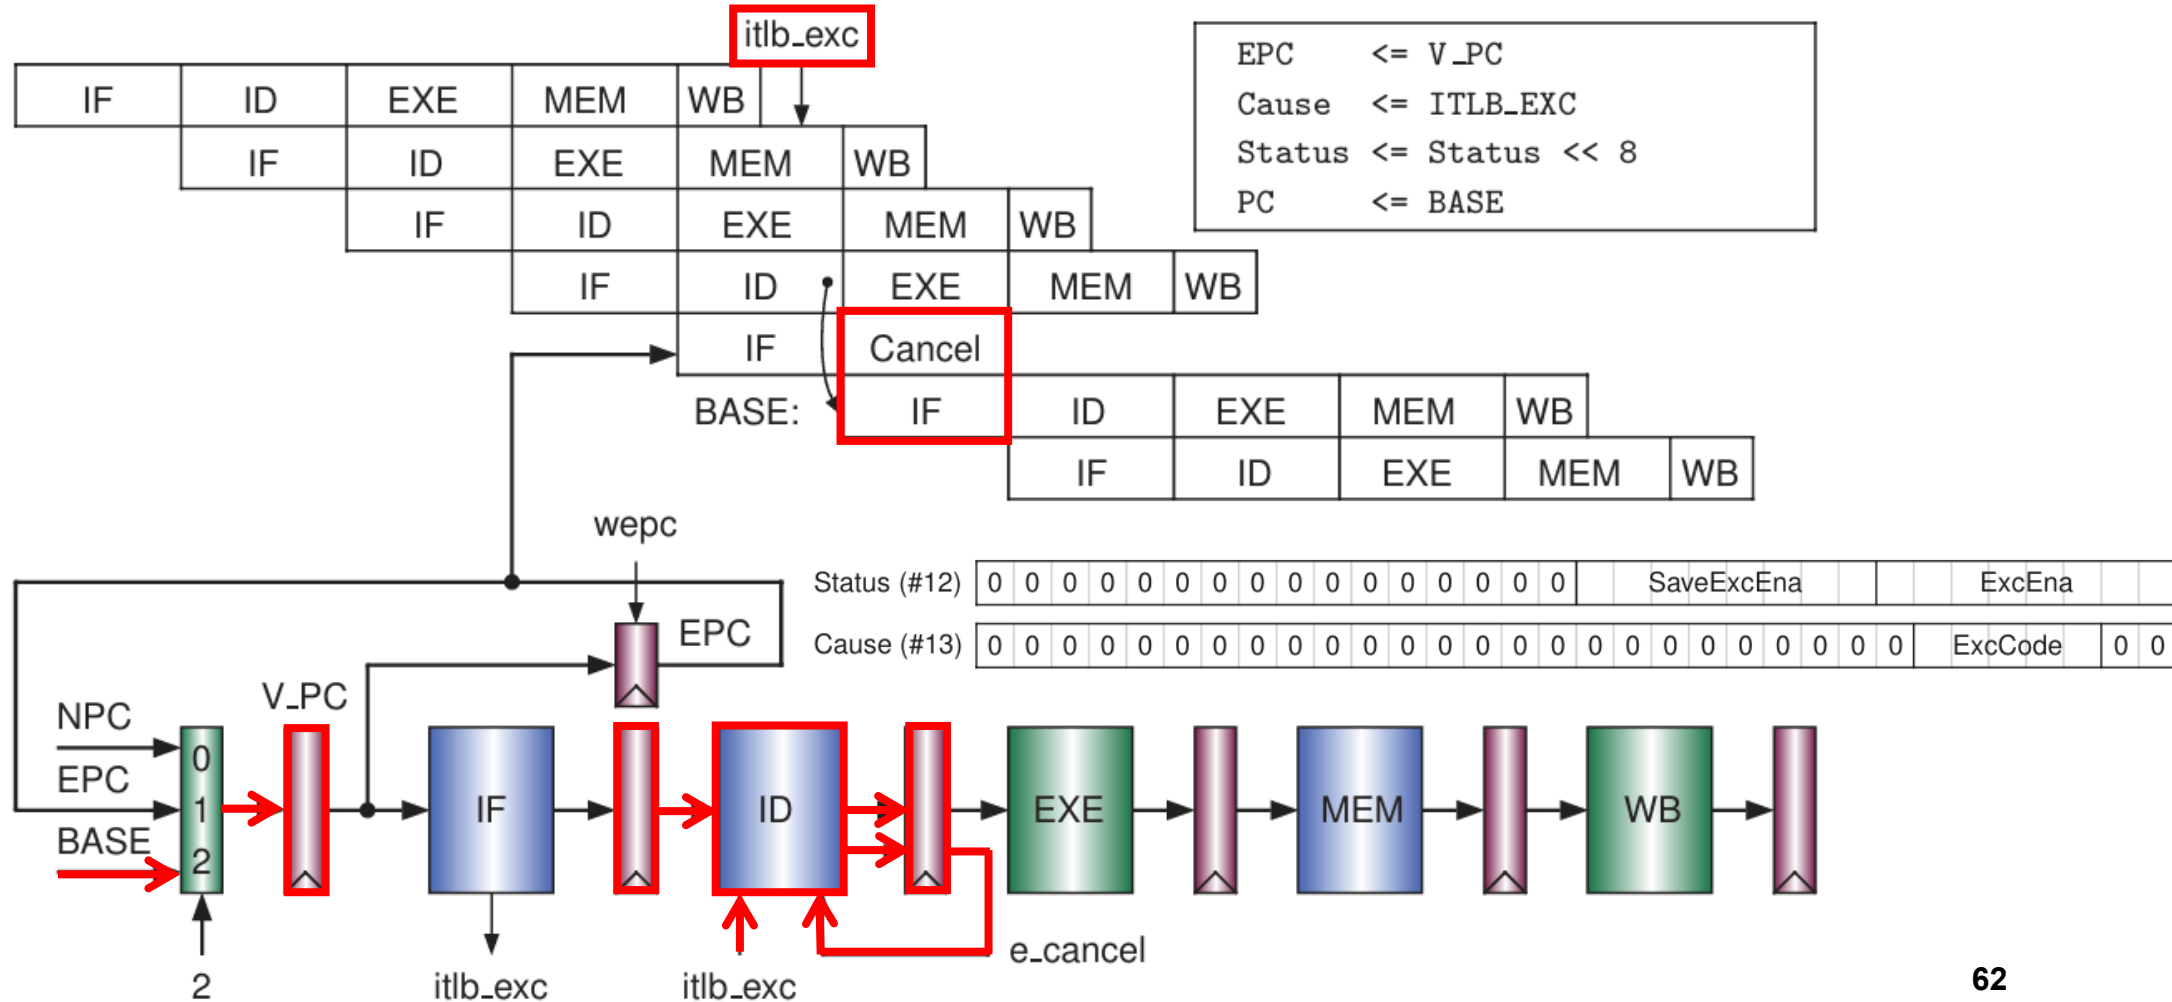
# Design of Pipelined CPU with Caches and TLB

- Registers related to ITLB miss exceptions



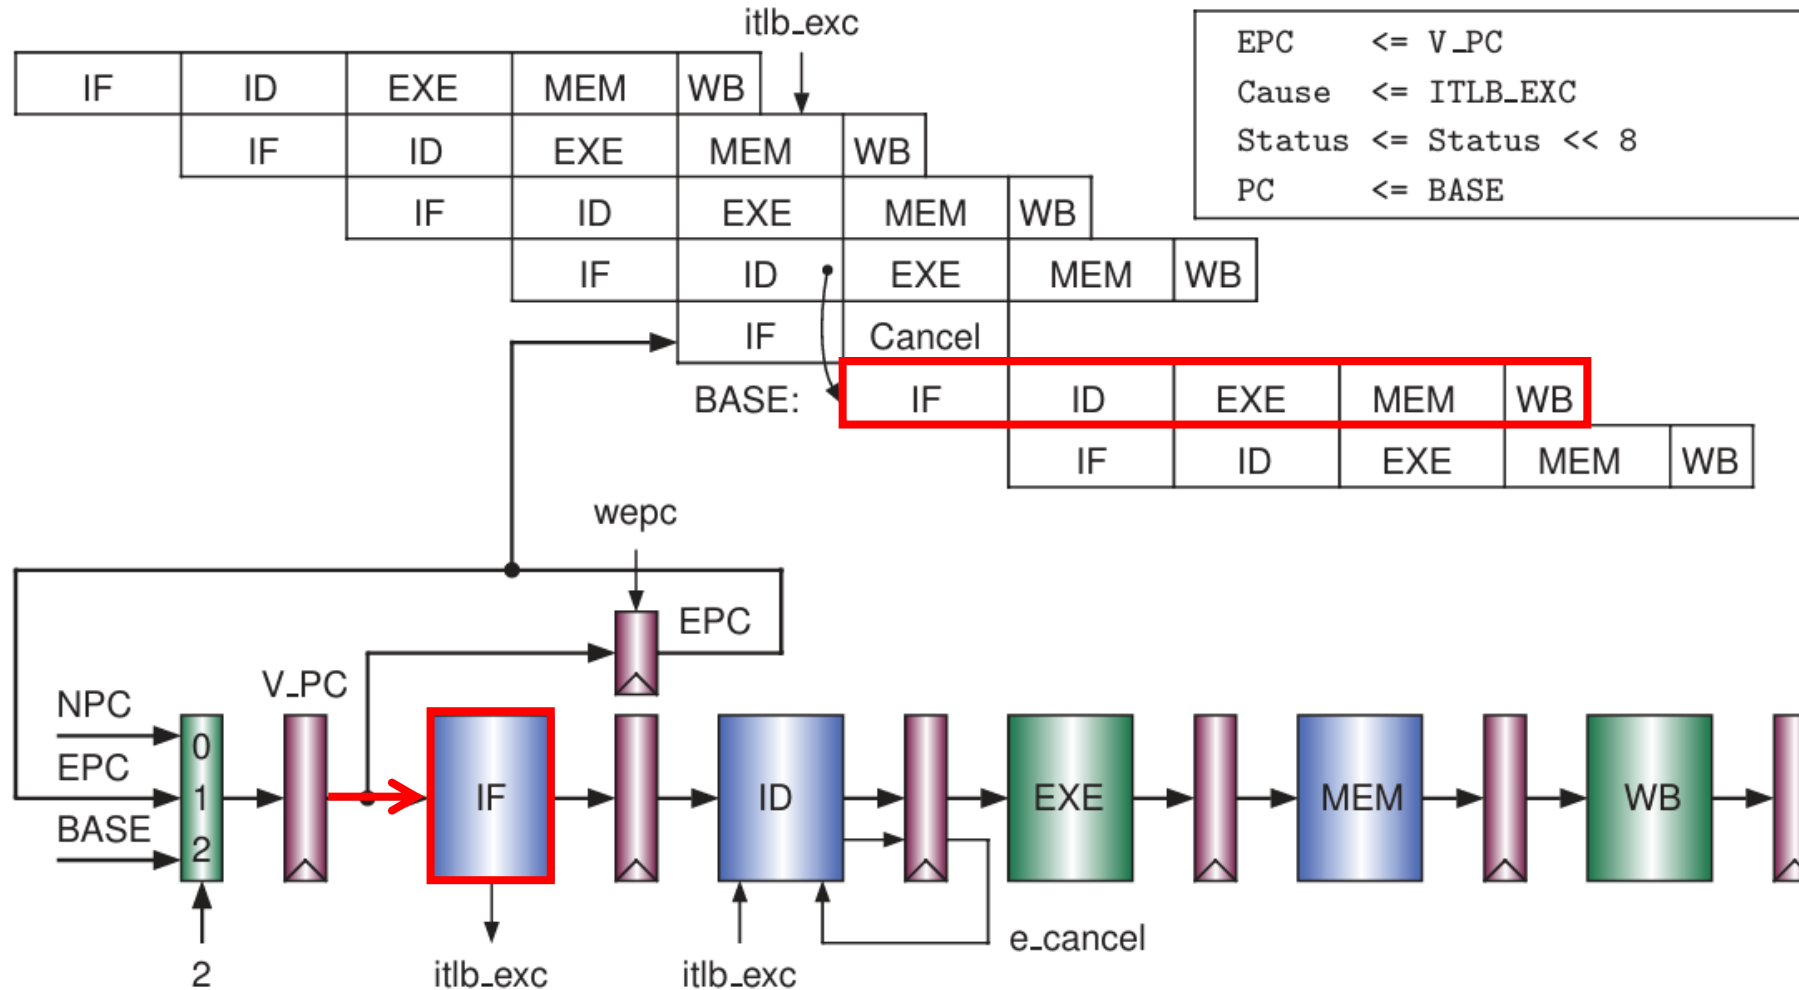
# Design of Pipelined CPU with Caches and TLB

- Registers related to ITLB miss exceptions



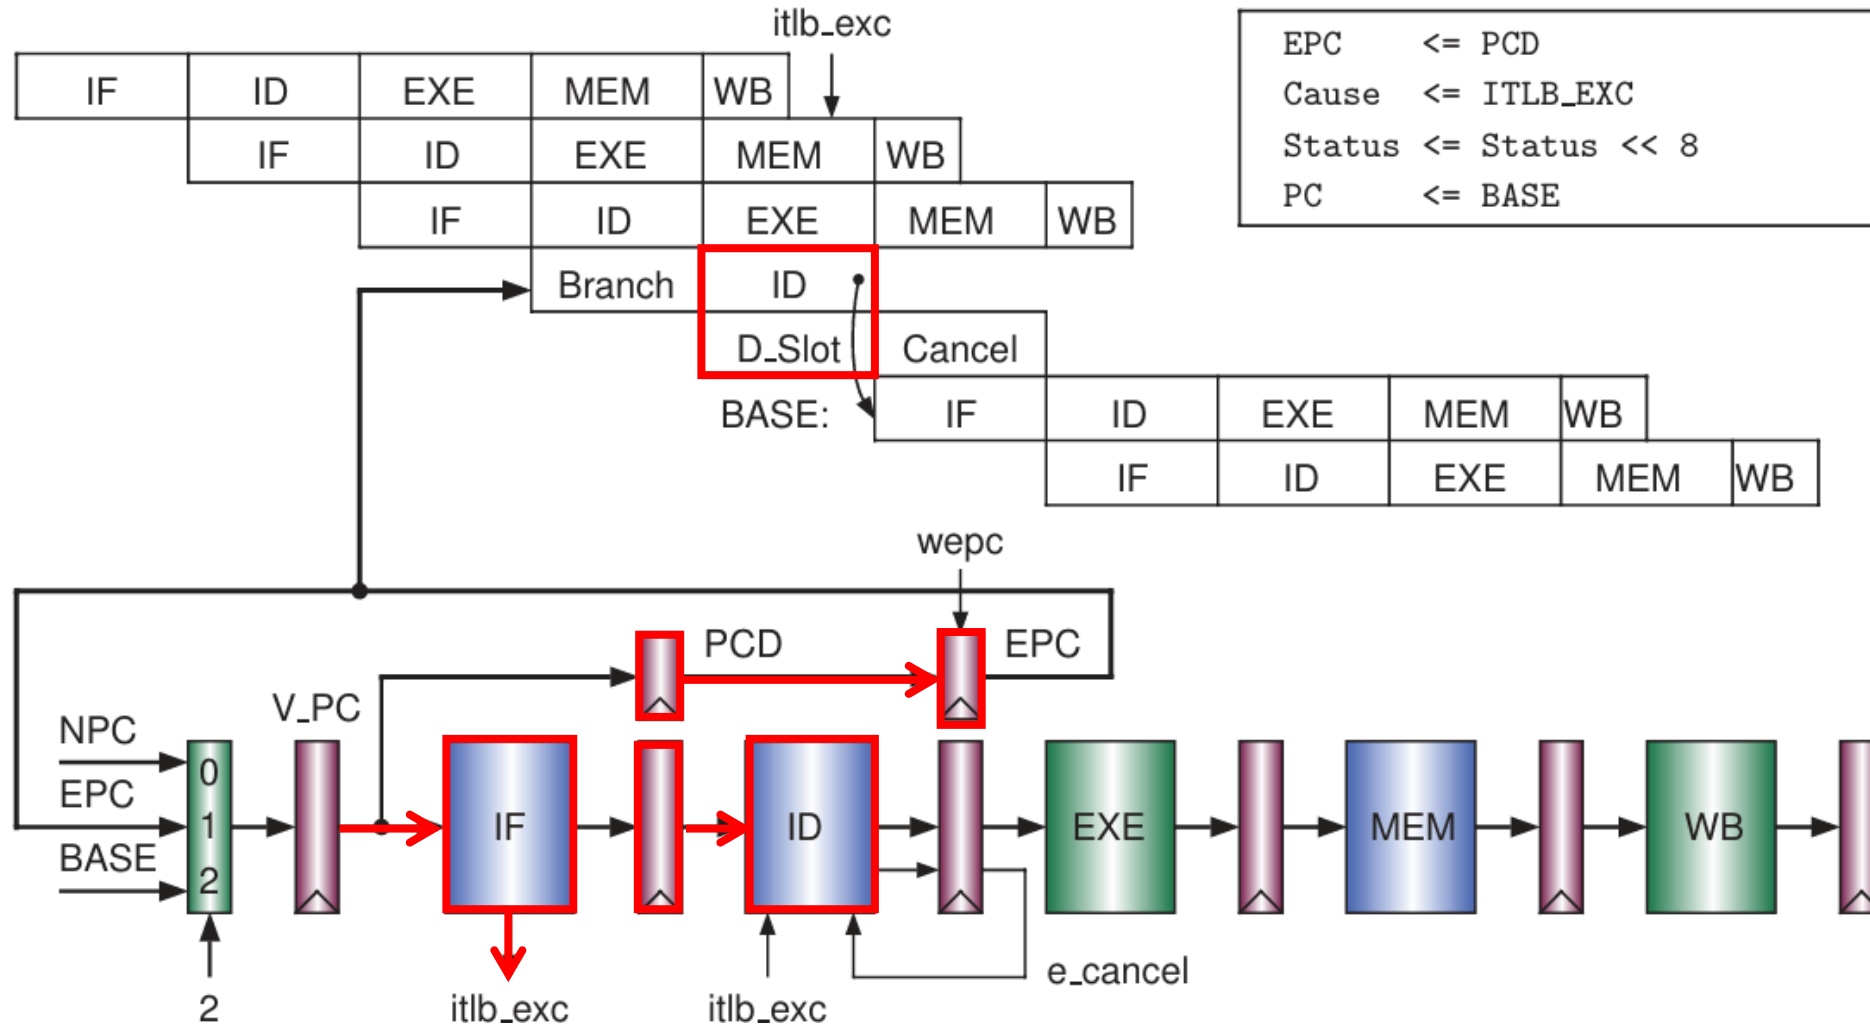
# Design of Pipelined CPU with Caches and TLB

- Registers related to ITLB miss exceptions



# Design of Pipelined CPU with Caches and TLB

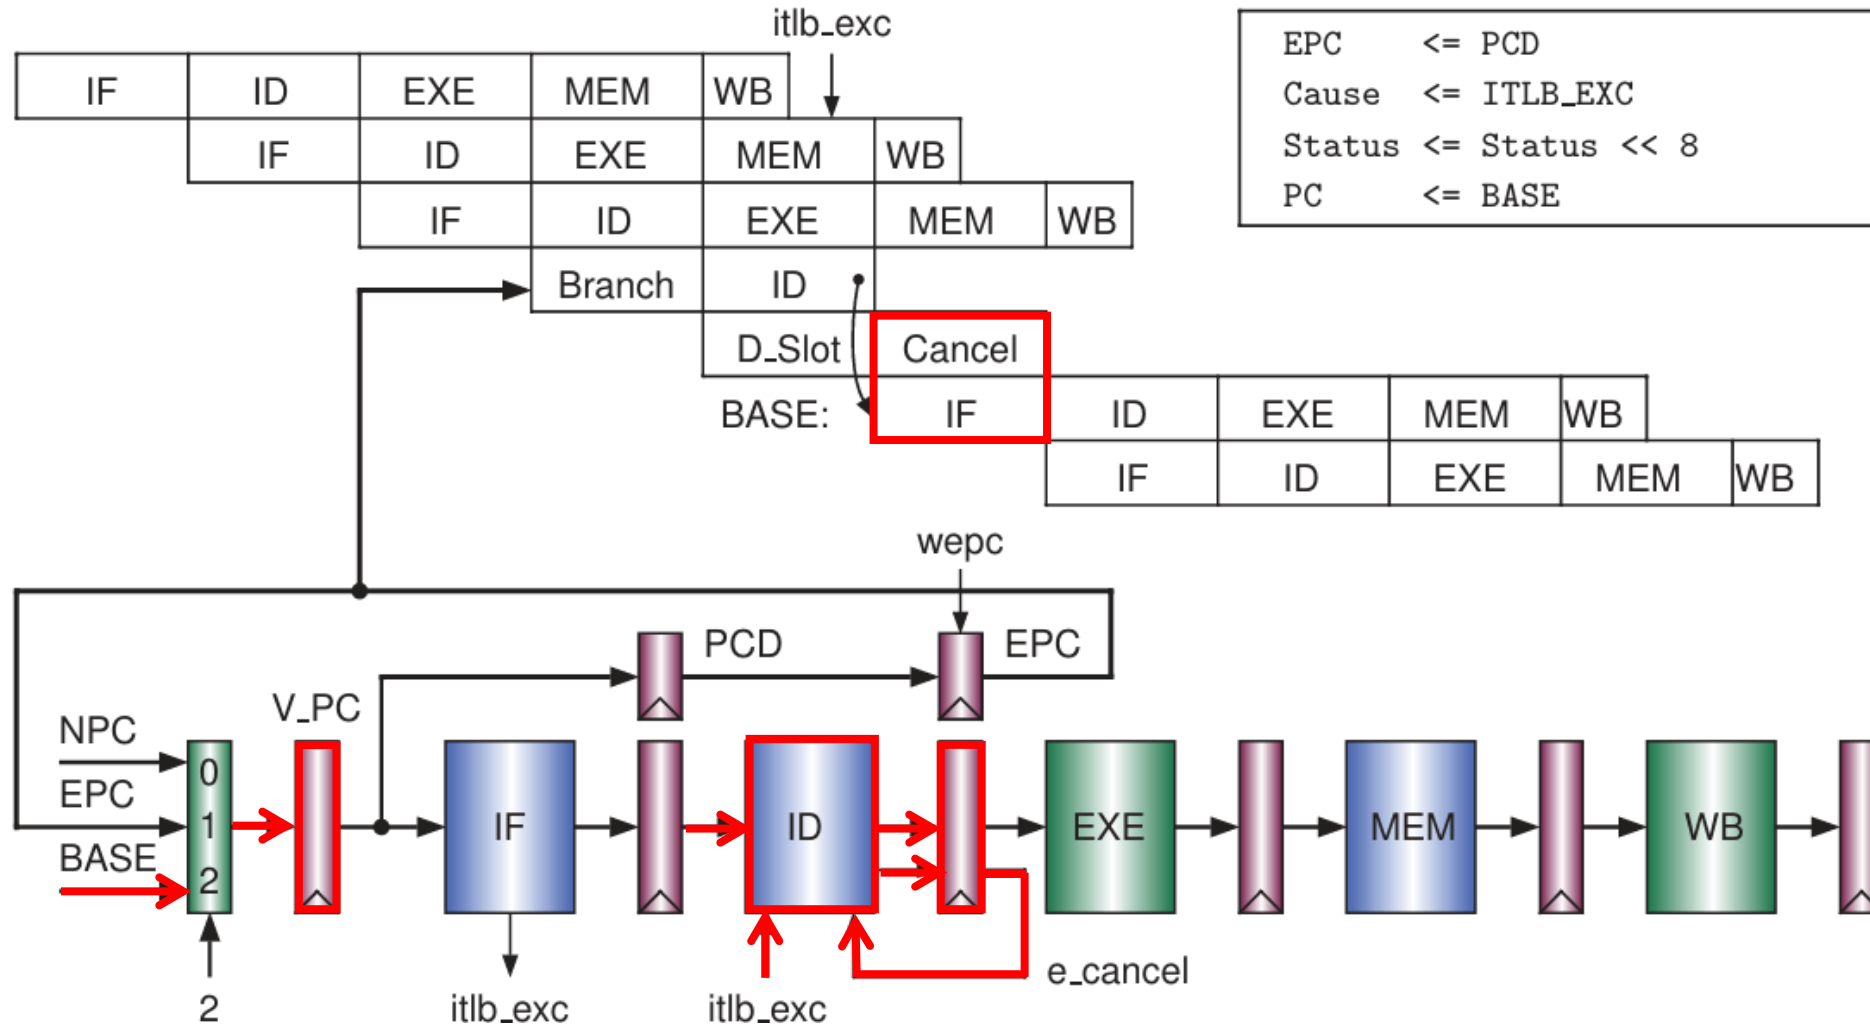
- Registers related to ITLB miss exceptions





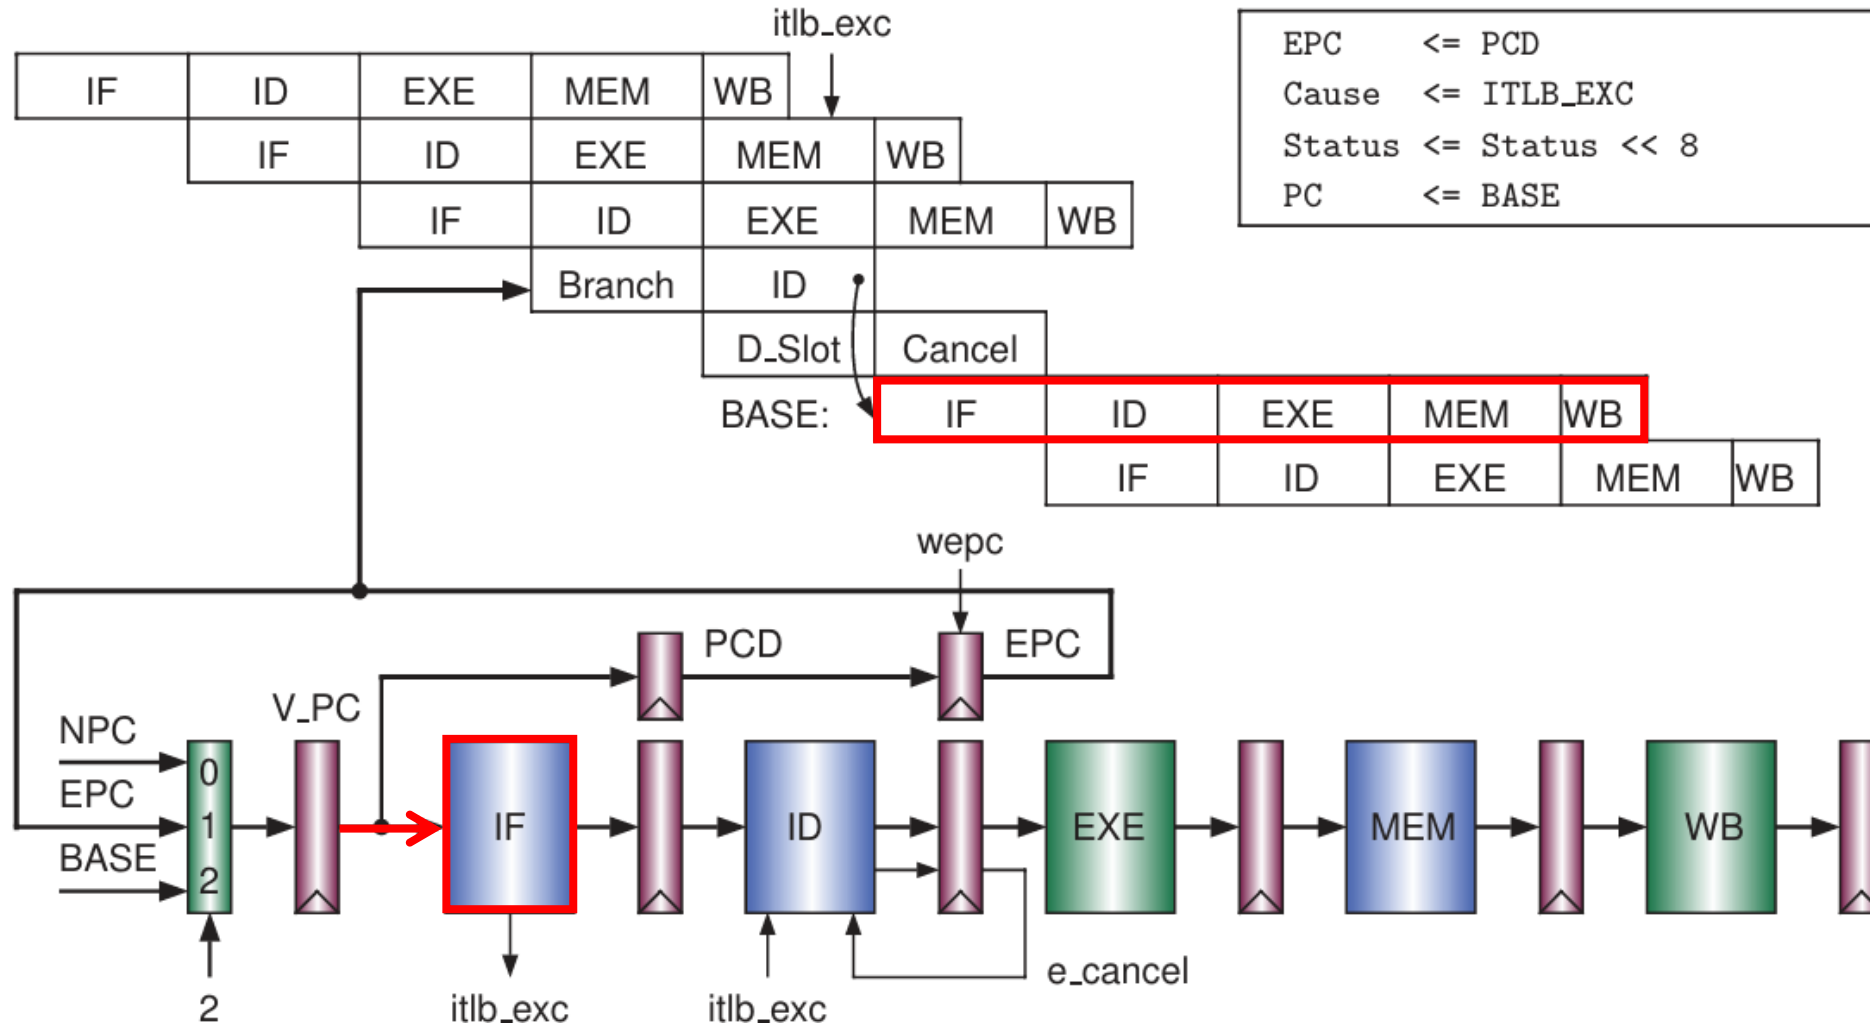
# Design of Pipelined CPU with Caches and TLB

- Registers related to ITLB miss exceptions



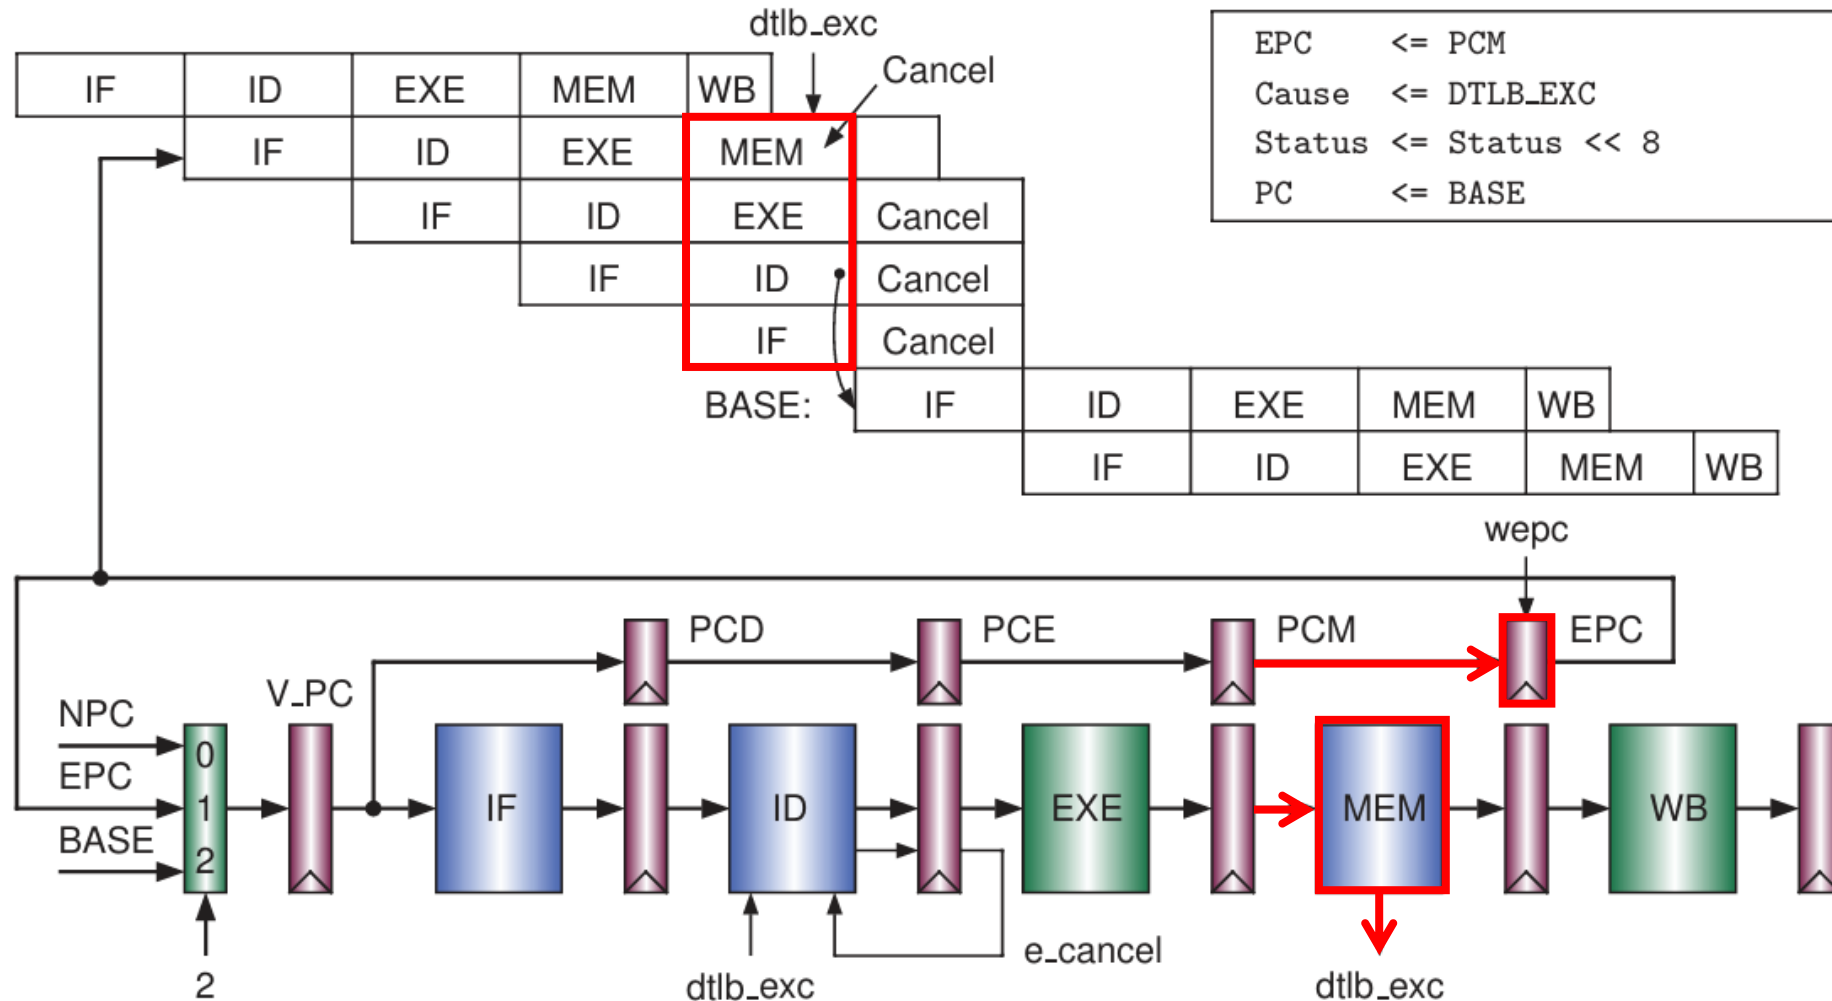
# Design of Pipelined CPU with Caches and TLB

- Registers related to ITLB miss exceptions



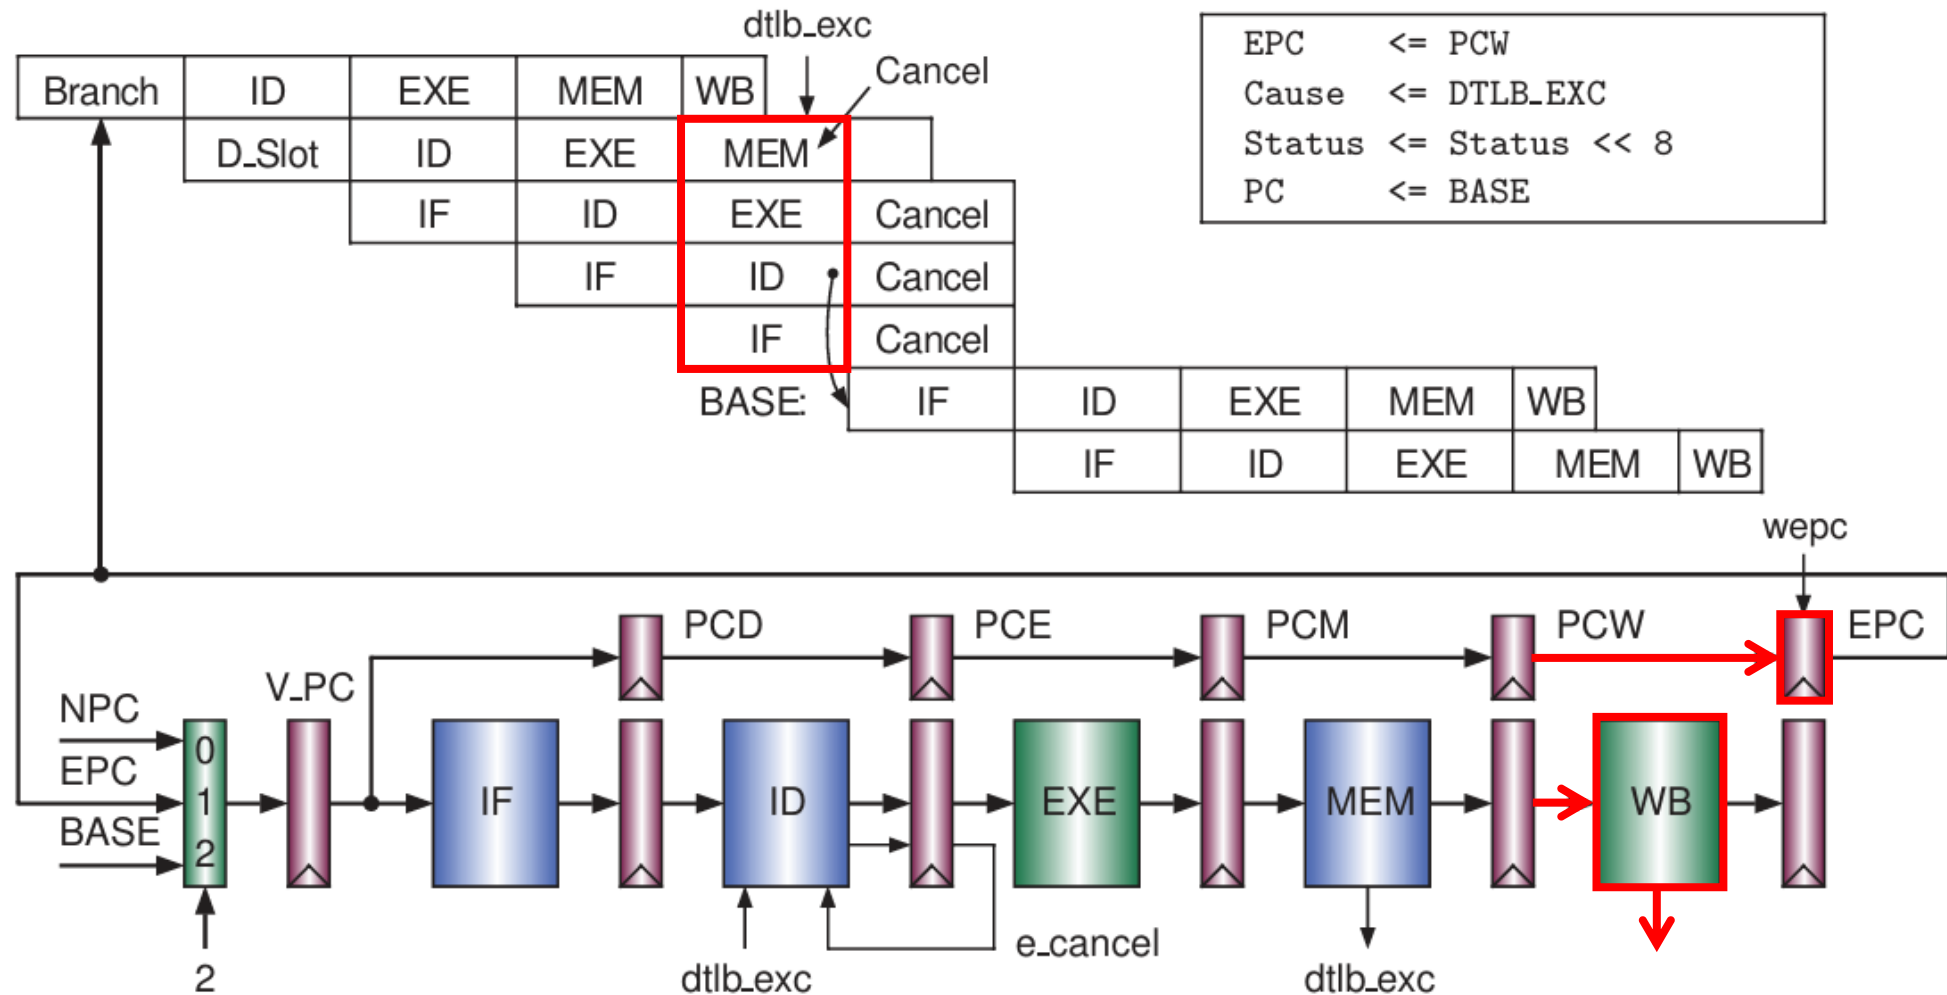
# Design of Pipelined CPU with Caches and TLB

- Registers related to DTLB miss exceptions



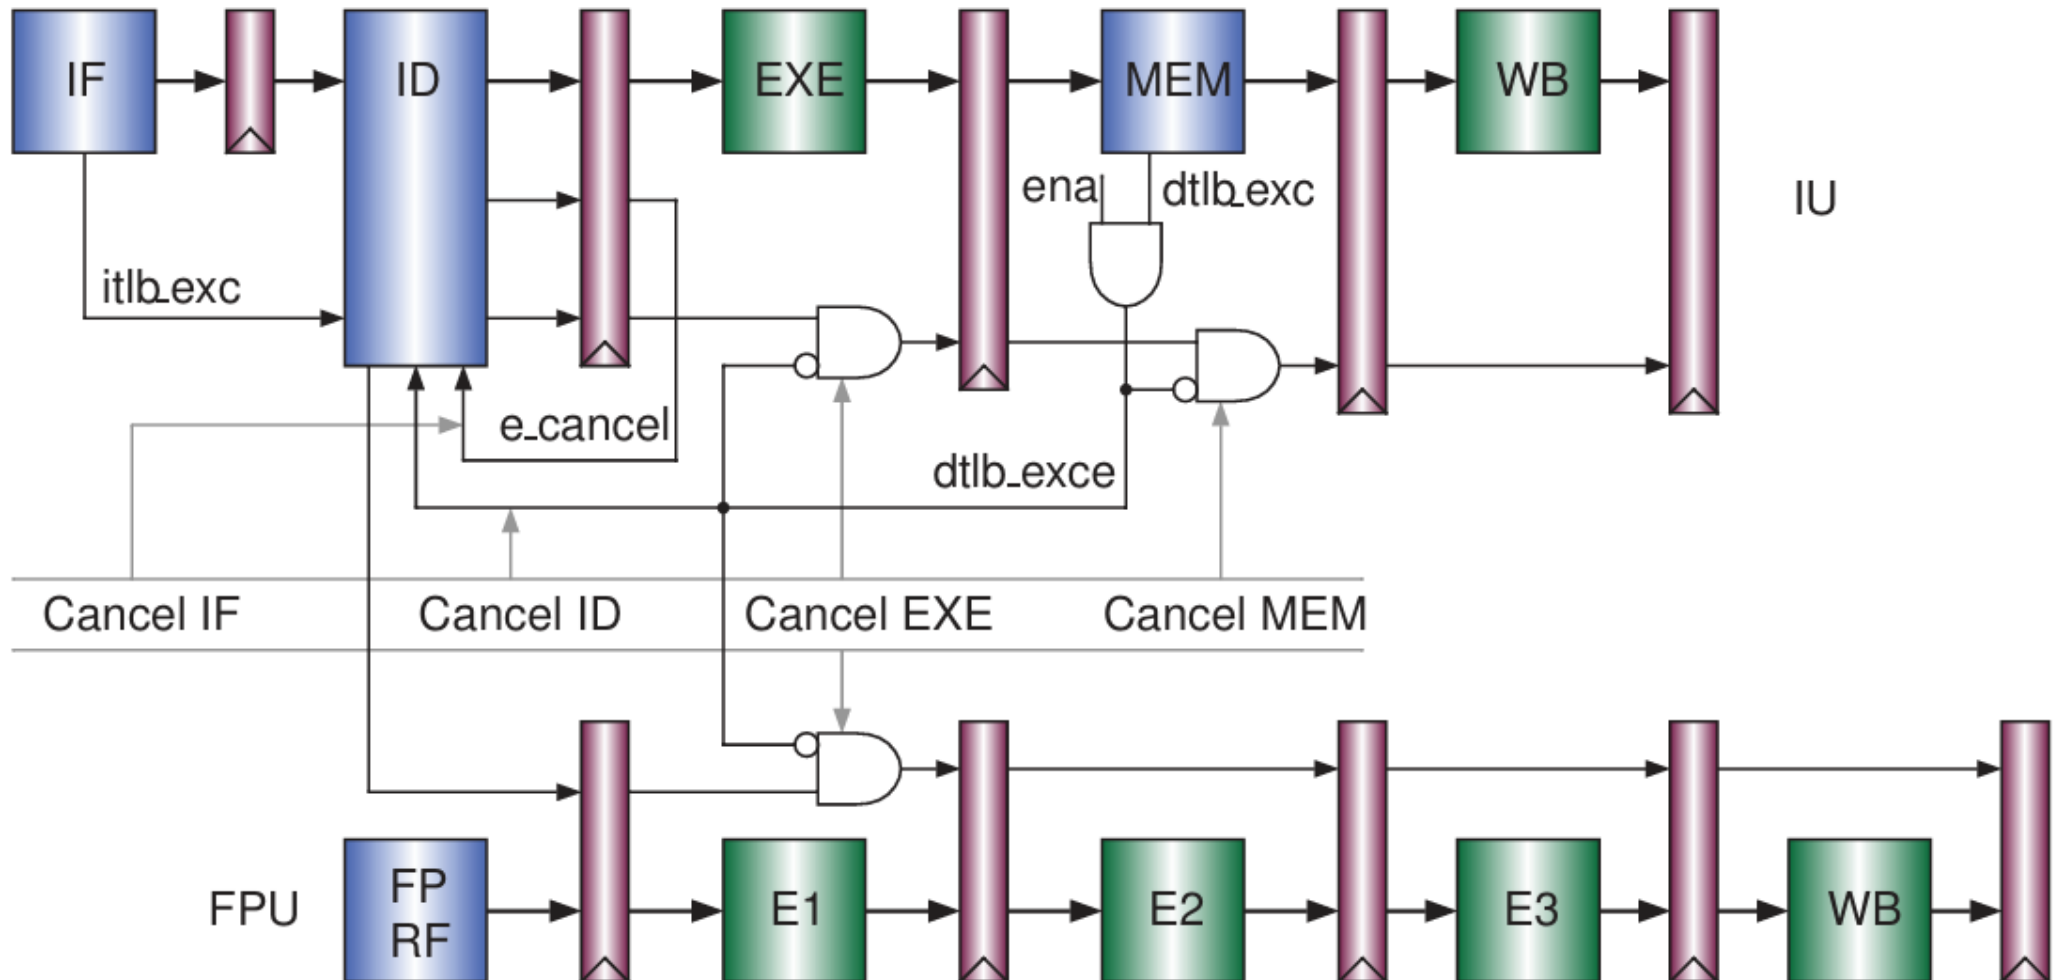
# Design of Pipelined CPU with Caches and TLB

- Registers related to DTLB miss exceptions




# Design of Pipelined CPU with Caches and TLB

- Cancel instructions due to DTLB miss



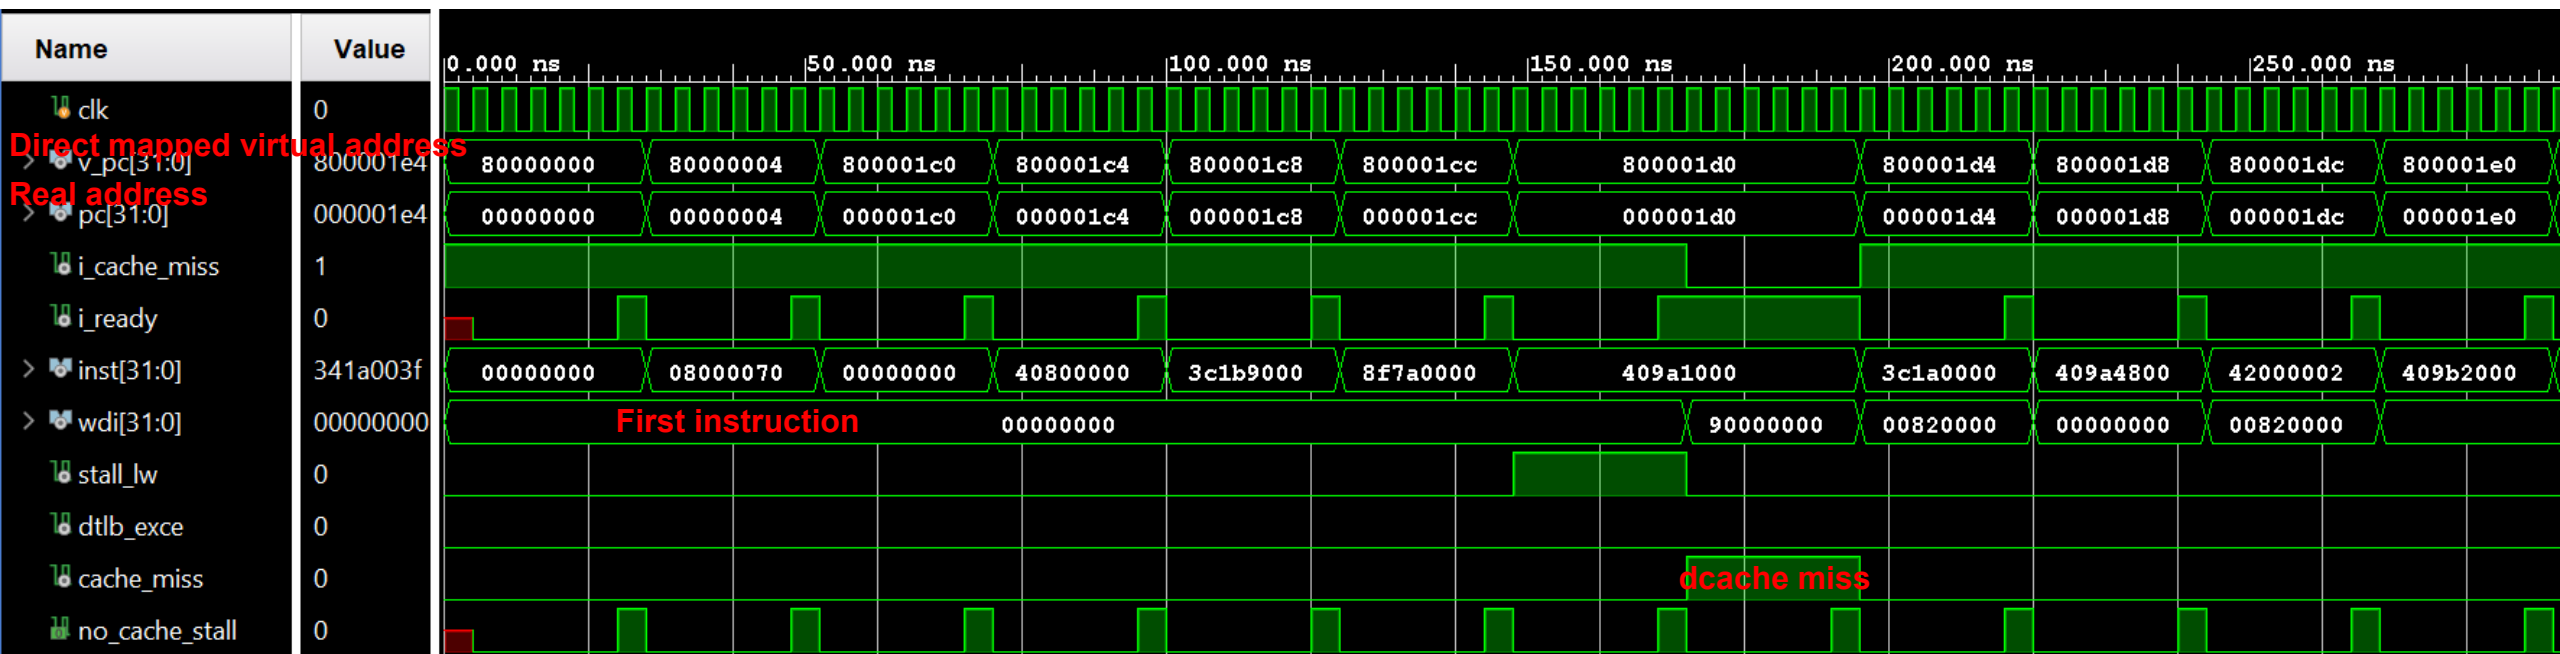
# Design of Pipelined CPU with Caches and TLB

- Structure of CPU with caches and TLBs



# Simulation

## Waveform of initializing ITLB

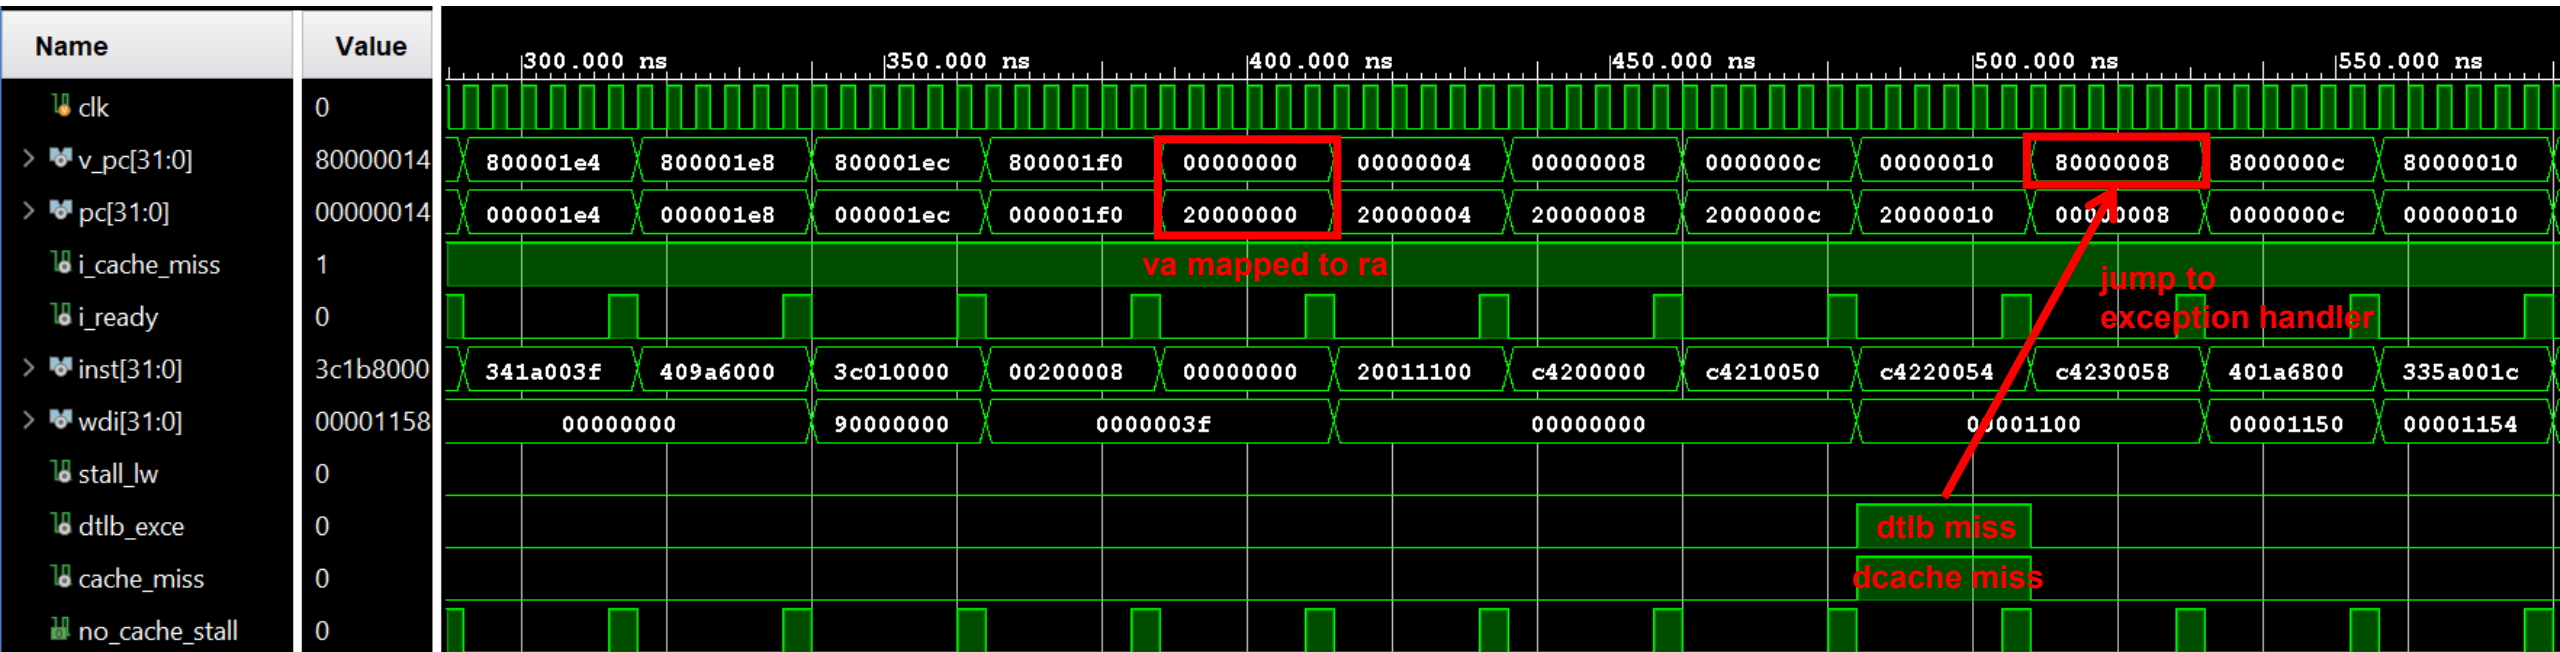


```

80000000 : 08000070; % j initialize_itlb # jump to init itlb %
80000004 : 00000000; % nop %
800001c0 : 40800000; % mtc0 $0, c0_index # c0_index <-- 0 (itlb[0]) %
800001c4 : 3c1b9000; % lui $27, 0x9000 # page table base %
800001c8 : 8f7a0000; % lw $26, 0x0($27) # 1st entry of page table %
800001cc : 409a1000; % mtc0 $26, c0_entry_lo # c0_entrylo <-- v,d,c,pfn %
800001d0 : 3c1a0000; % lui $26, 0x0 # va (=0) for c0_entry_hi %
800001d4 : 409a4800; % mtc0 $26, c0_entry_hi # c0_entry_hi <-- vpn (0) %
800001d8 : 42000002; % tlbwi # write itlb for user prog %
800001dc : 409b2000; % mtc0 $27, c0_context # c0_context <-- ptebase %
800001e0 : 341a003f; % ori $26, $0, 0x3f # enable exceptions %
  
```

# Simulation

## Waveform of dealing with DTLB miss exception

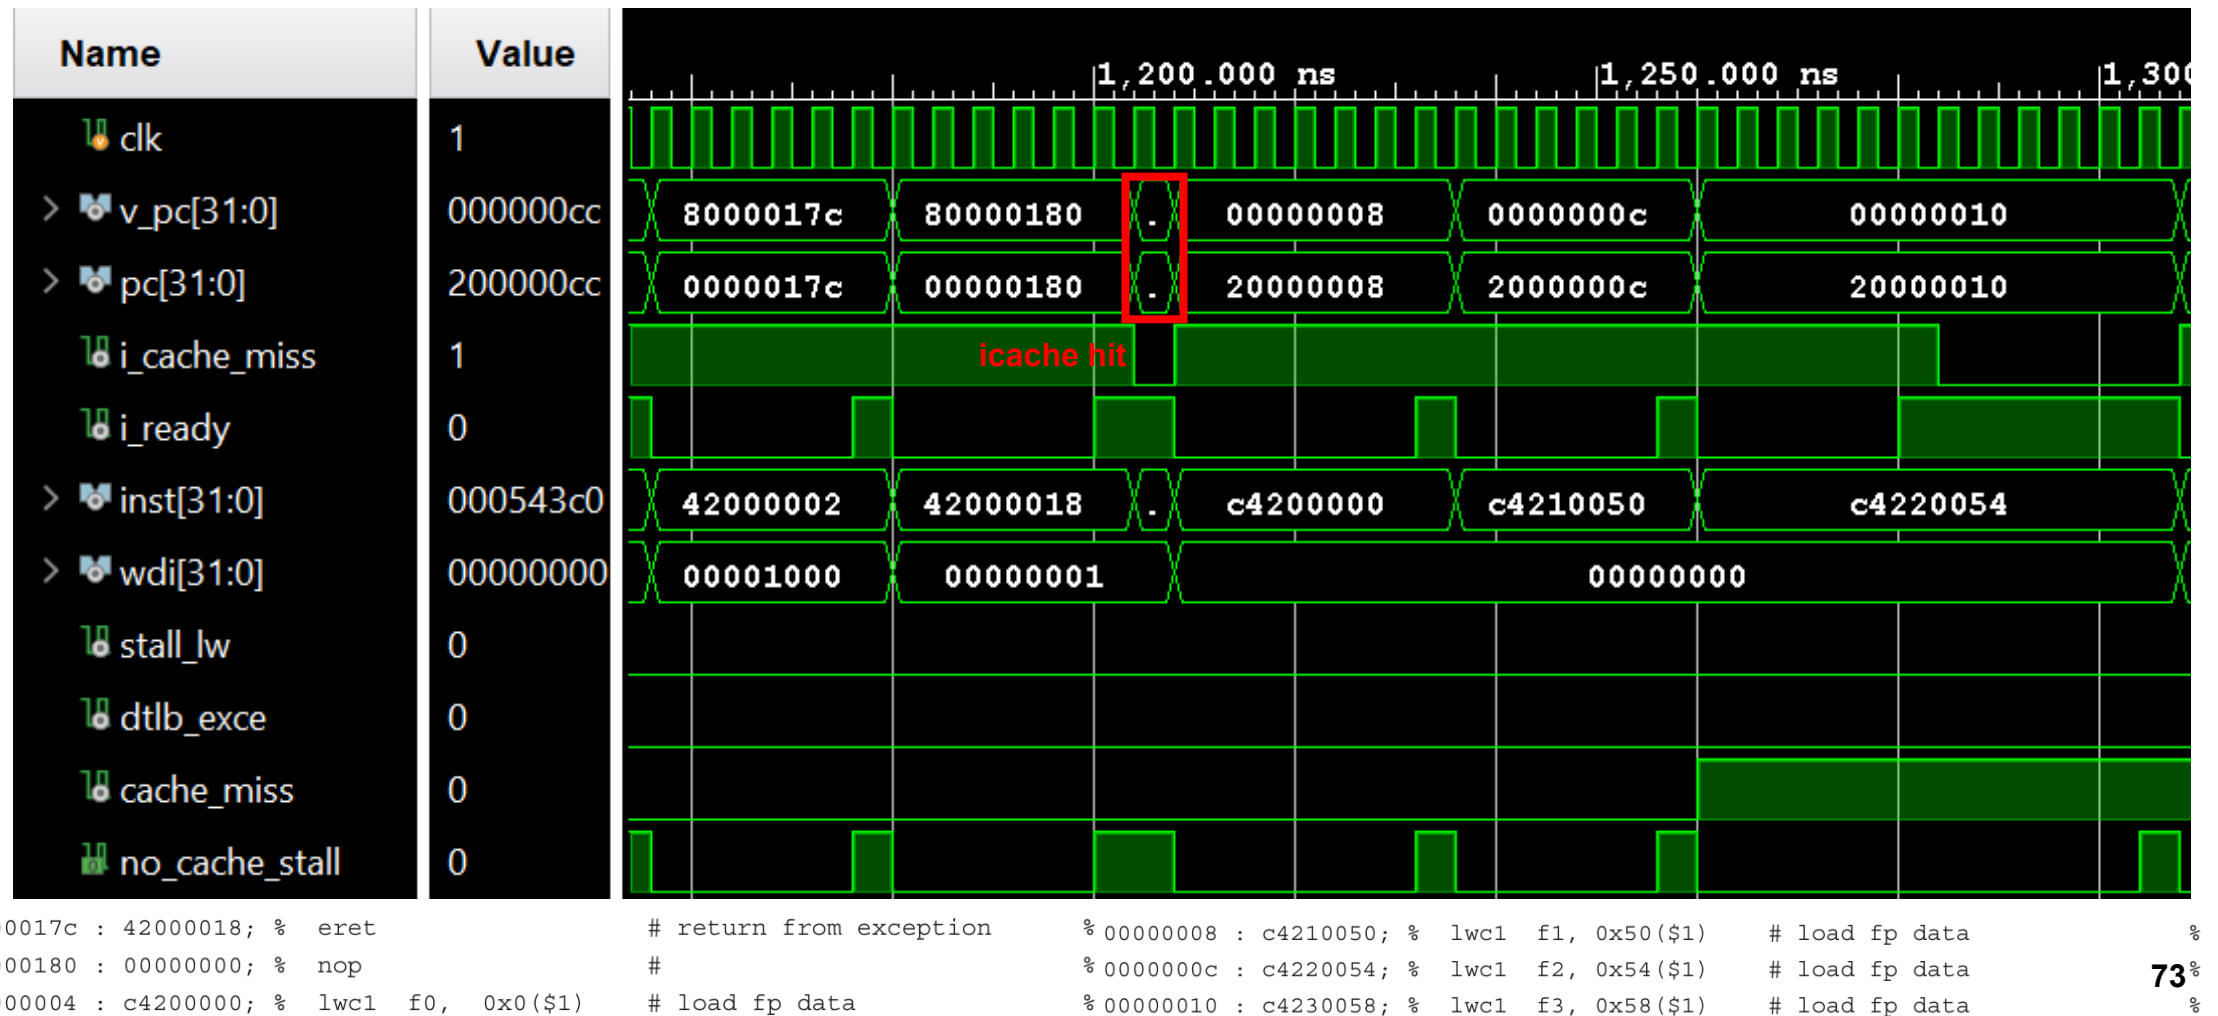


800001e4 : 409a6000; % mtc0 \$26, c0_status # c0_status <-- 0..01111111 %	00000000 : 20011100; % addi \$1, \$0, 0x1100 # address of data[0] %
800001e8 : 3c010000; % lui \$1, 0x0 # va = 0x0000_0000 %	00000004 : c4200000; % lwc1 f0, 0x0(\$1) # load fp data %
800001ec : 00200008; % jr \$1 # jump to user program %	00000008 : c4210050; % lwc1 f1, 0x50(\$1) # load fp data %
800001f0 : 00000000; % nop #	0000000c : c4220054; % lwc1 f2, 0x54(\$1) # load fp data %
	00000010 : c4230058; % lwc1 f3, 0x58(\$1) # load fp data %
	80000008 : 401a6800; % mfc0 \$26, c0_cause # read cp0 cause reg %
	8000000c : 335a001c; % andi \$26, \$26, 0x1c # get exccode, 3 bits %



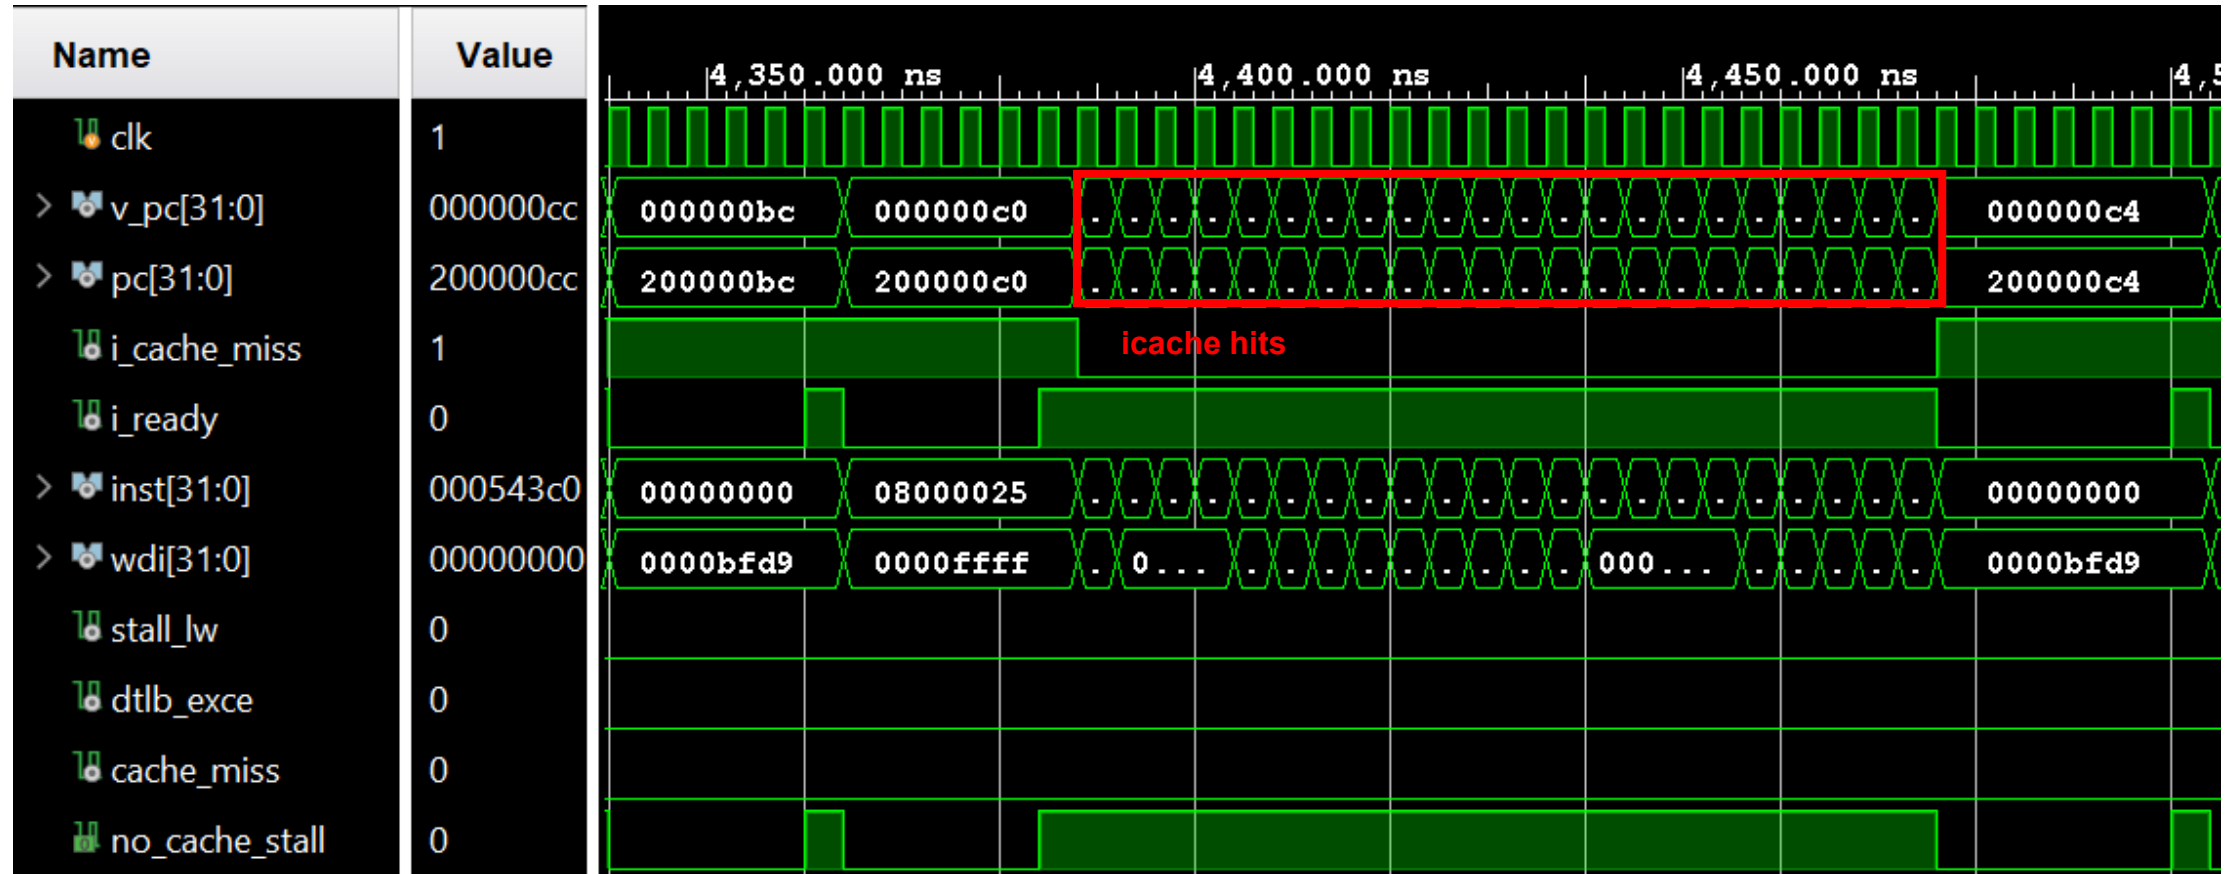
# Simulation

- Waveforms of returning from exception and cache hits



# Simulation

- Waveforms of returning from exception and cache hits



# References

---

[1] Yamin Li, *Computer Principles and Design in Verilog HDL*, Wiley, 2016.